

Approximation Methods



SciPy 2011

Maximum *a posteriori* (MAP)

$$\hat{\theta}_{MAP} = \arg \max P(\theta|y)$$

$$P(\theta|y) = \frac{P(y|\theta)P(\theta)}{P(y)}$$

Bayes Formula

$$P(\theta|y) = \frac{P(y|\theta)P(\theta)}{P(y)}$$

Bayes Formula

$$\Rightarrow \log P(\theta|y) = \log P(y|\theta) + \log P(\theta) + C$$

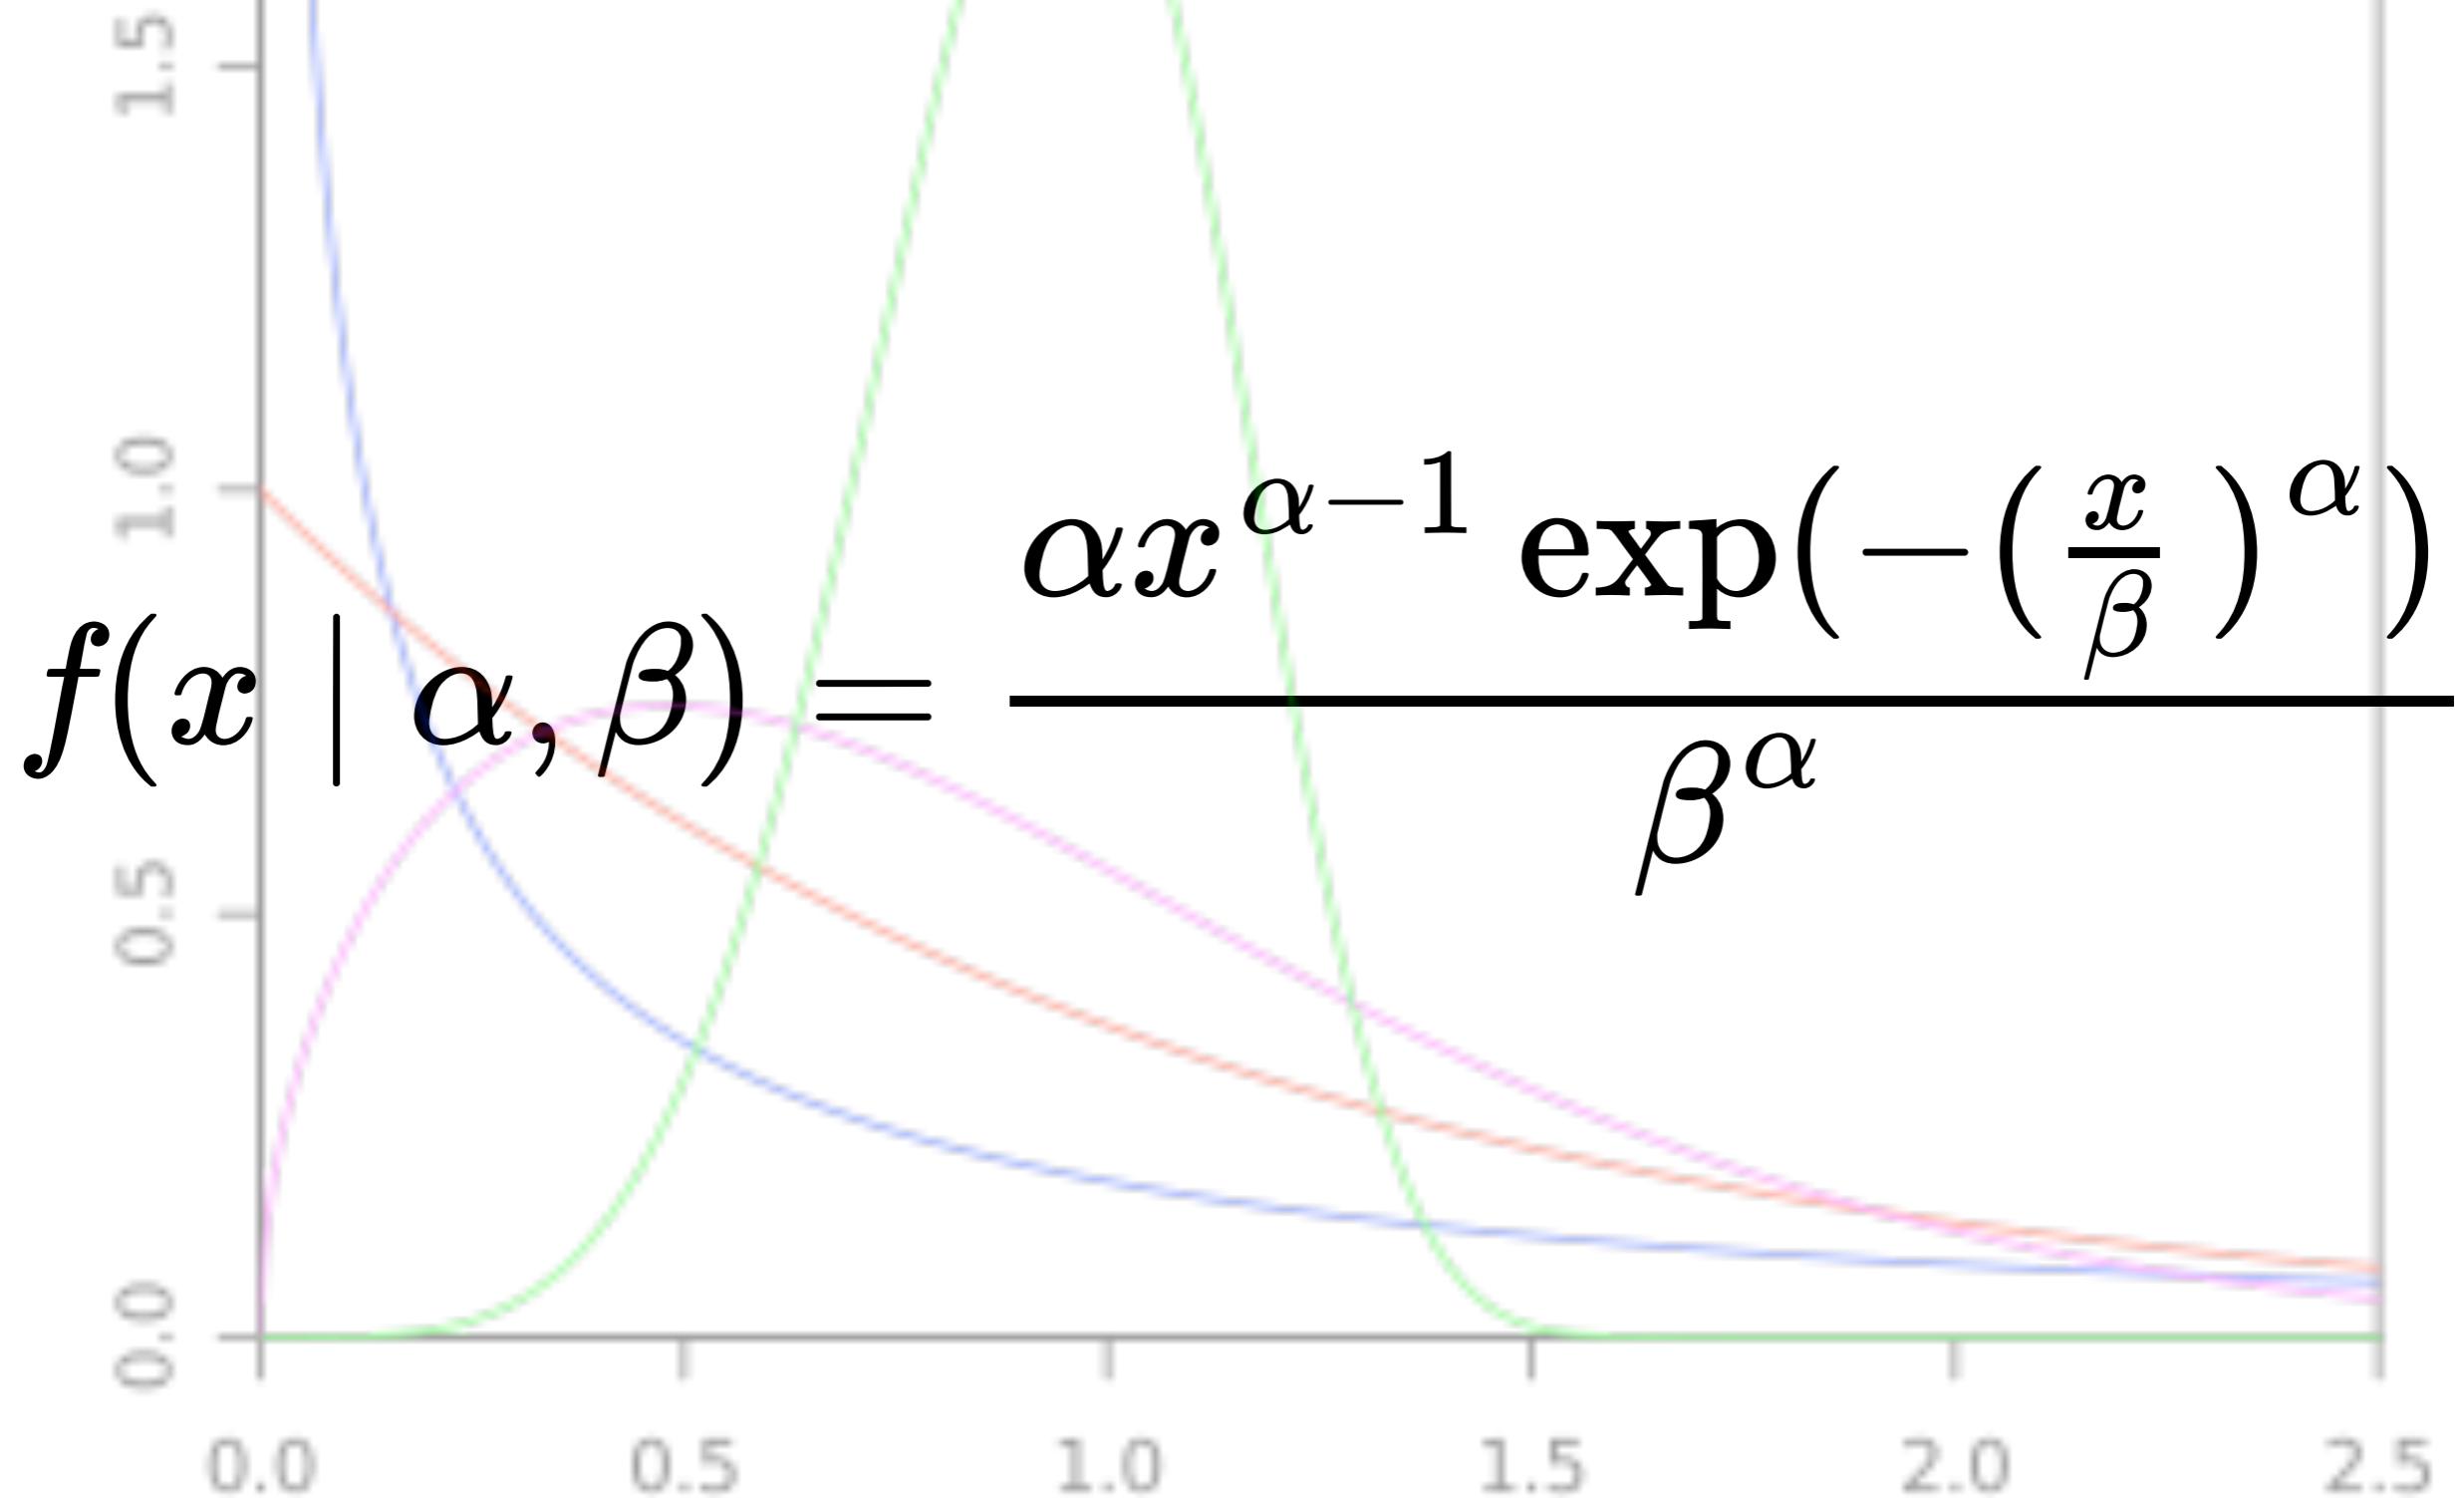
$$P(\theta|y) = \frac{P(y|\theta)P(\theta)}{P(y)}$$

Bayes Formula

$$\Rightarrow \log P(\theta|y) = \log P(y|\theta) + \log P(\theta) + C$$

$$\Rightarrow \hat{\theta}_{MAP} = \arg \max [\log P(y|\theta) + \log P(\theta)]$$

example



Weibull distribution

The screenshot shows a Python code editor window titled "weibull.py". The code is written in Pymc, a probabilistic programming language. It starts by importing Pymc, generating some fake data, and defining parameters alpha and beta. It then creates a dataset using Pymc's rweibull function. The code continues to define a model where parameters a and b are uniform distributions, and a likelihood variable like is a Weibull distribution with observed data from the dataset. The code editor interface includes a status bar at the bottom showing "Line: 12 Column: 42", a Python icon, and various toolbars.

```
1 import pymc-
2 -
3 # Some fake data-
4 alpha = 3-
5 beta = 5-
6 N = 100-
7 dataset = pymc.rweibull(alpha,beta, N)-
8 -
9 # Model-
10 a = pymc.Uniform('a', Lower=0, upper=10, value=5, doc='Weibull alpha
11 . parameter')-
12 b = pymc.Uniform('b', Lower=0, upper=10, value=5, doc='Weibull beta
13 . parameter')-
14 like = pymc.Weibull('like', alpha=a, beta=b, value=dataset, observed=True)
```

```
>>> M = pymc.MAP(weibull)
>>> M.fit()
>>> M.a.value
array(3.2281850543716857)
>>> M.b.value
array(4.7273717106950723)
```

```
N.fit(method='fmin', iterlim=1000, tol=.001):
```

Causes the normal approximation object to fit itself.

method: May be one of the following, from the `scipy.optimize` package:

- `fmin_l_bfgs_b`
- `fmin_ncg`
- `fmin_cg`
- `fmin_powell`
- `fmin`

The screenshot shows a Python code editor window titled "weibull.py". The code is written in Python and uses the PyMC library to create a Weibull distribution model. The code includes importing PyMC, generating a dataset of 100 fake data points from a Weibull distribution with parameters alpha=3 and beta=5, defining two uniform priors for the Weibull parameters (alpha and beta) ranging from 0 to 10, and finally defining a likelihood variable "like" which is a Weibull distribution with observed data points from the dataset.

```
1 import pymc-
2 -
3 # Some fake data-
4 alpha = 3-
5 beta = 5-
6 N = 100-
7 dataset = pymc.rweibull(alpha,beta, N)-
8 -
9 # Model-
10 a = pymc.Uniform('a', Lower=0, upper=10, value=5, doc='Weibull alpha
11 . parameter')-
12 b = pymc.Uniform('b', Lower=0, upper=10, value=5, doc='Weibull beta
13 . parameter')-
14 like = pymc.Weibull('like', alpha=a, beta=b, value=dataset, observed=True)
```

Line: 12 Column: 42 Python Soft Tabs: 4

```
>>> M.AIC
369.57863603595712
>>> M.BIC
374.78897640793332
```

Normal approximation

$$P(\theta|y) \approx N(E(\theta|y), C(\theta|y))$$

Taylor expansion

$$\log f(\theta|y) \approx \log f(\hat{\theta}|y) + \frac{1}{2} (\theta - \hat{\theta})' \left[\frac{d^2}{d\theta^2} \log f(\theta|y) \right]_{\theta=\hat{\theta}} (\theta - \hat{\theta})$$

$$I(\hat{\theta}) = \left[\frac{d^2}{d\theta^2} \log f(\theta|y) \right]_{\theta=\hat{\theta}}$$

Fisher information

Modal approximation

$$f(\theta|y) \approx f(\hat{\theta}|y) \exp\left[-\frac{1}{2} (\theta - \hat{\theta})' I(\hat{\theta})(\theta - \hat{\theta})\right]$$

$$P(\theta|y) \approx N(\hat{\theta}|y, I(\hat{\theta})^{-1})$$

The screenshot shows a Python code editor window titled "weibull.py". The code uses the PyMC library to create a Weibull distribution model. The editor has a dark theme with syntax highlighting for different code elements. The status bar at the bottom shows "Line: 12 Column: 17" and "Python".

```
1 import pymc-
2 -
3 # Some fake data-
4 alpha = 3-
5 beta = 5-
6 N = 100-
7 dataset = pymc.rweibull(alpha,beta, N)-
8 -
9 # Model-
10 a = pymc.Uniform('a', Lower=0, upper=10, value=5, doc='Weibull alpha
11 . parameter')-
12 b = pymc.Uniform('b', Lower=0, upper=10, value=5, doc='Weibull beta
13 . parameter')-
14 like = pymc.Weibull('like', alpha=a, beta=b, value=dataset, observed=True)
```

```
>>> N = pymc.NormApprox(weibull)
>>> N.fit()
>>> N.mu[ [N.a, N.b] ]
array([ 3.22820192,  4.72739324])
>>> N.C[ [N.a, N.b] ]
matrix([[ 0.06548929,  0.01247454],
       [ 0.01247454,  0.02382098]])
```

Extending PyMC



SciPy 2011

PyMC Objects

Node

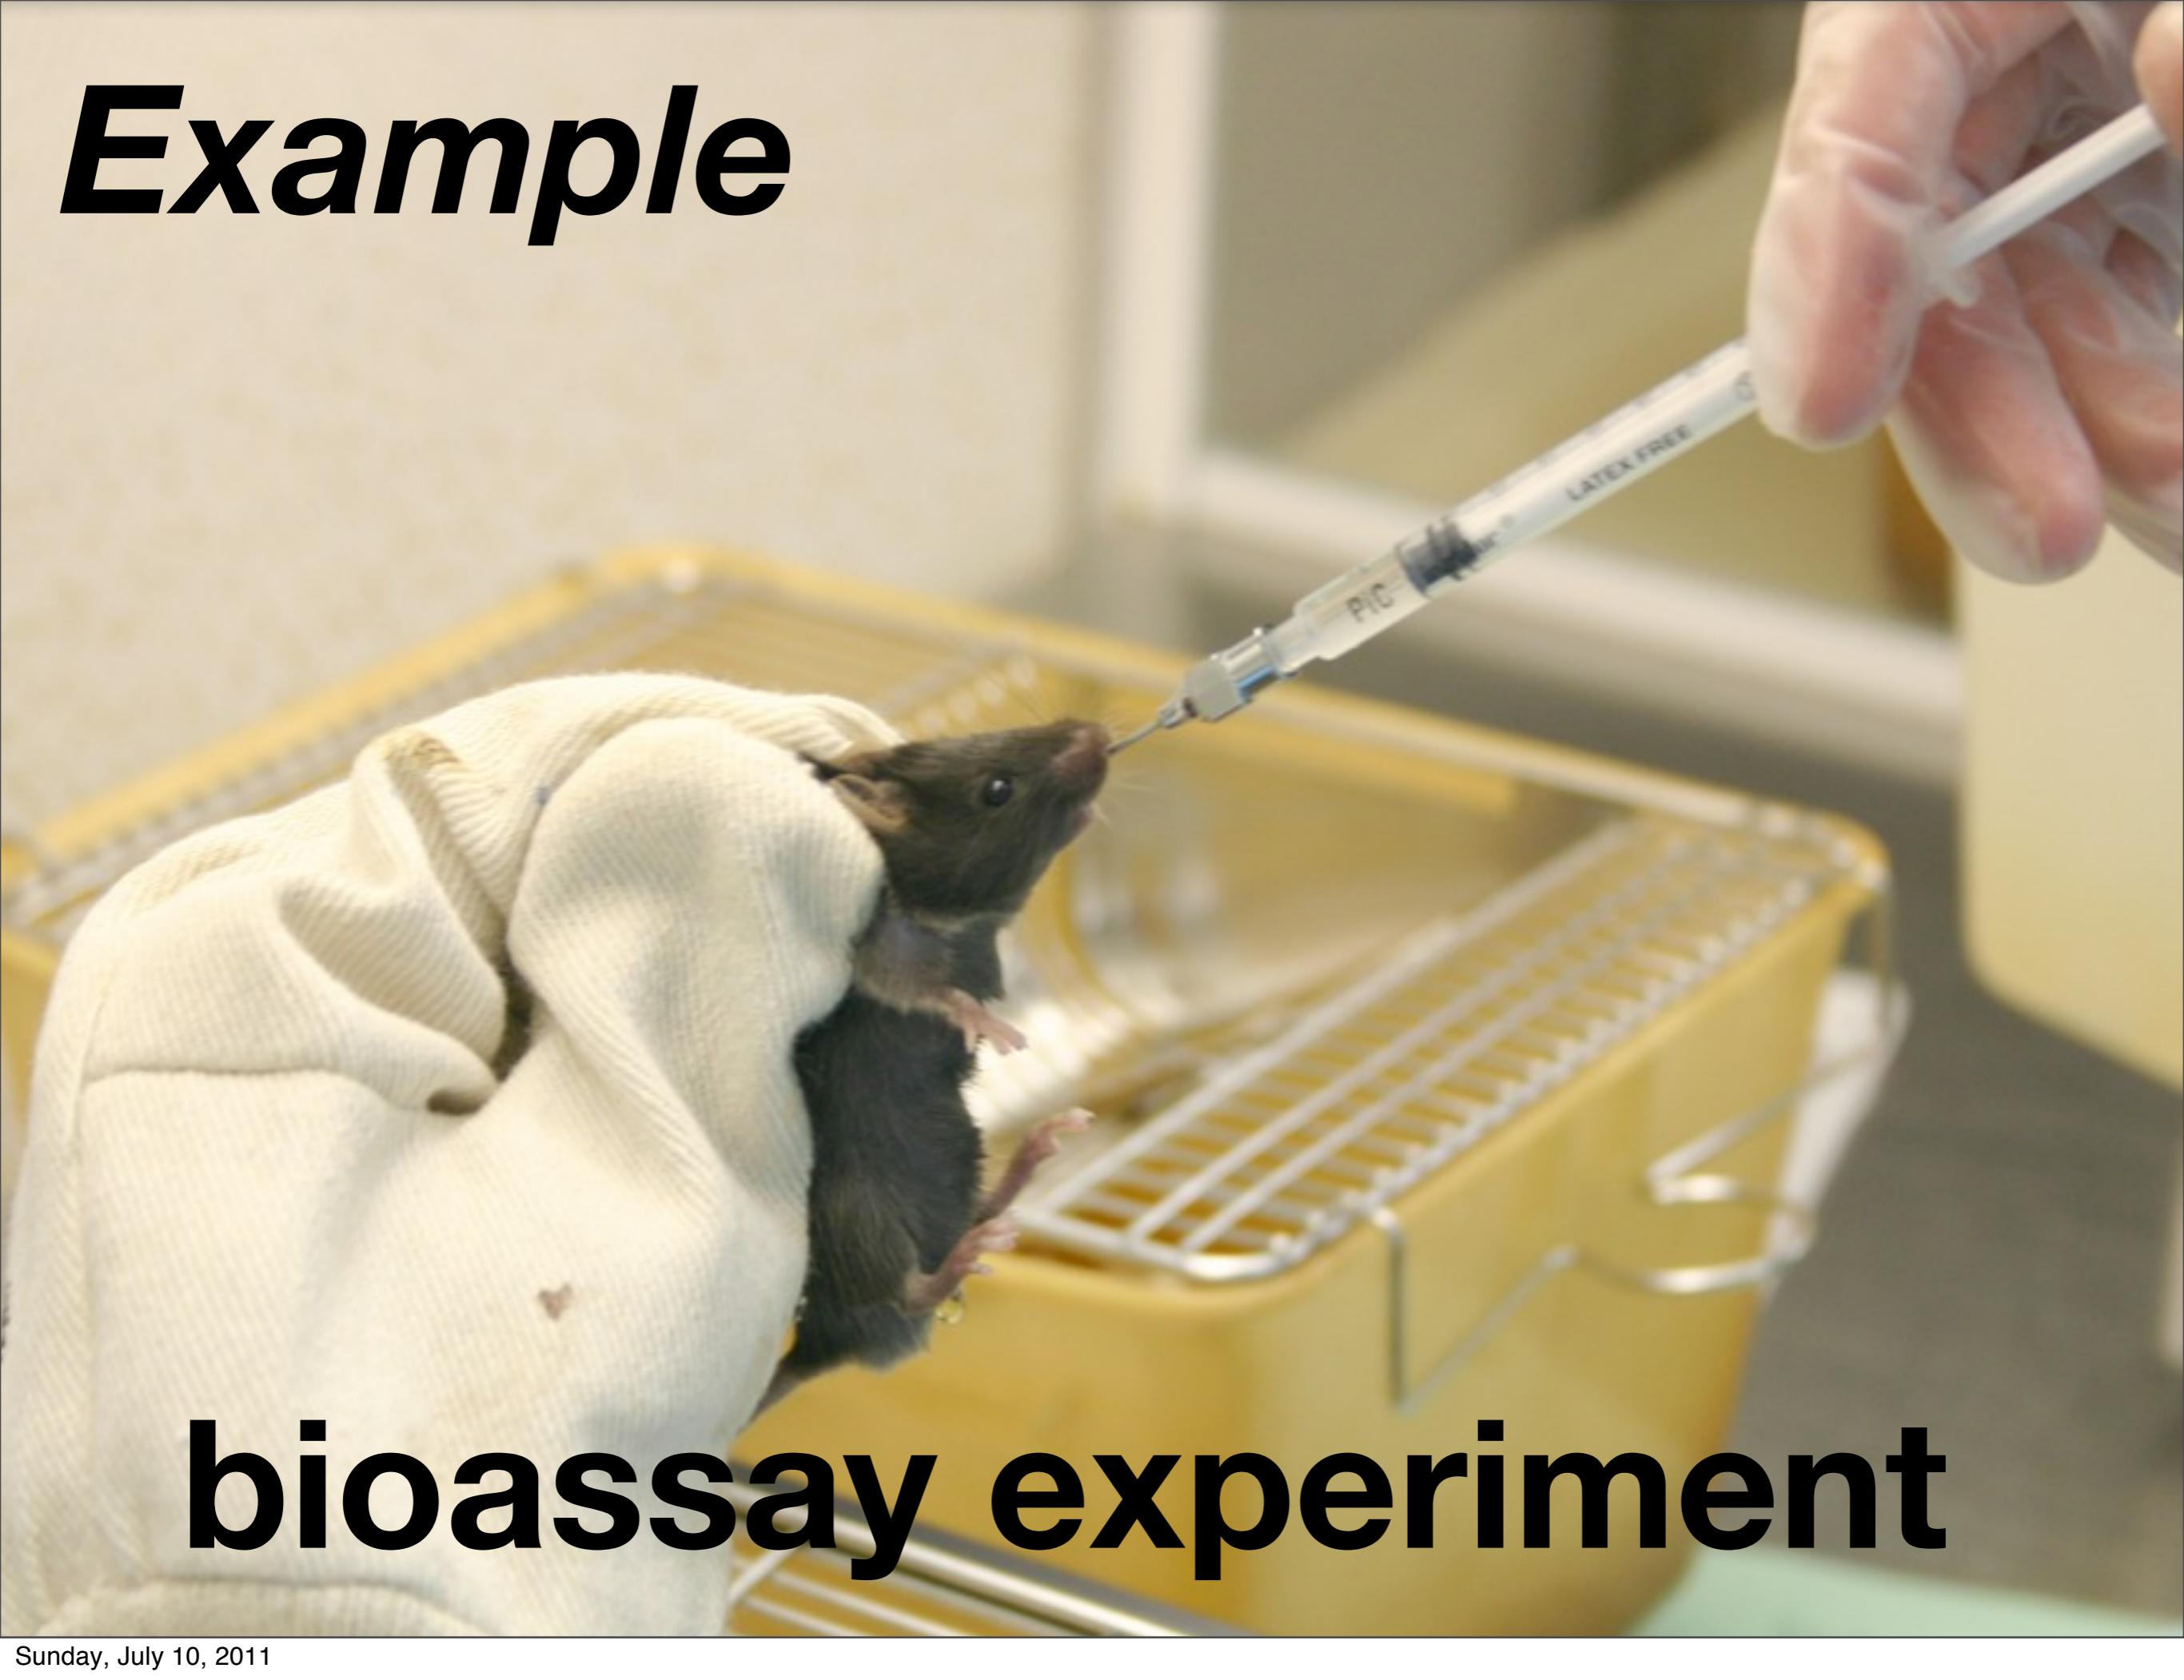
Stochastic

Model

Sampler

StepMethod

Example



bioassay experiment

Dose <i>(log g/ml)</i>	Animals	Deaths
-0.863	5	0
-0.296	5	1
-0.053	5	3
0.727	5	5

bioassay data

$$y_i|\theta_i \sim \text{Bin}(n_i,\theta_i)$$

$$\text{logit}(\theta_i) = \alpha + \beta x_i$$

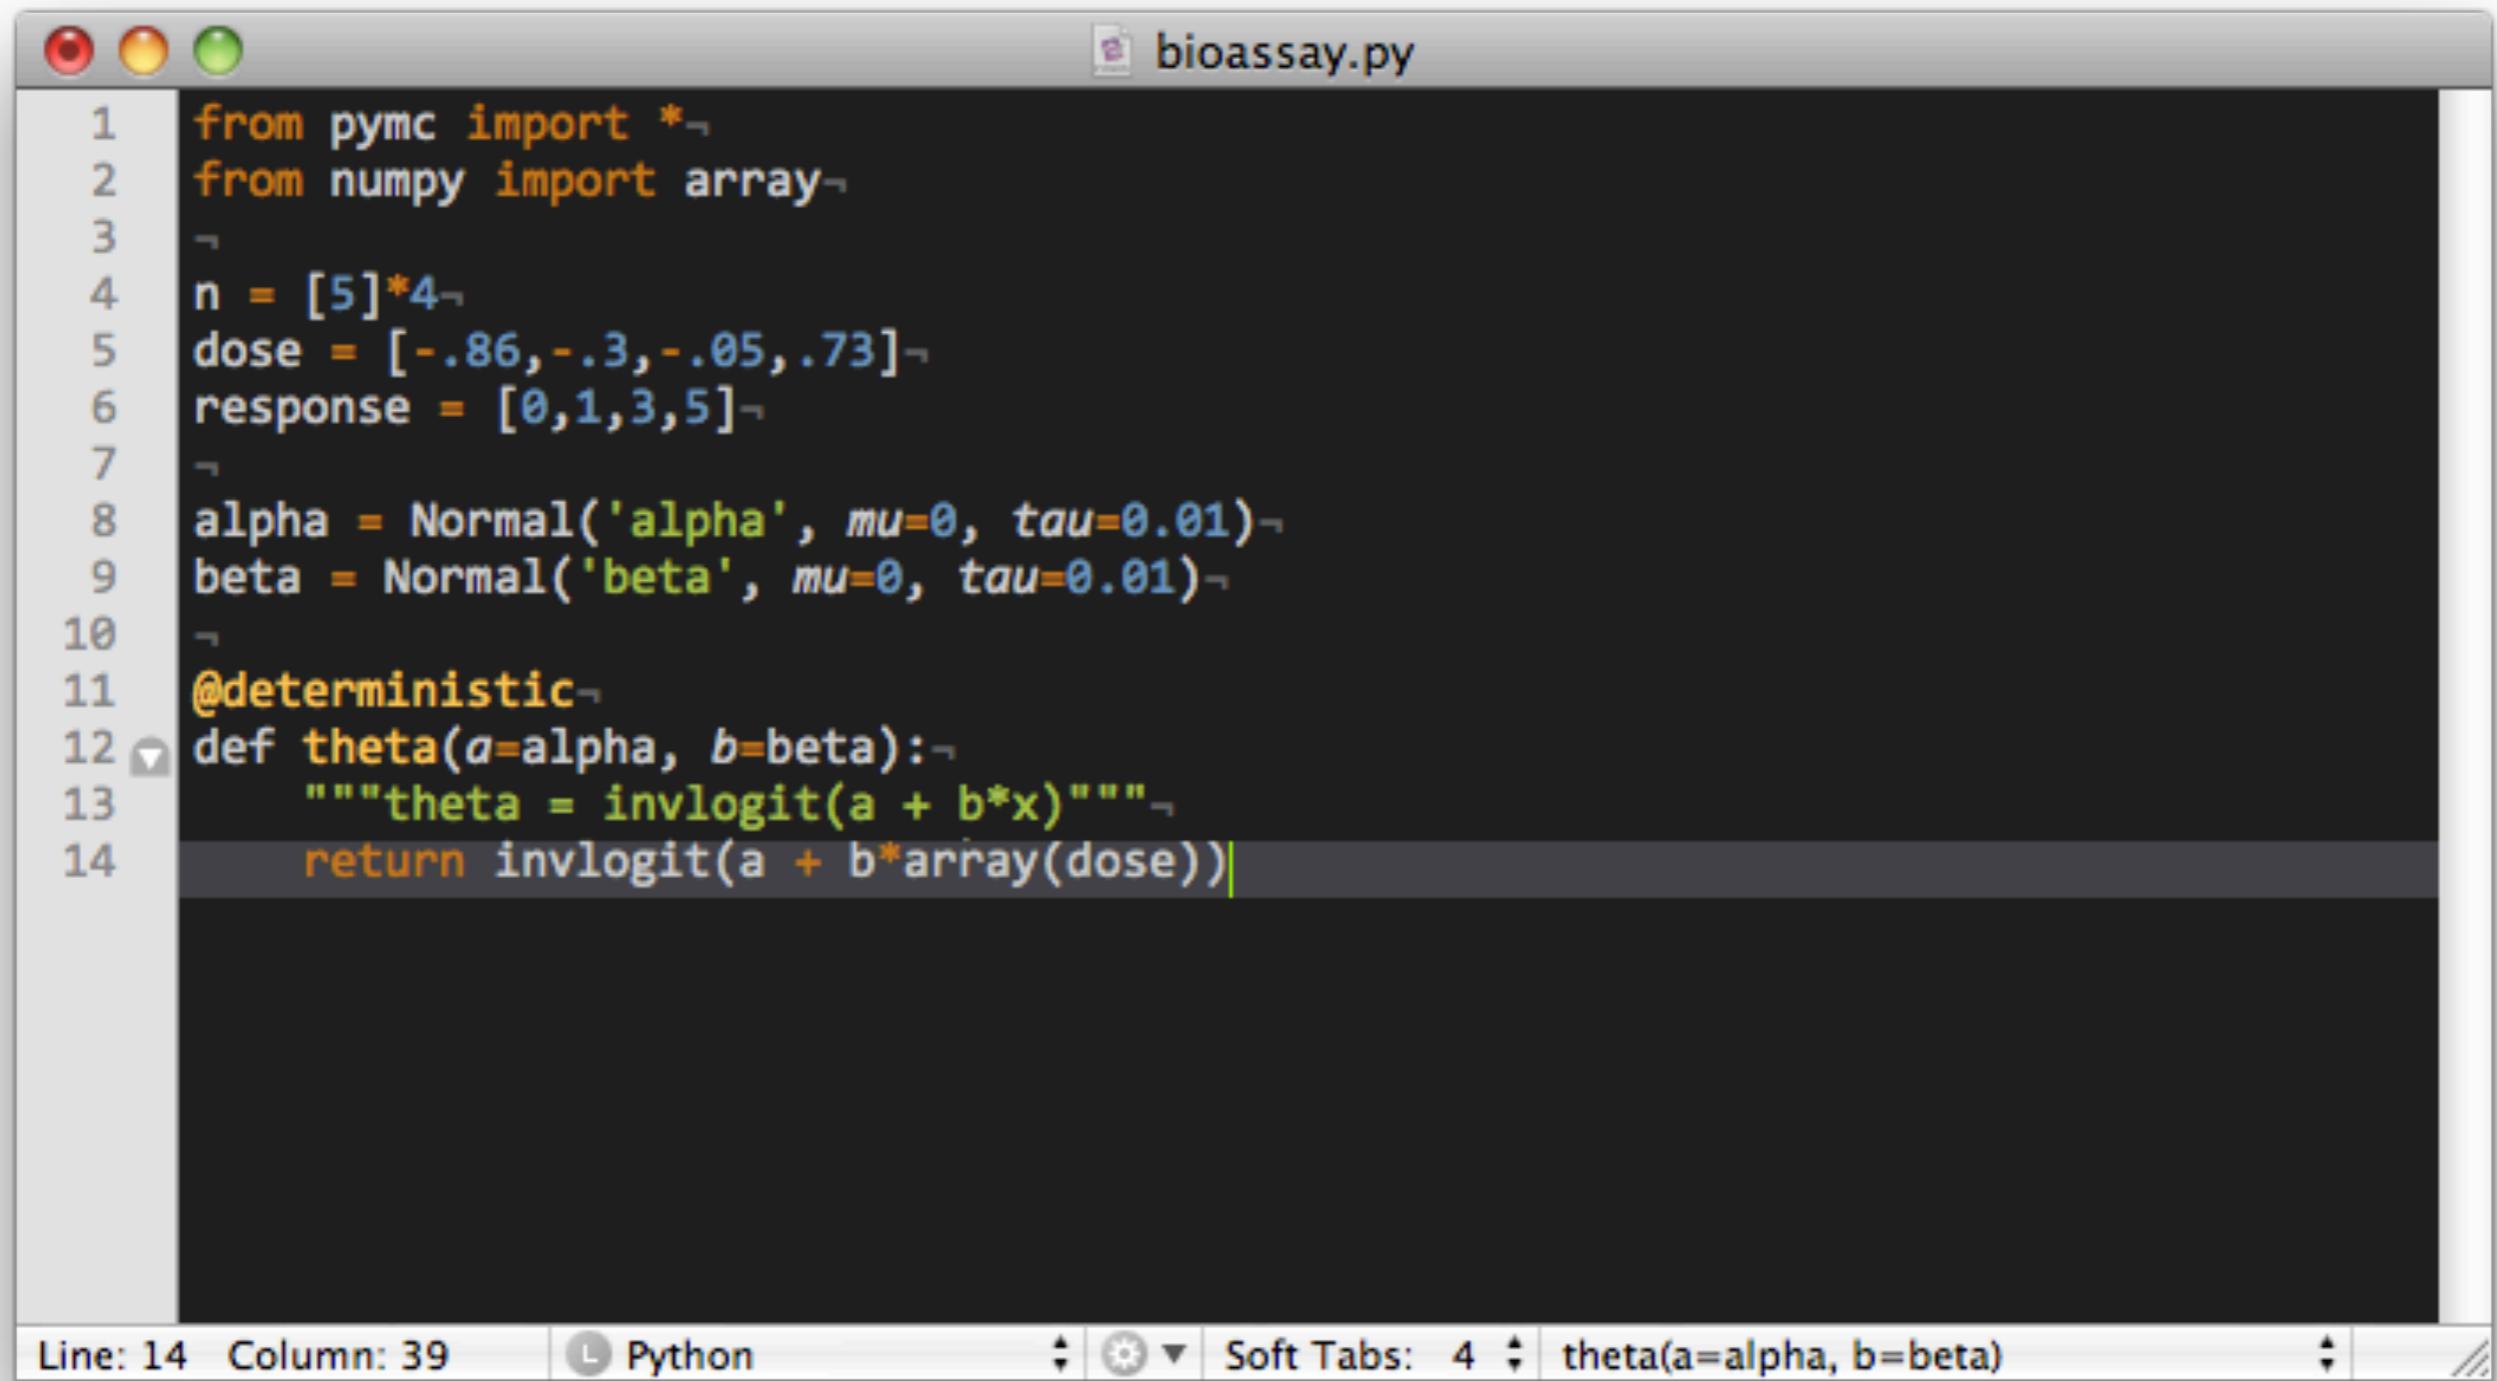
stochastic

A screenshot of a Python code editor window titled "bioassay.py". The code in the editor is as follows:

```
1 from pymc import *
2 from numpy import array
3
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
```

The editor interface includes a toolbar at the top with red, yellow, and green circular icons, and a status bar at the bottom showing "Line: 9 Column: 38" and "Python". The status bar also includes icons for file operations and settings, and a "Soft Tabs: 4" setting.

deterministic



A screenshot of a Python code editor showing a file named "bioassay.py". The code uses the PyMC library to define variables and a function. The function "theta" is annotated with the "@deterministic" decorator. The code editor interface includes a toolbar at the top, a status bar at the bottom with information about the current line and column, and a menu bar with tabs for "Python" and "File".

```
1 from pymc import *
2 from numpy import array
3 
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7 
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
10 
11 @deterministic
12 def theta(a=alpha, b=beta):
13     """theta = invlogit(a + b*x)"""
14     return invlogit(a + b*array(dose))
```

Line: 14 Column: 39 Python Soft Tabs: 4 theta(a=alpha, b=beta)

deterministic

The screenshot shows a window titled "bioassay.py" containing Python code. The code imports pymc and numpy, defines variables n, dose, and response, and creates three normal distributions for alpha, beta, and theta. The theta variable is defined as a Lambda function that takes a value and applies a logistic regression model with parameters a and b.

```
1 from pymc import *
2 from numpy import array
3
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
10
11 theta = Lambda('theta', lambda a=alpha, b=beta: invlogit(a + b*array(dose)))
```

Line: 11 Column: 77 Python Soft Tabs: 4

data stochastic

```
bioassay.py
1 from pymc import *
2 from numpy import array
3
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
10
11 theta = Lambda('theta', lambda a=alpha, b=beta: invlogit(a + b*array(dose)))
12
13 @observed_
14 def deaths(value=response, n=n, p=theta):
15     """deaths ~ binomial_like(n, p)"""
16     return binomial_like(value, n, p)

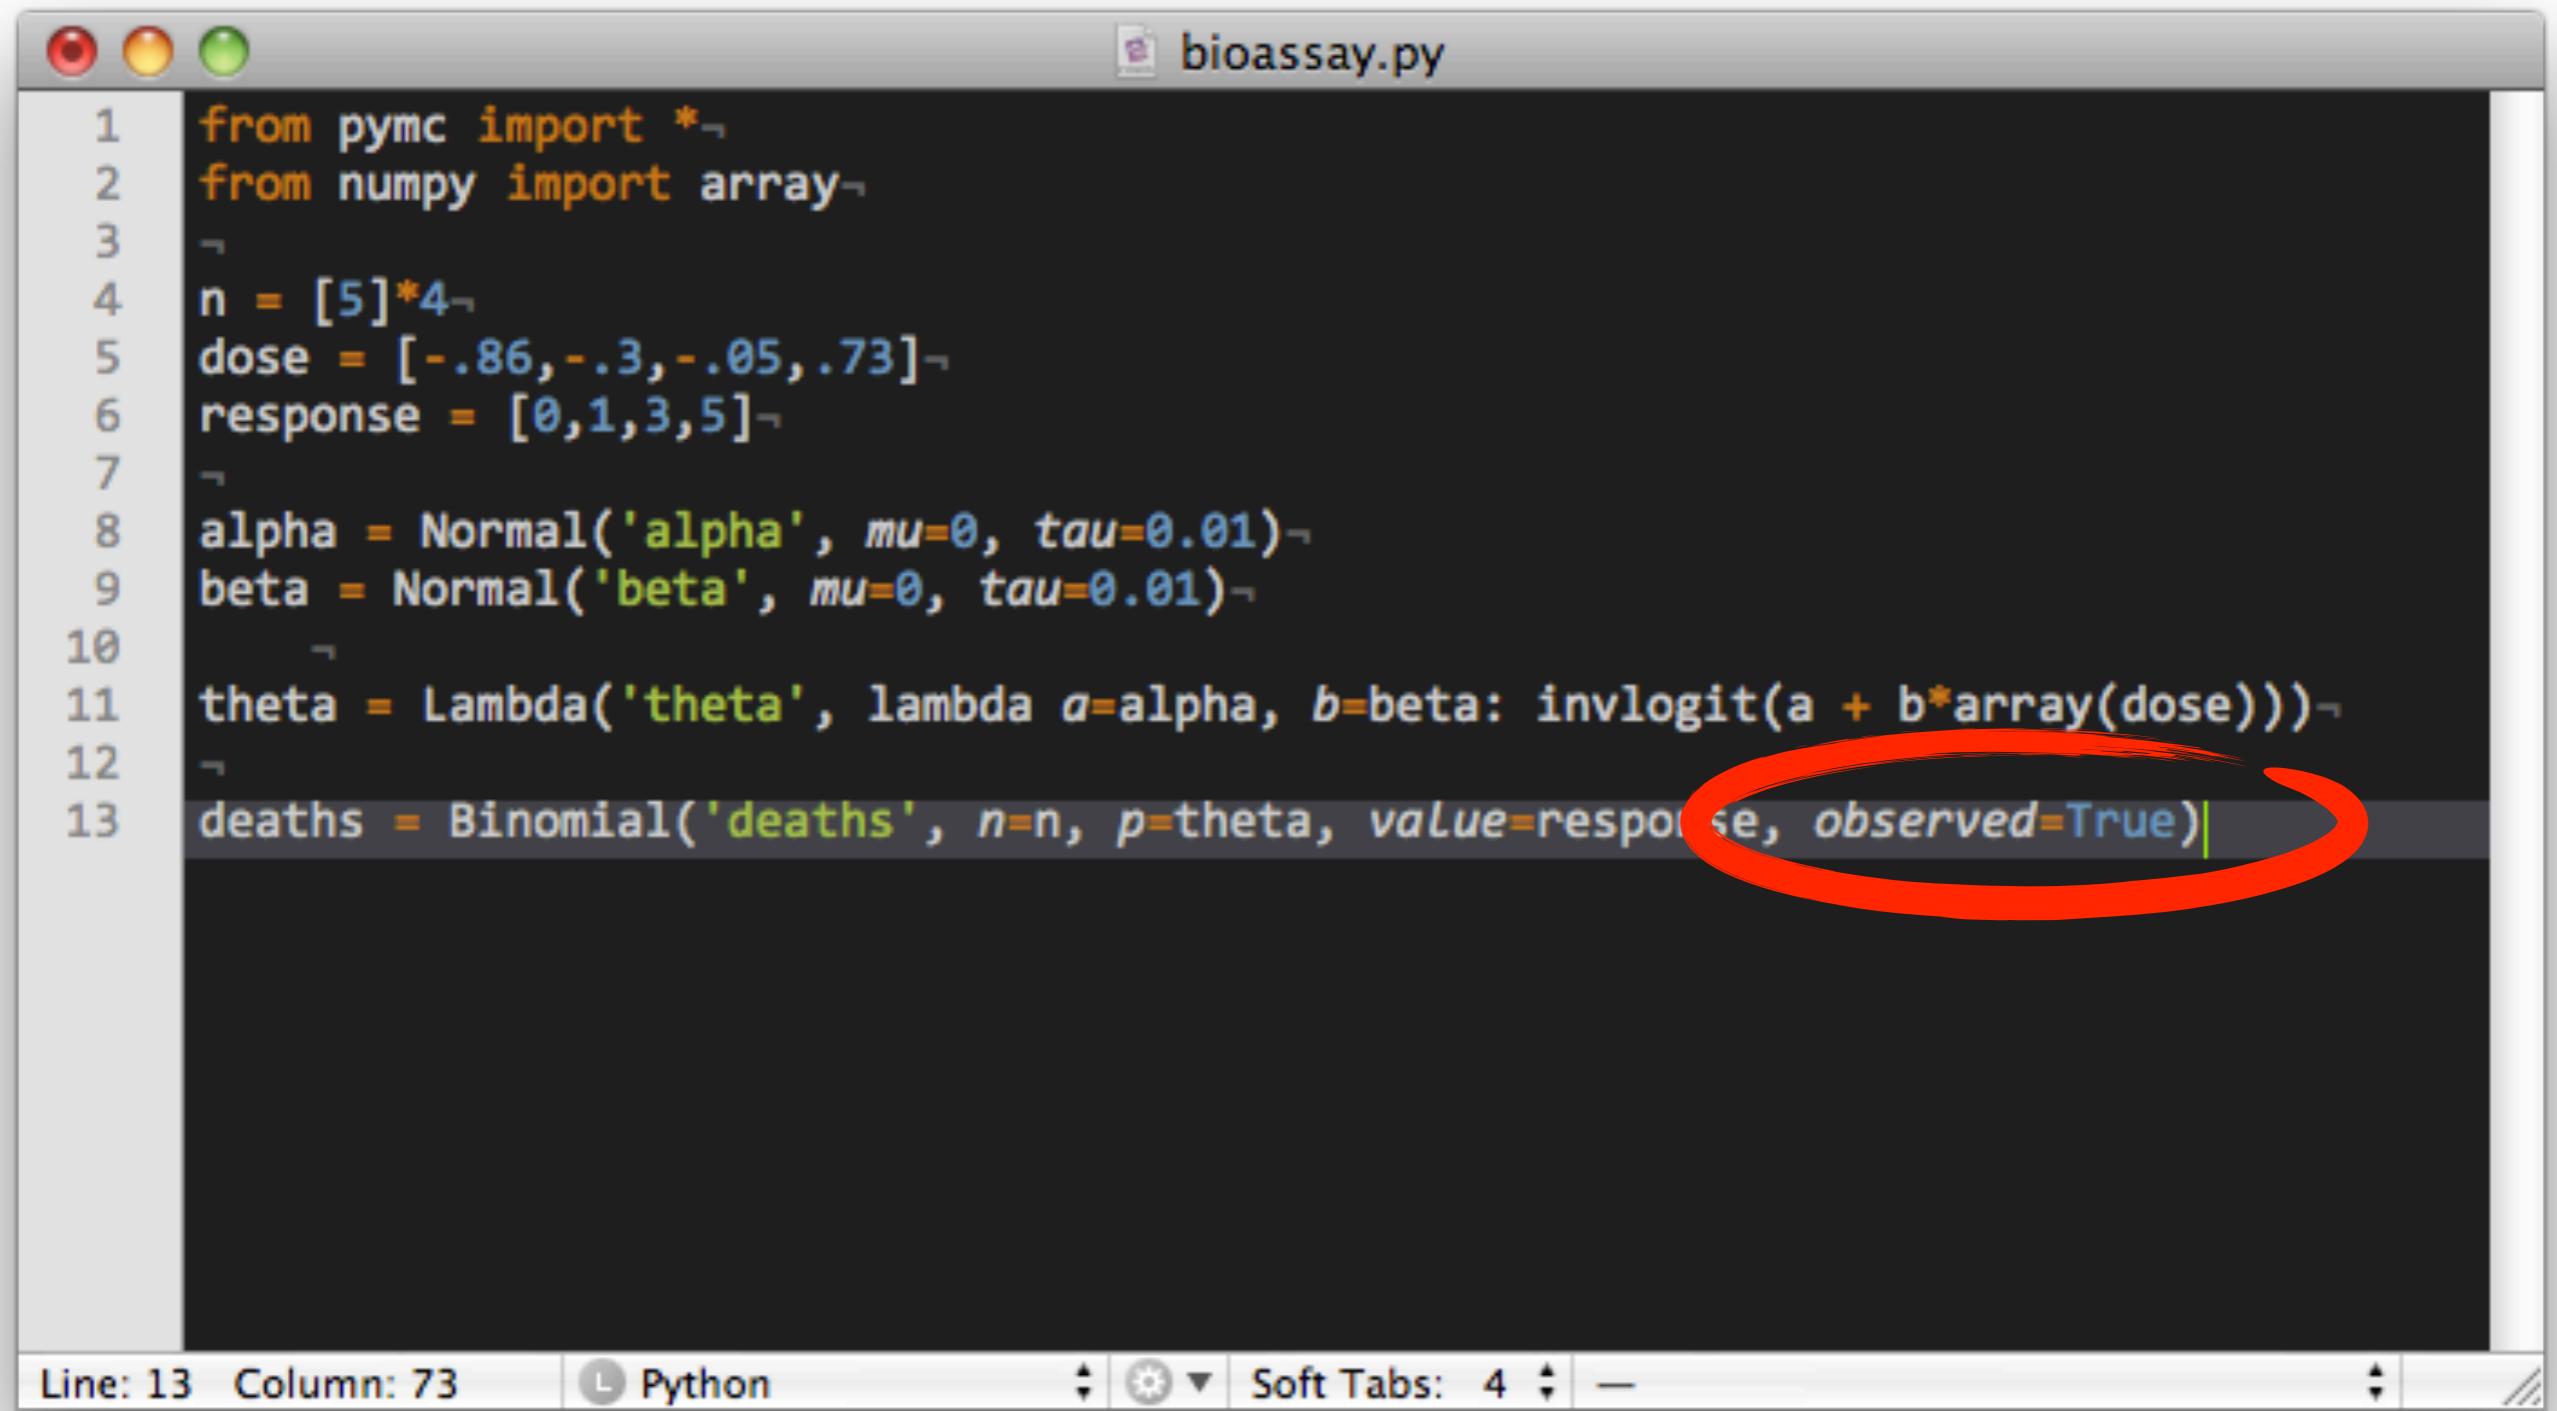
Line: 16  Column: 43 | Python | Soft Tabs: 4 | deaths(value=response, n=n, p=theta)
```

data stochastic

```
1 from pymc import *
2 from numpy import array
3
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
10
11 theta = Lambda('theta', lambda a=alpha, b=beta: invlogit(a + b*array(dose)))
12
13 deaths = Binomial('deaths', n=n, p=theta, value=response, observed=True)
```

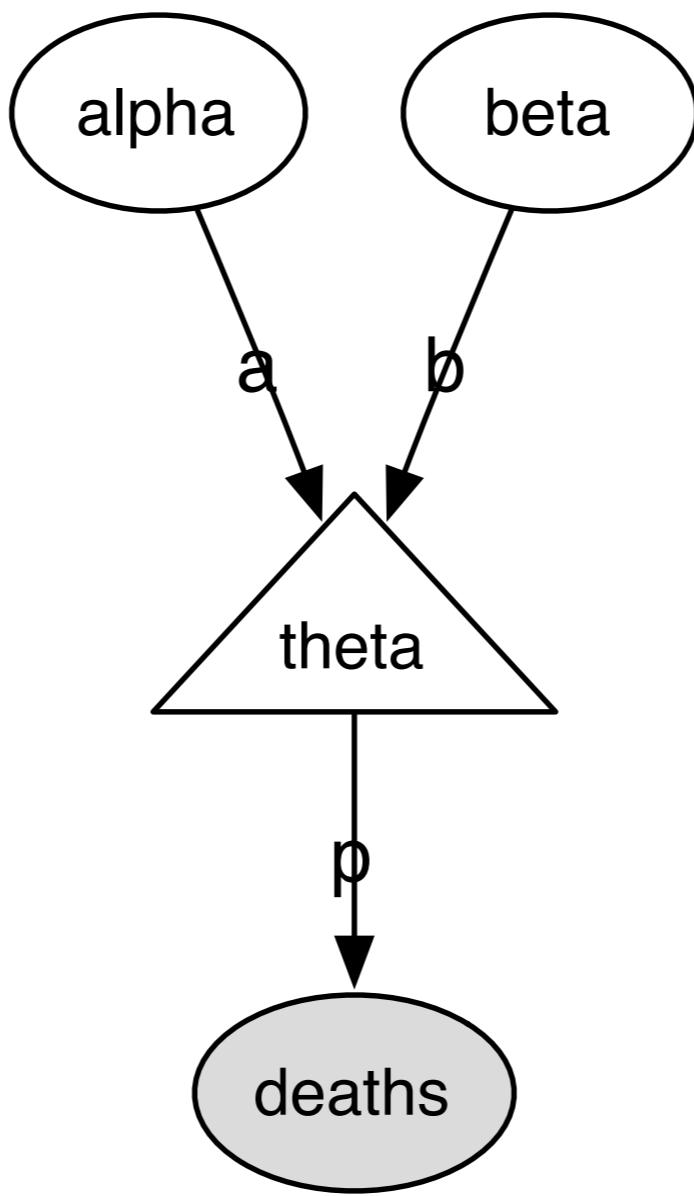
Line: 13 Column: 73 L Python Soft Tabs: 4

data stochastic



```
1 from pymc import *
2 from numpy import array
3
4 n = [5]*4
5 dose = [-.86, -.3, -.05, .73]
6 response = [0,1,3,5]
7
8 alpha = Normal('alpha', mu=0, tau=0.01)
9 beta = Normal('beta', mu=0, tau=0.01)
10
11 theta = Lambda('theta', lambda a=alpha, b=beta: invlogit(a + b*array(dose)))
12
13 deaths = Binomial('deaths', n=n, p=theta, value=response, observed=True)
```

Line: 13 Column: 73 L Python Soft Tabs: 4

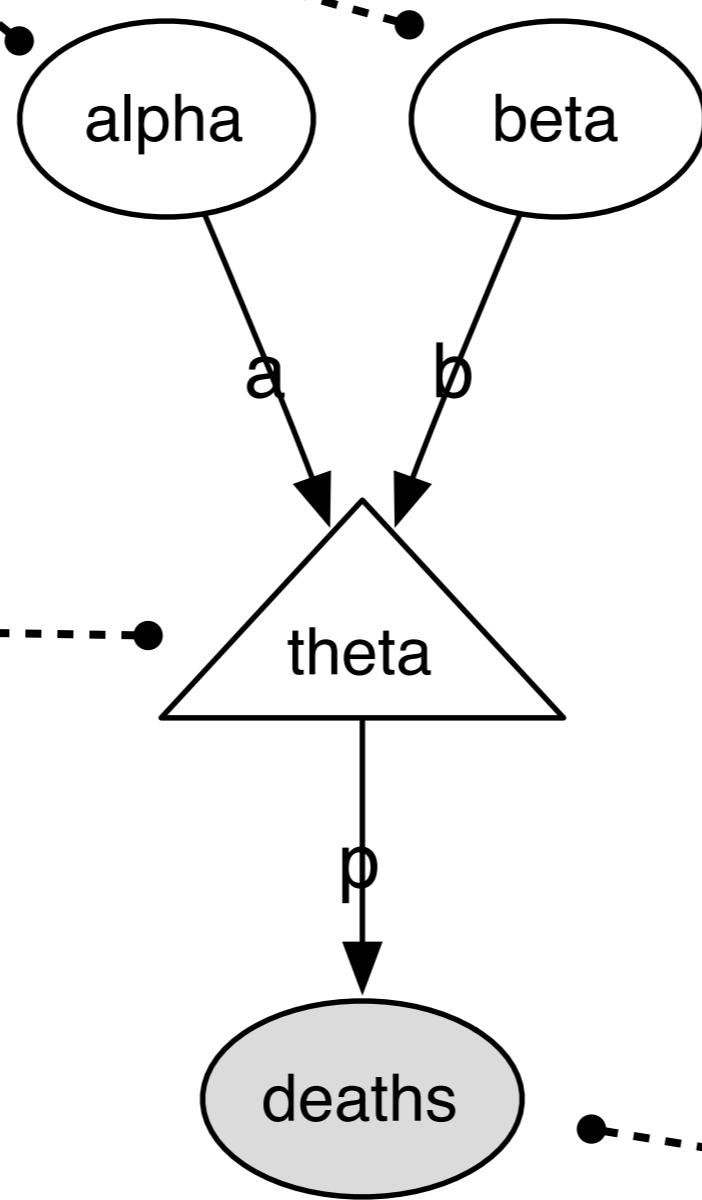


>>> graph.dag(M)

stochastic

deterministic

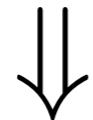
data
stochastic



>>> graph.dag(M)

factor potential

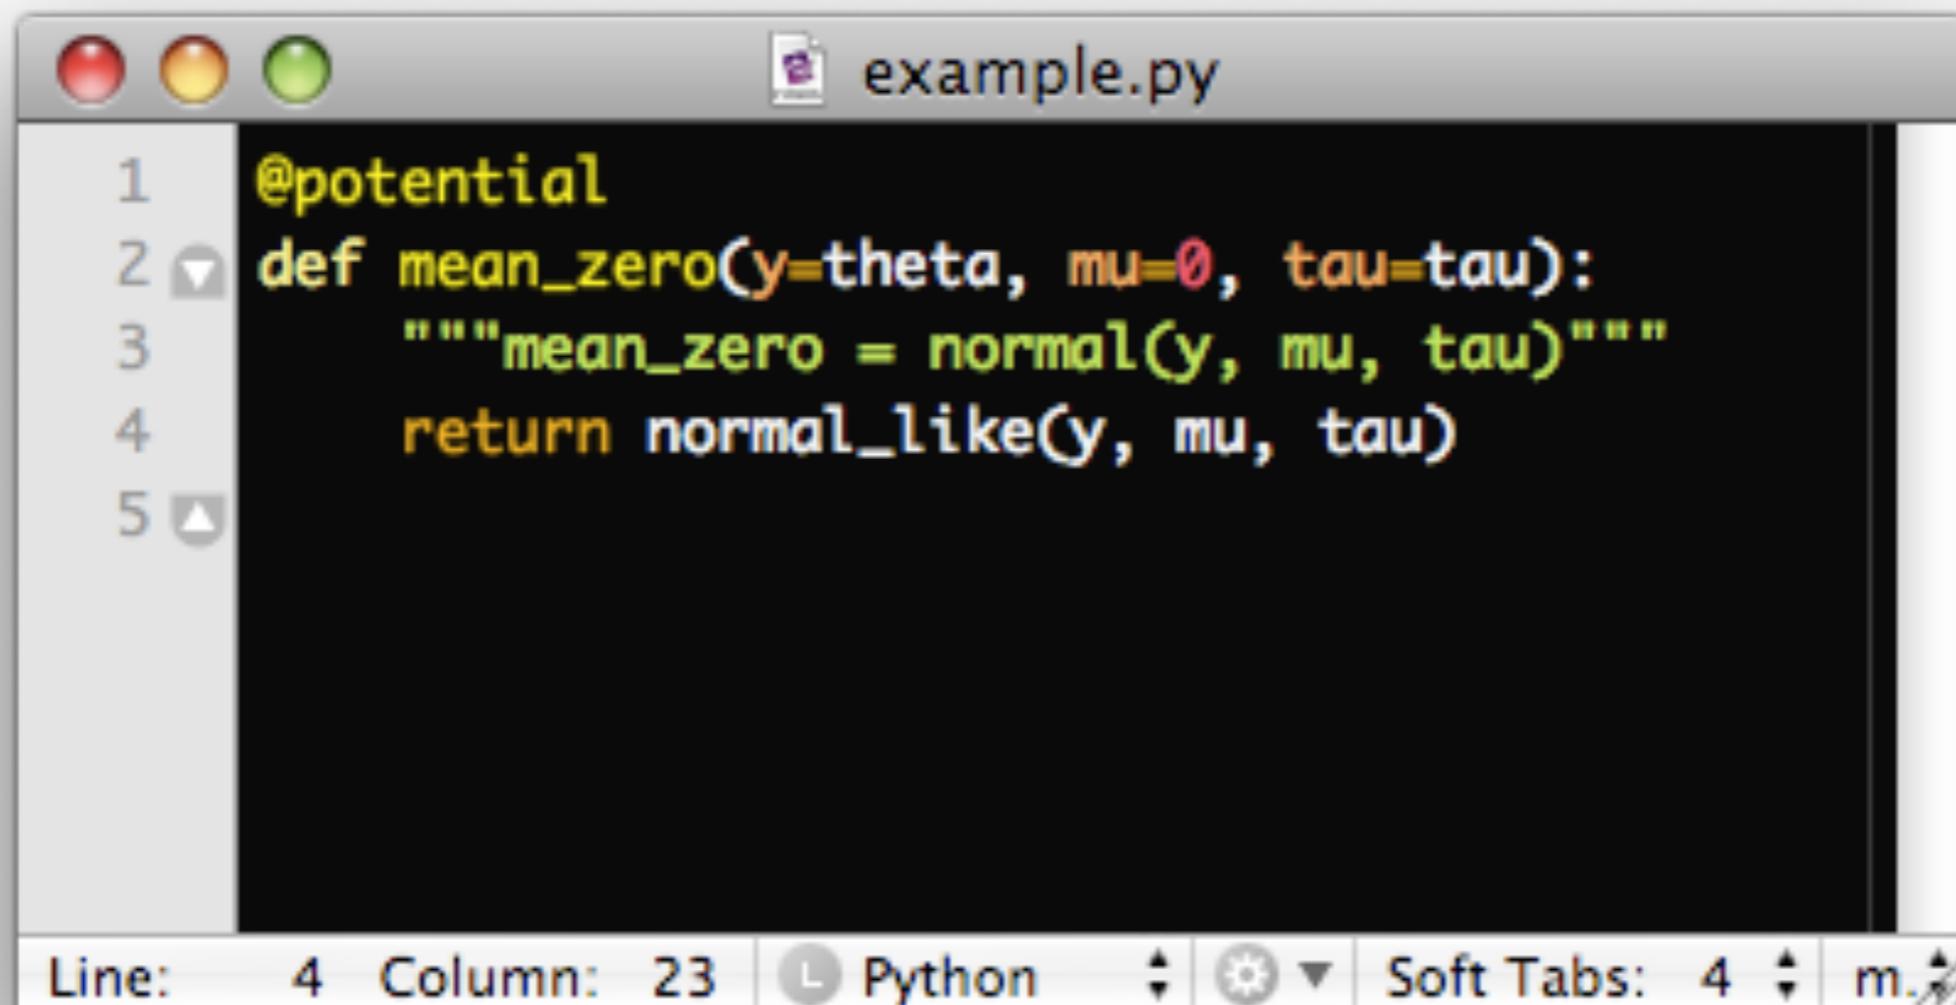
$$p(\theta, \phi | y) \propto p(y|\theta)p(\theta|\phi)p(\phi)$$



$$p(\theta, \phi, \tau | y) \propto p(y|\theta)p(\theta|\phi)p(\theta|\tau)p(\tau)p(\phi)$$

arbitrary log-probability terms

factor potential



The screenshot shows a Python code editor window titled "example.py". The code defines a function "mean_zero" with a docstring and a return statement. The editor interface includes a toolbar with red, yellow, and green buttons, a status bar at the bottom showing "Line: 4 Column: 23 Python", and a menu icon in the top right corner.

```
1 @potential
2 def mean_zero(y=theta, mu=0, tau=tau):
3     """mean_zero = normal(y, mu, tau)"""
4     return normal_like(y, mu, tau)
5
```

missing data

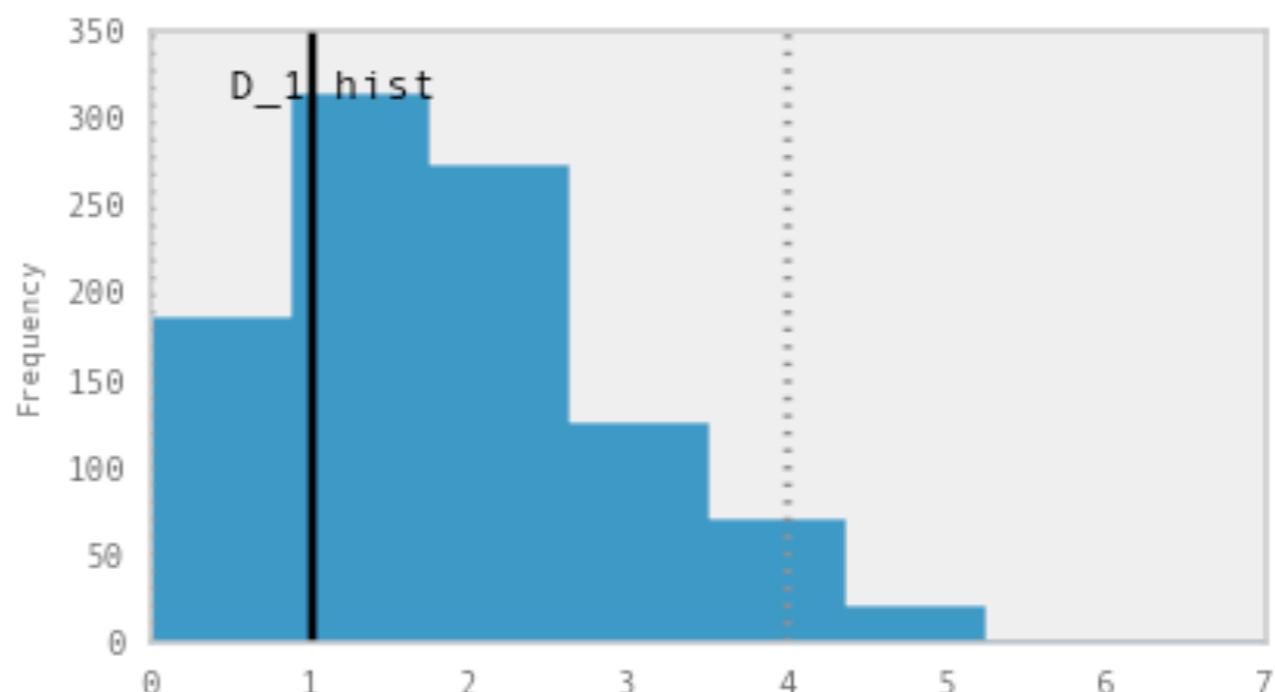
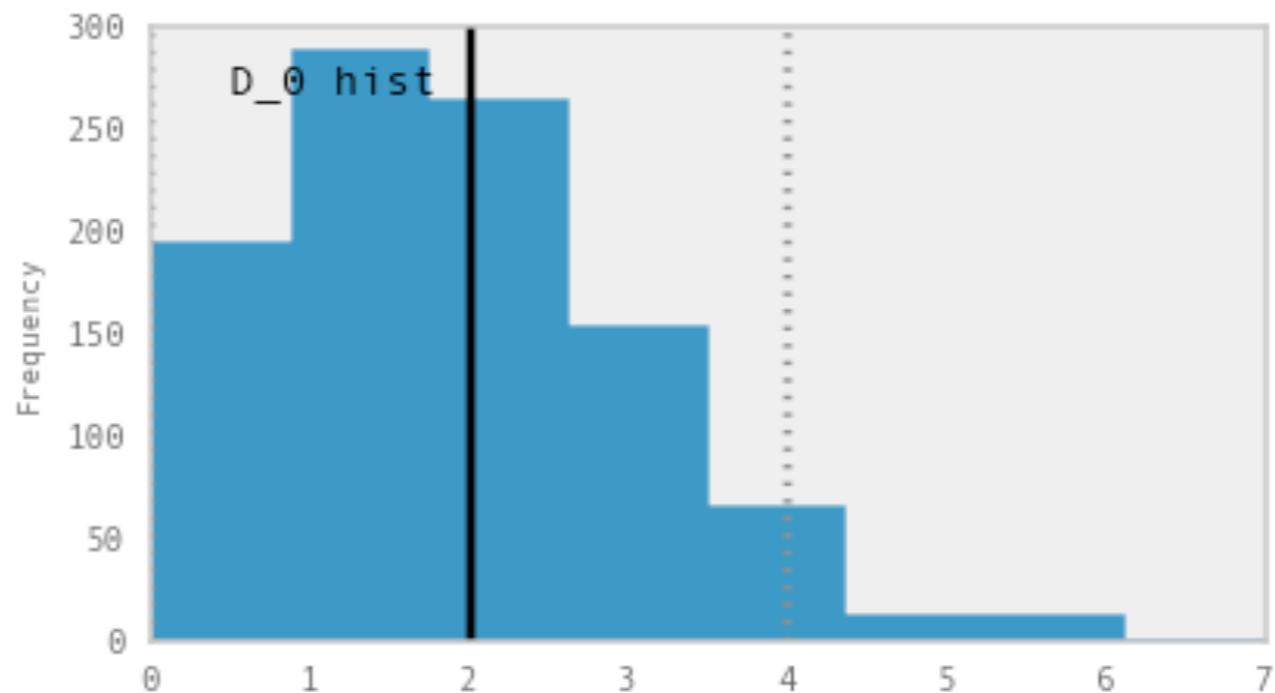
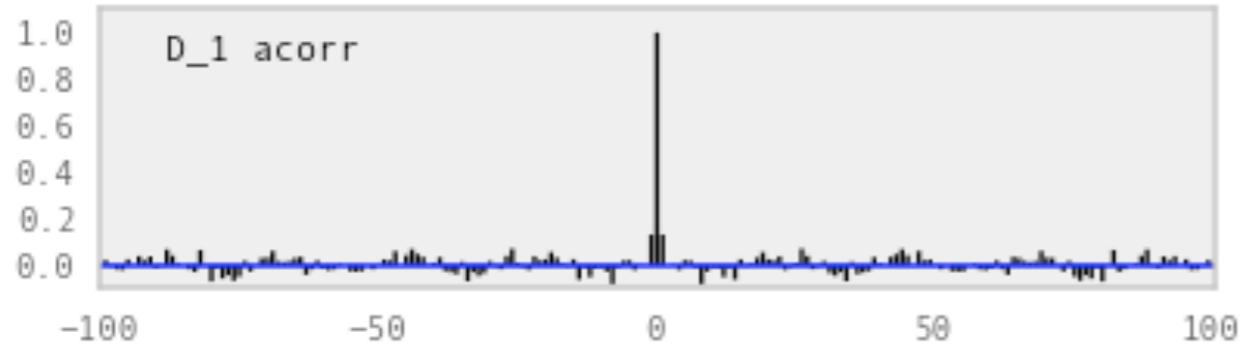
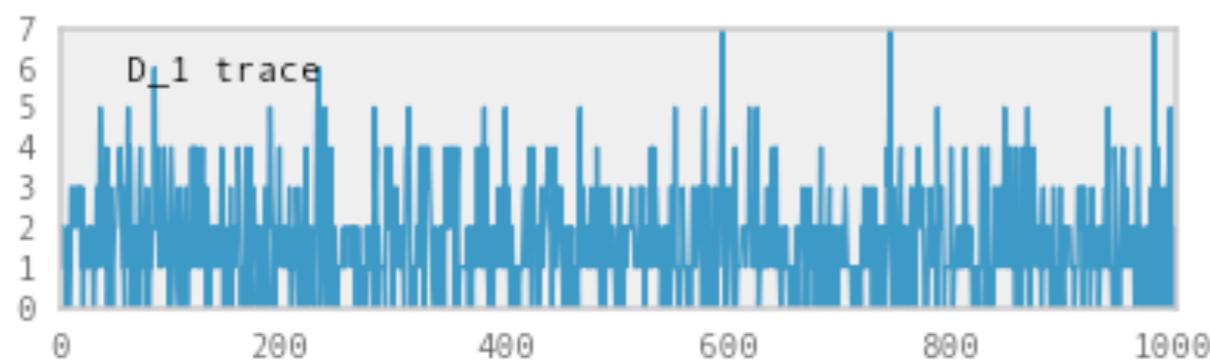
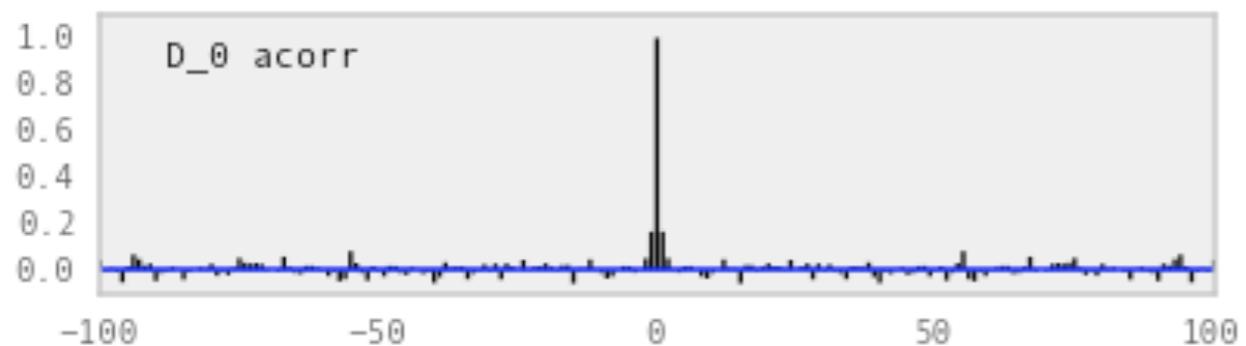
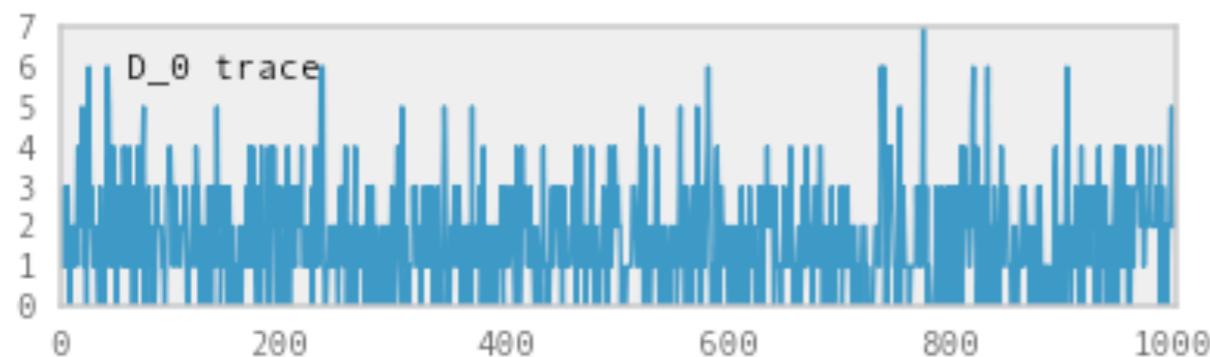
The screenshot shows a window titled "missing_count.py" containing Python code. The code uses masked arrays to handle missing data. It includes imports for numpy, masked_array, and pymc, defines a count array with missing values (-999), creates a masked array from it, and passes it to a Poisson stochastic variable.

```
1 import numpy as np
2 from numpy.ma import masked_array
3 from pymc import *
4
5 # Missing values indicated by negative-valued placeholders-
6 count = np.array([ 4, 5, 4, 0, 1, 4, 3, 4, 0, 6, 3, 3, 4, 0, 2, 6,
7 3, 3, 5, 4, 5, 3, 1, 4, 4, 1, 5, 5, 3, 4, 2, 5,
8 2, 2, 3, 4, 2, 1, 3, -999, 2, 1, 1, 1, 1, 3, 0, 0,
9 1, 0, 1, 1, 0, 0, 3, 1, 0, 3, 2, 2, 0, 1, 1, 1,
10 0, 1, 0, 1, 0, 0, 2, 1, 0, 0, 0, 1, 1, 0, 2,
11 3, 3, 1, -999, 2, 1, 1, 1, 2, 4, 2, 0, 0, 1, 4,
12 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1])-
13
14 # Mean count-
15 m = Exponential('mu', beta=1)
16
17 # Create masked array with data-
18 masked_values = masked_array(count, mask=count== -999)
19
20 # Pass masked array to data stochastic, and it does the right thing-
21 D = Poisson('D', mu=m, value=masked_values, observed=True)
```

Line: 18 Column: 1

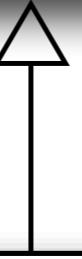
Python

Soft Tabs: 4



Model

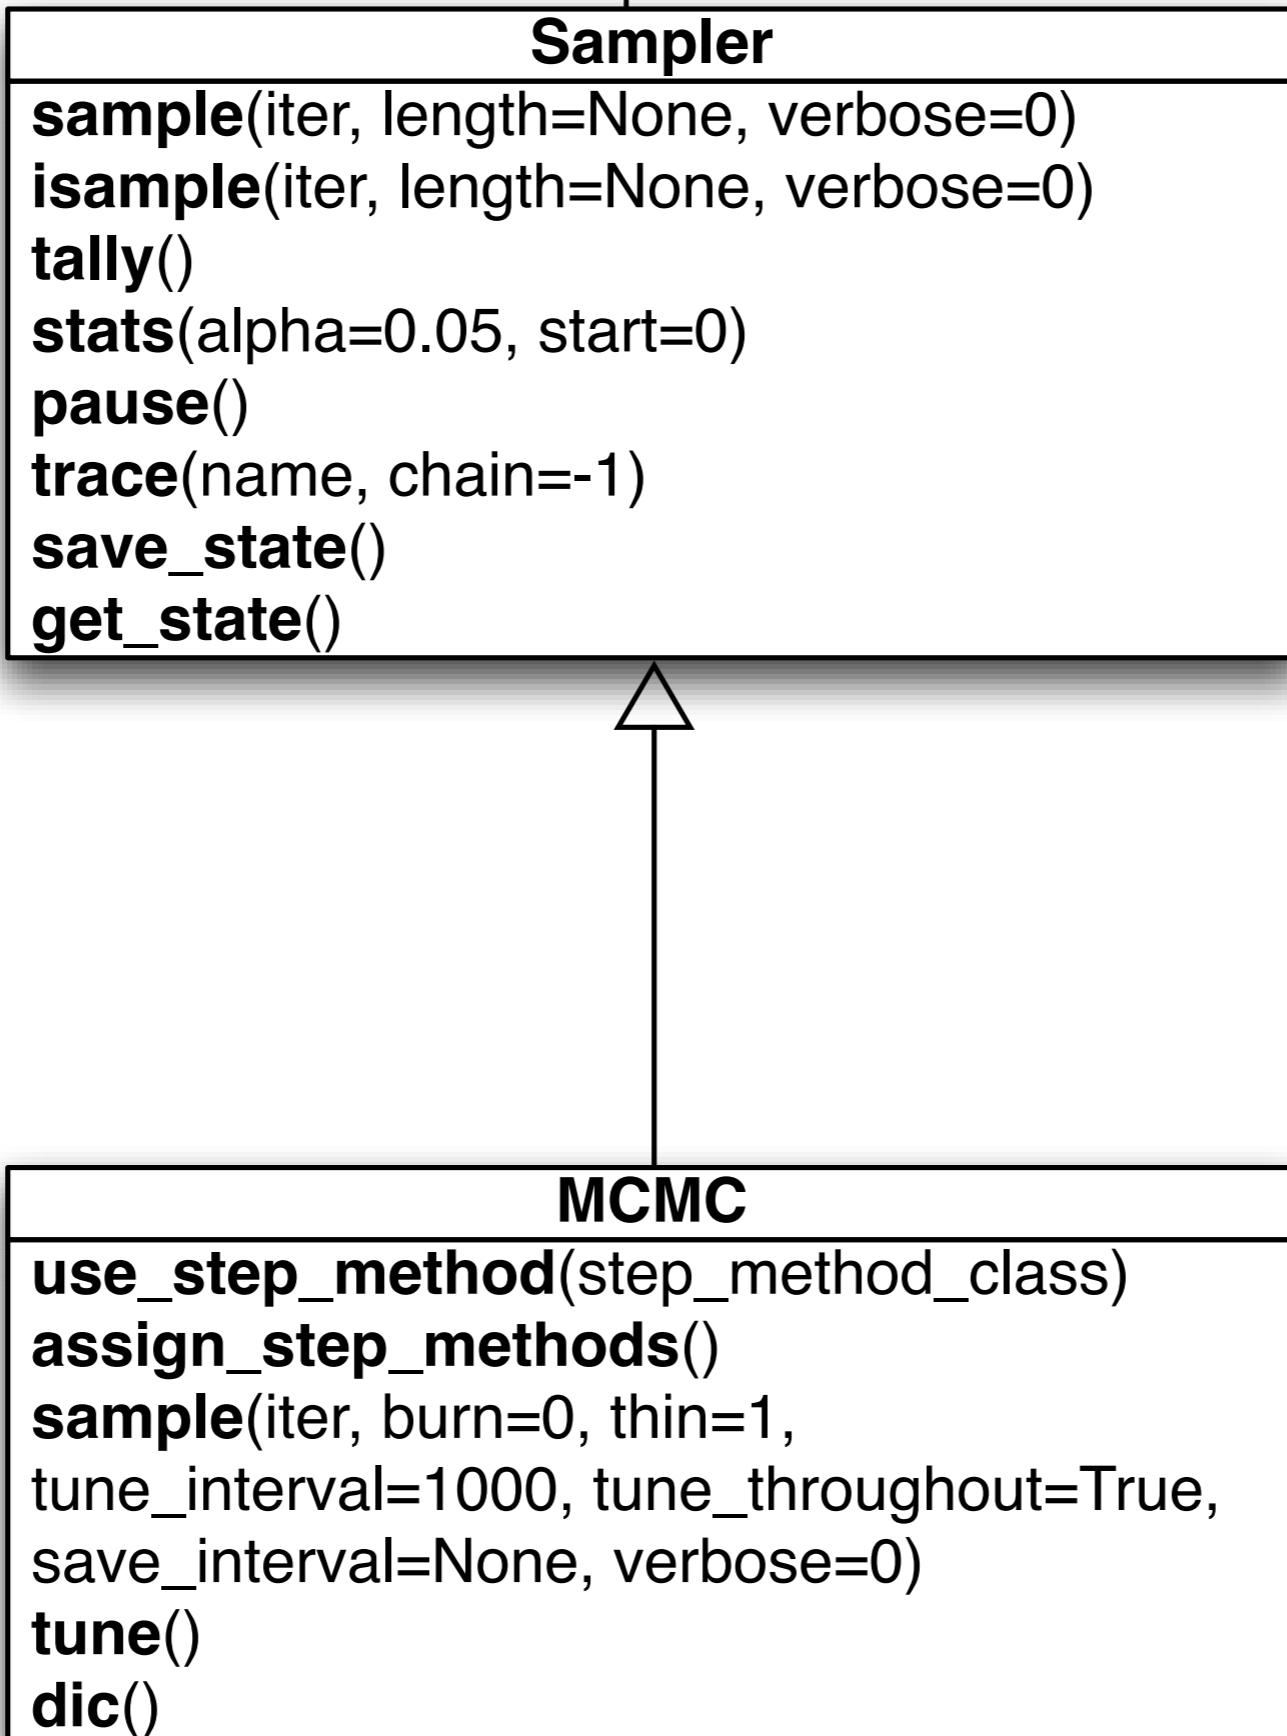
deterministics
stochastics
data
variables
potentials
containers
nodes
all_objects
moral_neighbors
extended_children
generations
seed()
get_node(node_name)
draw_from_prior()



Sampler

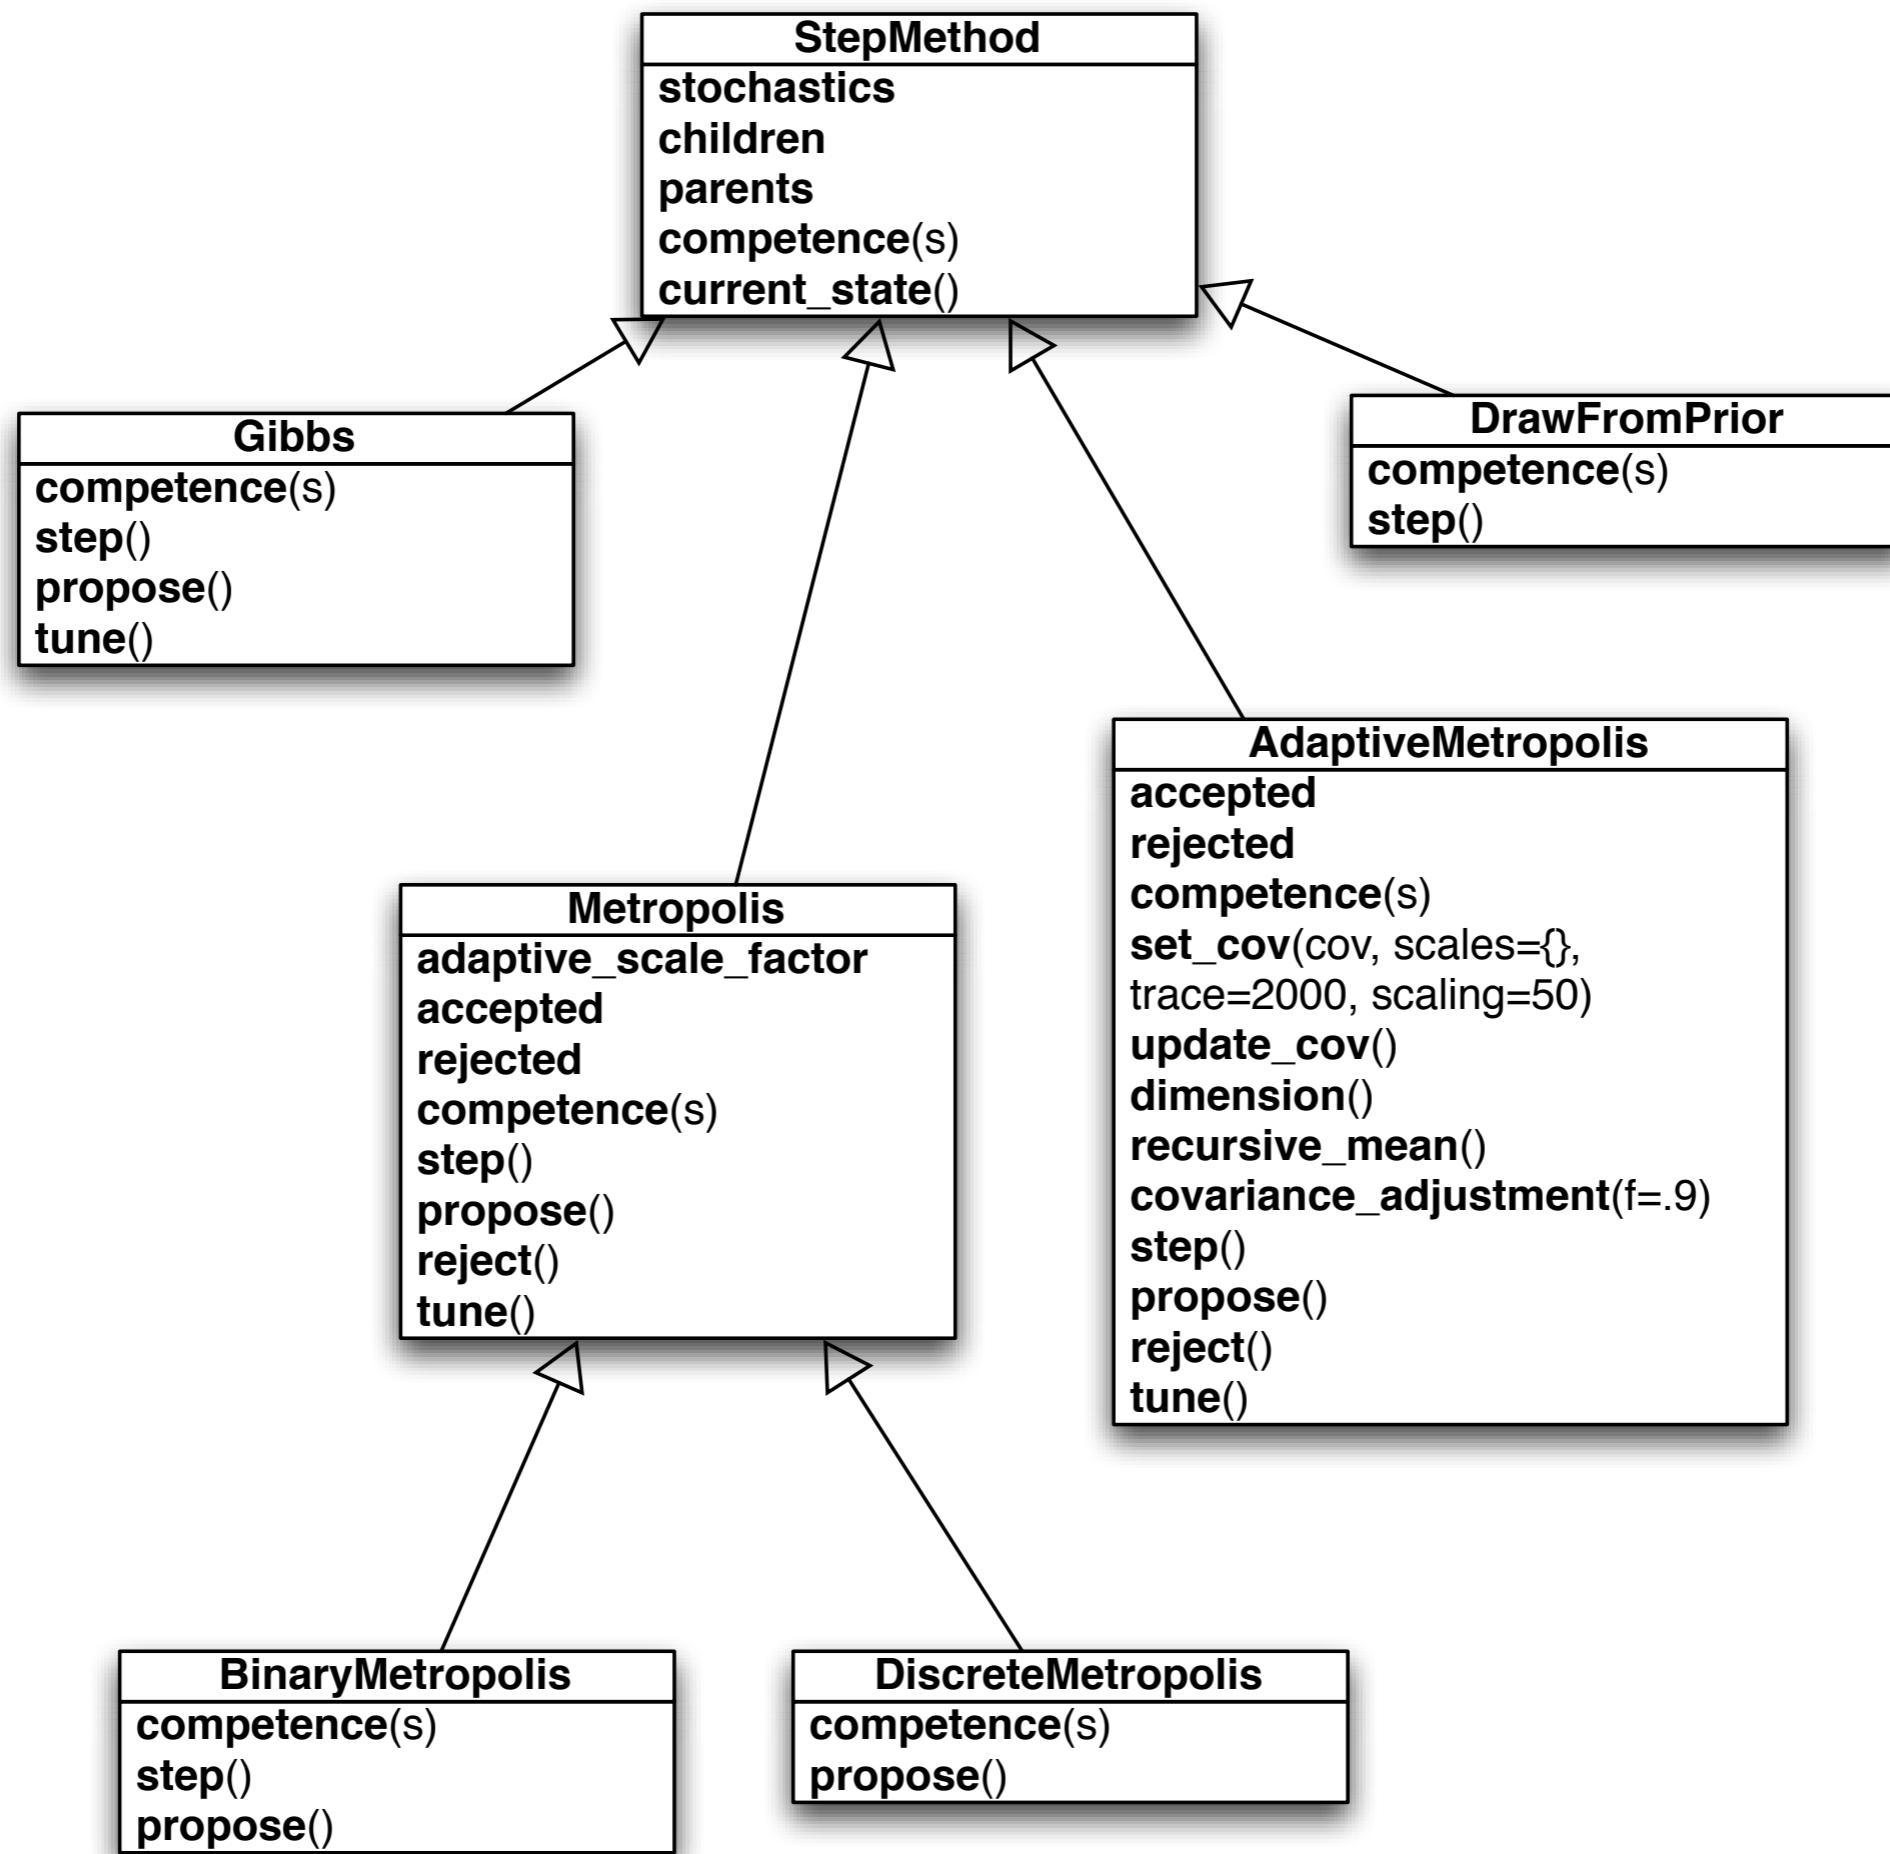
sample(iter, length=None, verbose=0)
isample(iter, length=None, verbose=0)
tally()
stats(alpha=0.05, start=0)
pause()
trace(name, chain=-1)
save_state()
get_state()







Step Methods



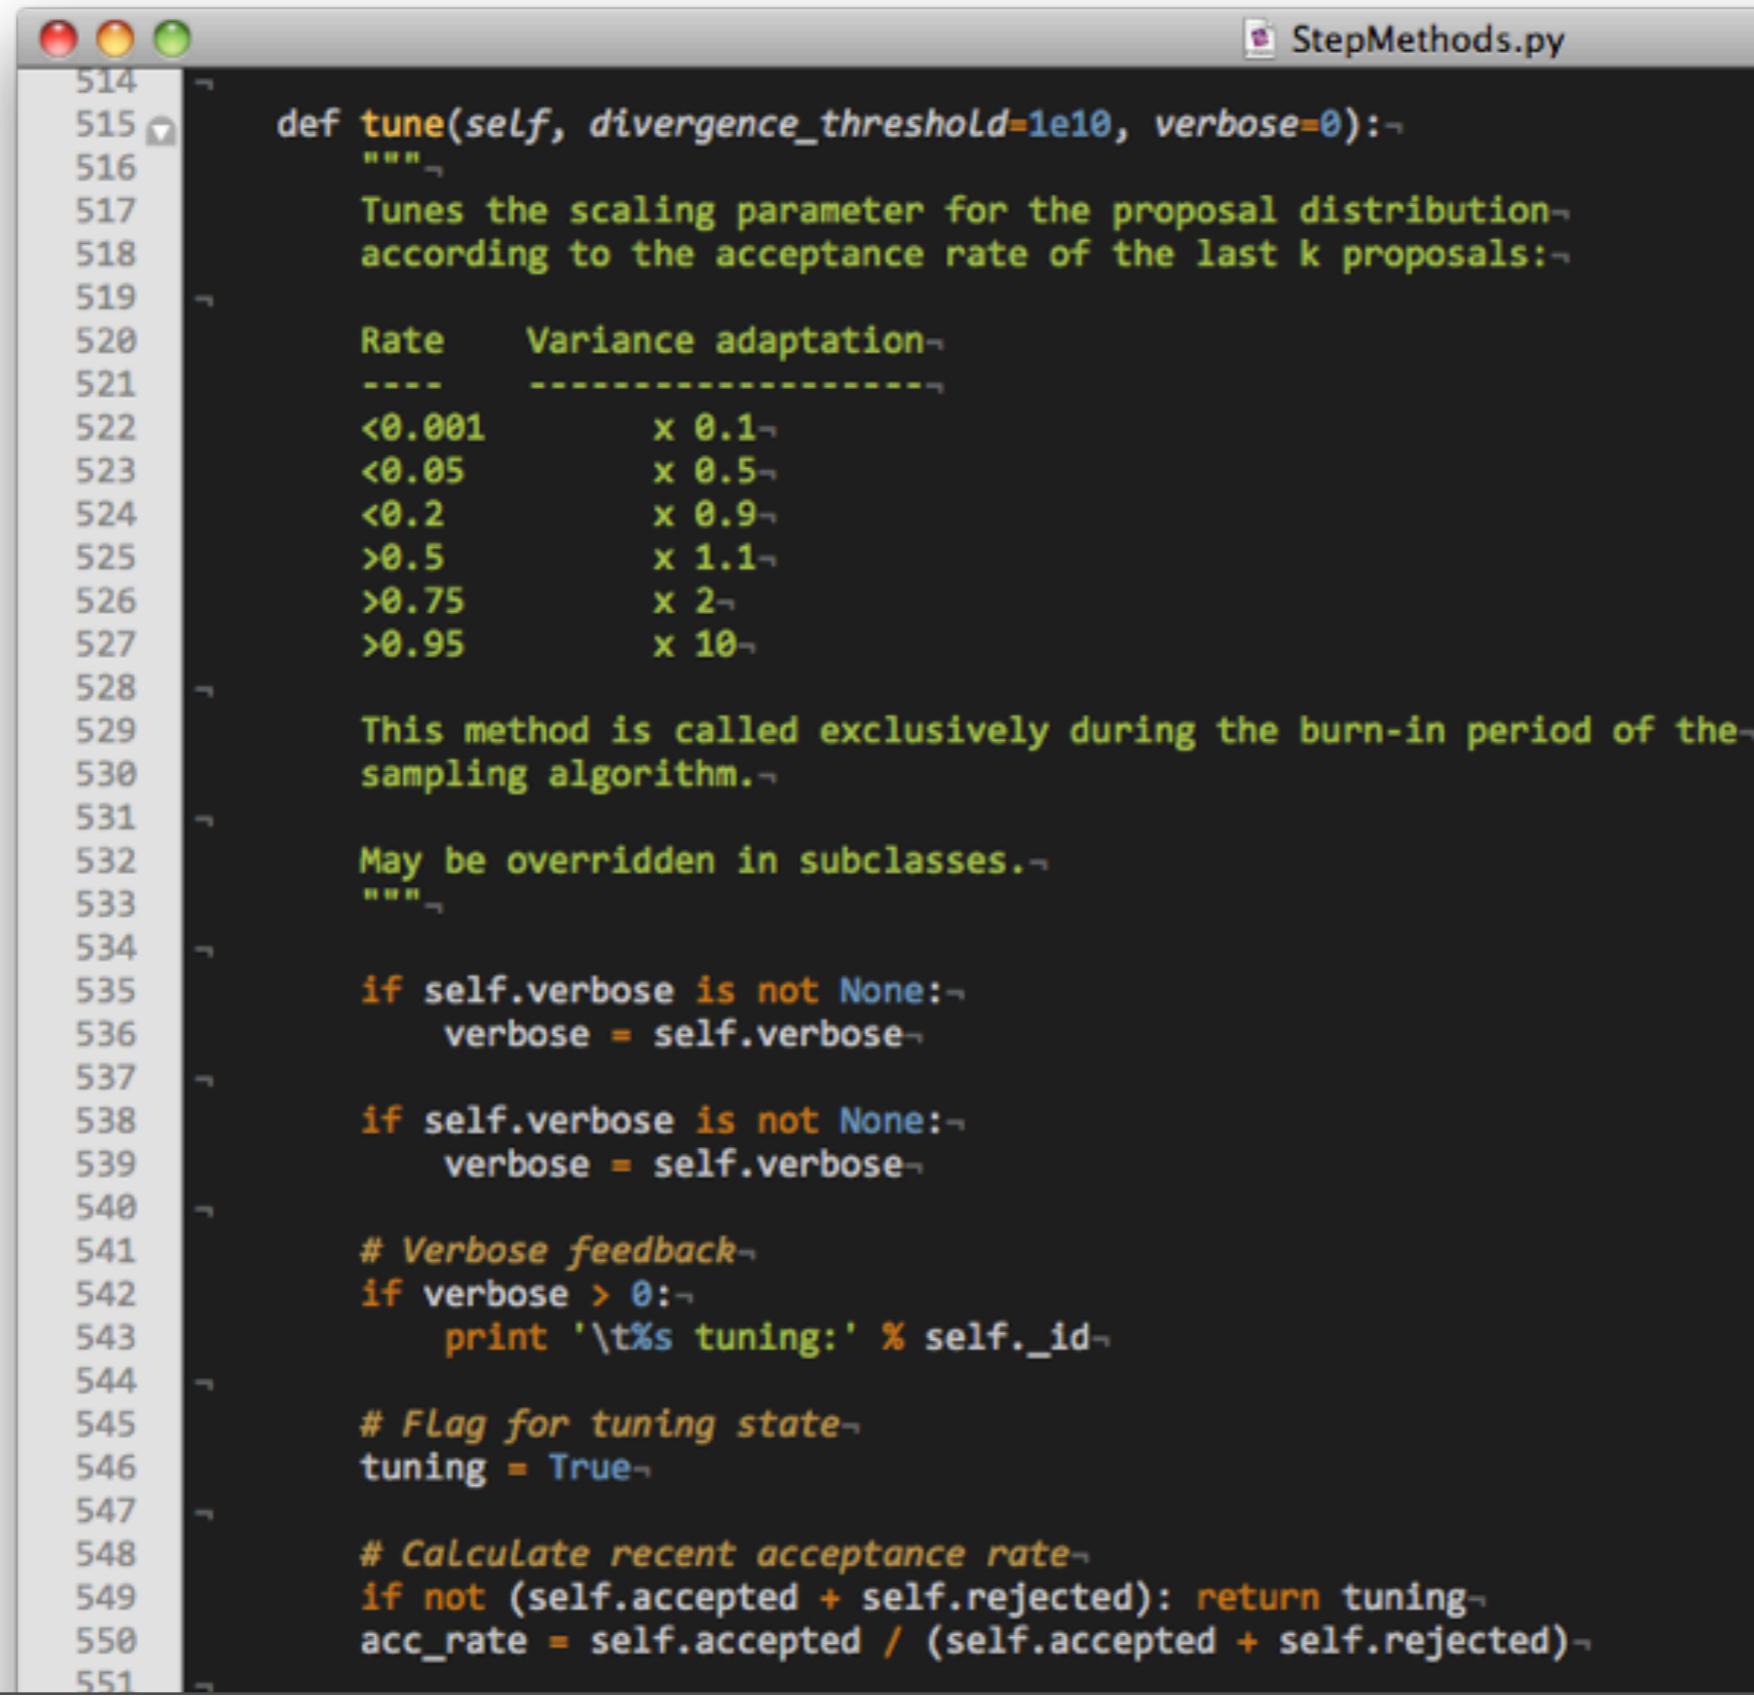
StepMethod methods

The screenshot shows a Python code editor with the file `StepMethods.py` open. The code defines a class with a method `step` that handles the proposal and acceptance of a candidate value. It includes comments explaining the logic and handles exceptions like `ZeroProbability`.

```
421     return s.~
422
423     def step(self):~
424         """~  
425             The default step method applies if the variable is floating-point-~  
426             valued, and is not being proposed from its prior.~  
427         """~
428
429         # Probability and Likelihood for s's current value:~
430
431         if self.verbose>2:~
432             print~
433             print self._id + ' getting initial logp.'~
434
435         if self.proposal_distribution == "Prior":~
436             logp = self.loglike~
437         else:~
438             logp = self.logp_plus_loglike~
439
440         if self.verbose>2:~
441             print self._id + ' proposing.'~
442
443         # Sample a candidate value~
444         self.propose()|~
445
446         # Probability and Likelihood for s's proposed value:~
447         try:~
448             if self.proposal_distribution == "Prior":~
449                 logp_p = self.loglike~
450                 # Check for weirdness before accepting jump~
451                 if self.check_before_accepting:~
452                     self.stochastic.logp~
453                 else:~
454                     logp_p = self.logp_plus_loglike~
455
456             except ZeroProbability:~
457
458                 # Reject proposal~
```

```
422
423     def step(self):-
424         """
425             The default step method applies if the variable is floating-point-
426             valued, and is not being proposed from its prior.
427         """
428
429         # Probability and Likelihood for s's current value:-
430
431         if self.verbose>2:-
432             print-
433                 print self._id + ' getting initial logp.'-
434
435         if self.proposal_distribution == "Prior":-
436             logp = self.loglike-
437         else:-_
438             logp = self.logp_plus_loglike-
439
440         if self.verbose>2:-
441             print self._id + ' proposing.'-
442
443         # Sample a candidate value-
444         self.propose()-
445
446         # Probability and Likelihood for s's proposed value:-
447         try:-_
448             if self.proposal_distribution == "Prior":-
449                 logp_p = self.loglike-
450                 # Check for weirdness before accepting jump-
451                 if self.check_before_accepting:-_
452                     self.stochastic.logp-
453                 else:-_
454                     logp_p = self.logp_plus_loglike-
455
456         except ZeroProbability-
```

StepMethod methods

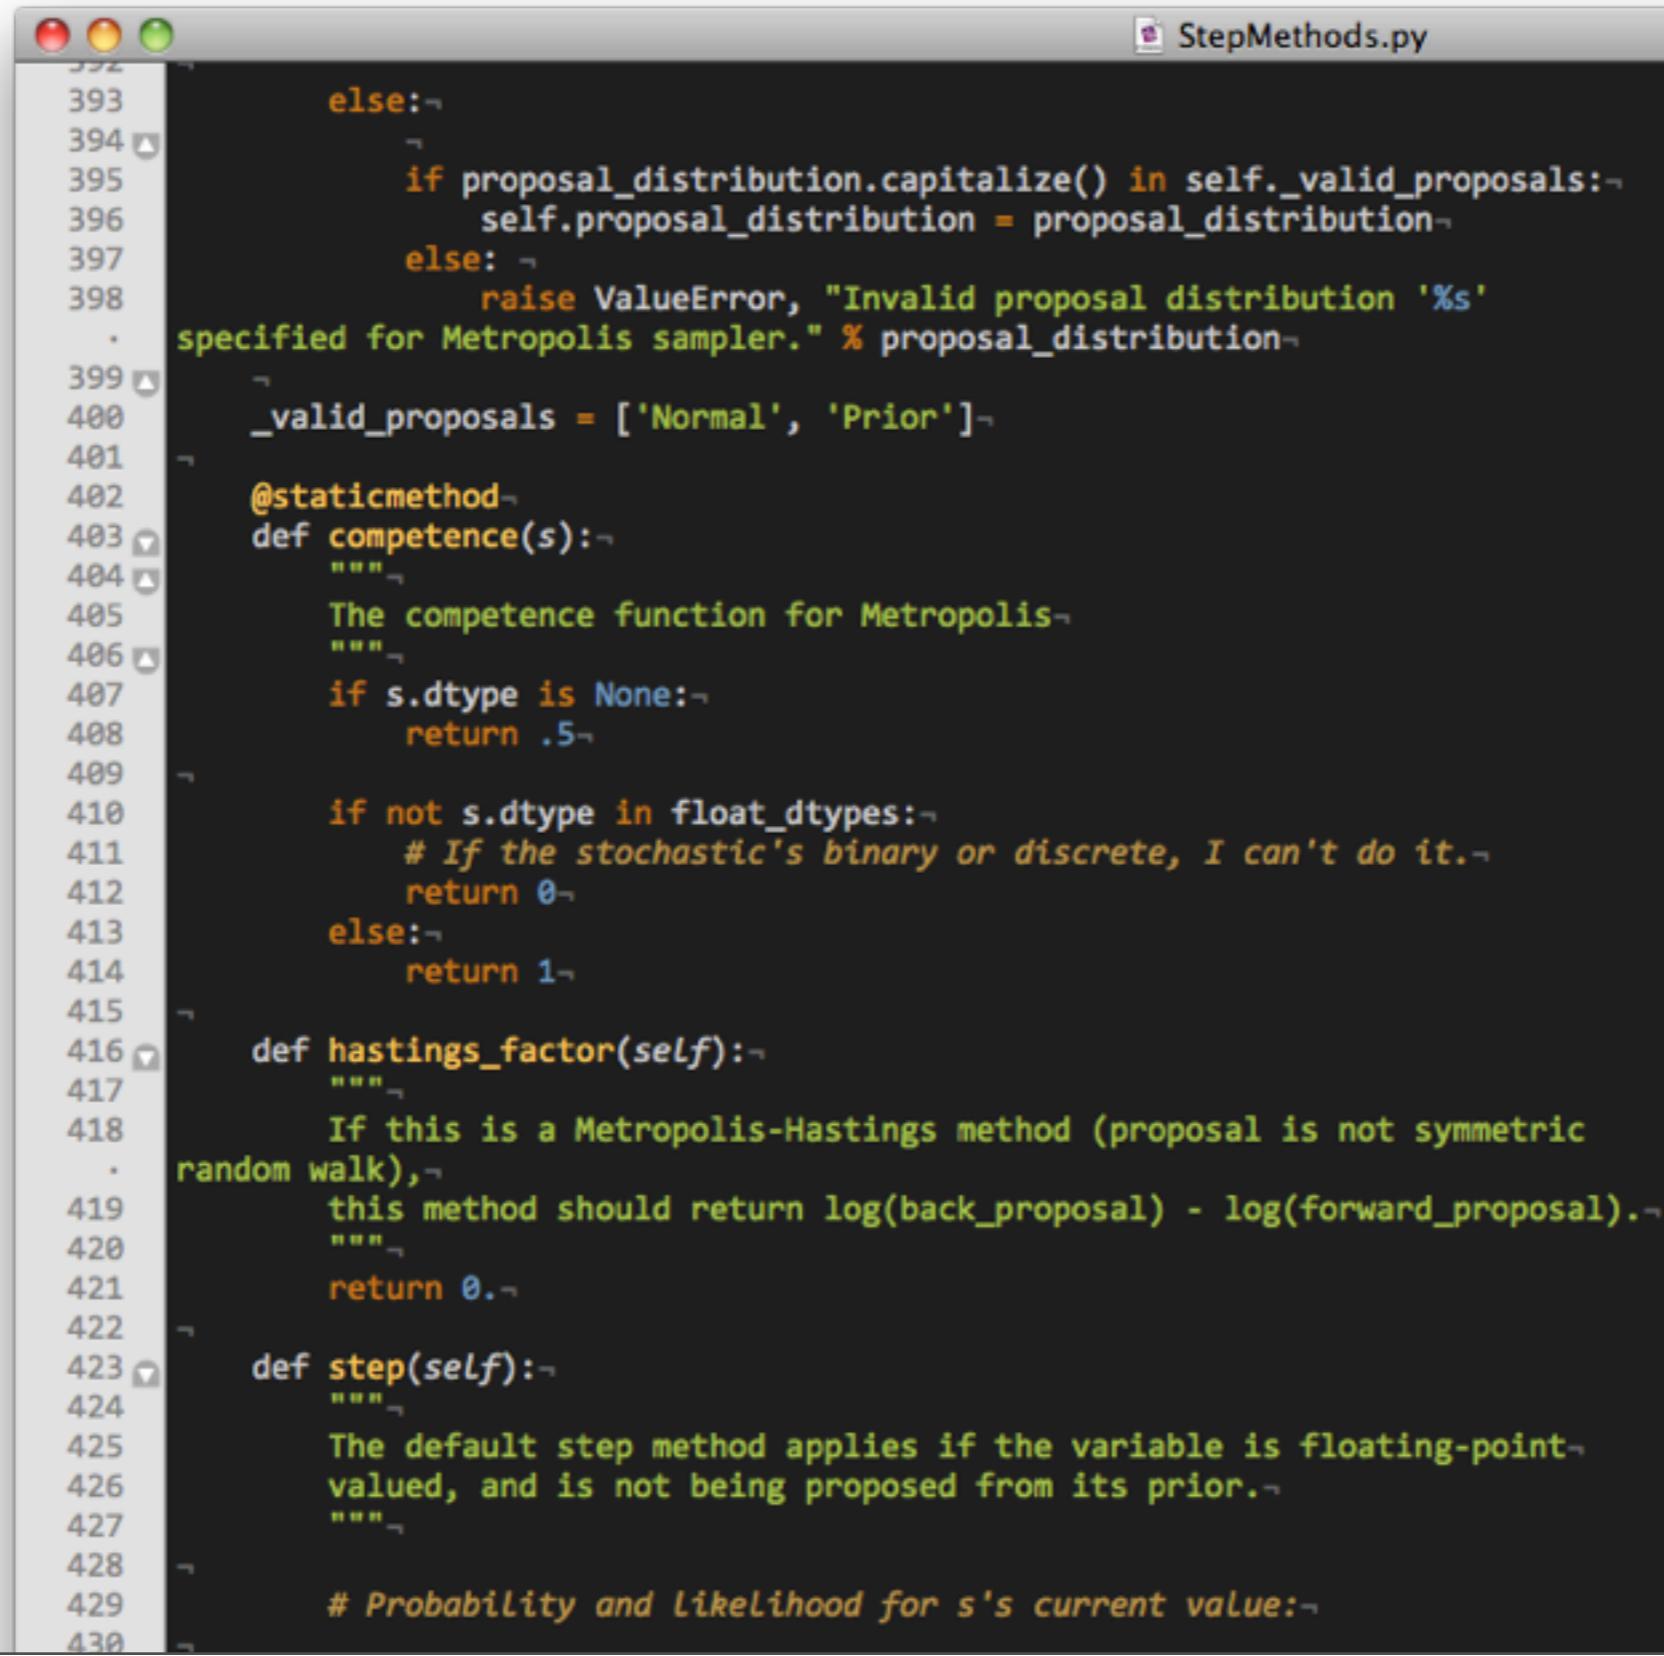


The screenshot shows a Python code editor window titled "StepMethods.py". The code is a part of a class definition for a "StepMethod" object. It includes a detailed docstring for the "tune" method, which performs variance adaptation based on acceptance rates. The code also handles verbose output and tuning flags.

```
514
515     def tune(self, divergence_threshold=1e10, verbose=0):-
516         """
517             Tunes the scaling parameter for the proposal distribution-
518             according to the acceptance rate of the last k proposals:-
519
520             Rate      Variance adaptation-
521             ----- -
522             <0.001      x 0.1-
523             <0.05       x 0.5-
524             <0.2        x 0.9-
525             >0.5        x 1.1-
526             >0.75       x 2-
527             >0.95       x 10-
528
529             This method is called exclusively during the burn-in period of the-
530             sampling algorithm.-.
531
532             May be overridden in subclasses.-.
533             """
534
535     if self.verbose is not None:-
536         verbose = self.verbose-
537
538     if self.verbose is not None:-
539         verbose = self.verbose-
540
541     # Verbose feedback-
542     if verbose > 0:-.
543         print '\t%s tuning:' % self._id-
544
545     # Flag for tuning state-
546     tuning = True-
547
548     # Calculate recent acceptance rate-
549     if not (self.accepted + self.rejected): return tuning-
550     acc_rate = self.accepted / (self.accepted + self.rejected)-
```

```
514
515     def tune(self, divergence_threshold=1e10, verbose=0):-
516         """
517             Tunes the scaling parameter for the proposal distribution-
518             according to the acceptance rate of the last k proposals:-
519
520             Rate      Variance adaptation-
521             ----      -----
522             <0.001      x 0.1-
523             <0.05       x 0.5-
524             <0.2        x 0.9-
525             >0.5        x 1.1-
526             >0.75       x 2-
527             >0.95       x 10-
528
529             This method is called exclusively during the burn-in period of the-
530             sampling algorithm.-.
531
532             May be overridden in subclasses.-.
533             """
534
535         if self.verbose is not None:-
536             verbose = self.verbose-
537
538         if self.verbose is not None:-
539             verbose = self.verbose-
540
541         # Verbose feedback-
542         if verbose > 0:-.
543             print '\t%s tuning:' % self._id-
544
545         # Flag for tuning state-
546         tuning = True-
547
548         # Calculate recent acceptance rate-
549         if not tuning: return
```

StepMethod methods



A screenshot of a Mac OS X desktop showing a code editor window titled "StepMethods.py". The code is written in Python and defines several methods for a "StepMethod" class. The methods include `proposal_distribution`, `competence`, `hastings_factor`, and `step`. The code uses standard Python syntax with docstrings and type annotations.

```
392
393     else:-
394         if proposal_distribution.capitalize() in self._valid_proposals:-
395             self.proposal_distribution = proposal_distribution-
396         else: -
397             raise ValueError, "Invalid proposal distribution '%s'-
398 specified for Metropolis sampler." % proposal_distribution-
399
400     _valid_proposals = ['Normal', 'Prior']-
401
402     @staticmethod-
403     def competence(s):-
404         """-
405             The competence function for Metropolis-
406         """
407         if s.dtype is None:-.
408             return .5-
409
410         if not s.dtype in float_dtypes:-
411             # If the stochastic's binary or discrete, I can't do it.-.
412             return 0-
413         else:-
414             return 1-
415
416     def hastings_factor(self):-
417         """-
418             If this is a Metropolis-Hastings method (proposal is not symmetric-
419             random walk),-
420             this method should return log(back_proposal) - log(forward_proposal).-
421         """
422         return 0.-.
423
424     def step(self):-
425         """-
426             The default step method applies if the variable is floating-point-
427             valued, and is not being proposed from its prior.-.
428         """
429
430         # Probability and Likelihood for s's current value:-.
```

```
401
402     @staticmethod
403     def competence(s):-
404         """
405             The competence function for Metropolis-
406         """
407         if s.dtype is None:-
408             return .5
409
410         if not s.dtype in float_dtypes:-
411             # If the stochastic's binary or discrete, I can't do it.-.
412             return 0
413         else:-.
414             return 1
415
416     def hastings_factor(self):-
417         """
418             If this is a Metropolis-Hastings method (proposal is not symmetric
random walk),-
419             this method should return log(back_proposal) - log(forward_proposal).-
420         """
421         return 0.
422
423     def step(self):-
424         """
425             The default step method applies if the variable is floating-point-
valued, and is not being proposed from its prior.-.
426         """
427
428
429         # Probability and Likelihood for s's current value:-
430
431         if self.verbose>2:-
432             print-
433             print self._id + ' getting initial logp.'-
434
435         if self.proposal_distribution == "Prior":-
```

Competence

- 0** = I cannot safely handle this variable.
- 1** = I can handle the variable about as well as the standard Metropolis step method.
- 2** = I can do better than Metropolis.
- 3** = I am the best step method you are likely to find for this variable in most cases.

```
S.use_step_method(better_method,  
                  my_stochastic,  
                  *args,  
                  **kwargs)
```

```
>>> mu = pymc.Normal('mu', 0, .01,  
value=0)  
  
>>> tau = pymc.Exponential('tau', .01,  
value=1)  
  
>>> cutoff = pymc.Exponential('cutoff', 1,  
value=1.3)  
  
>>> D = pymc.TruncatedNormal('D', mu, tau,  
-numpy.inf, cutoff, value=data,  
observed=True)
```

$$q(x', x^{(j)}) = \text{TruncatedNormal}(x, \sigma, \max(D), \infty)$$

truncated normal proposal

truncated_metropolis.py

```
1 class TruncatedMetropolis(pymc.Metropolis):-
2     def __init__(self, stochastic, Low_bound, up_bound, *args, **kwargs):-
3         self.low_bound = low_bound-
4         self.up_bound = up_bound-
5         pymc.Metropolis.__init__(self, stochastic, *args, **kwargs)-.
6
7     # Propose method generates proposal values-
8     def propose(self):-
9         tau = 1./(self.adaptive_scale_factor * self.proposal_sd)**2-
10        self.stochastic.value = \
11            pymc.rtruncnorm(self.stochastic.value, tau, self.low_bound,
12            self.up_bound)-.
13
14    # Hastings factor method accounts for asymmetric proposal distribution-
15    def hastings_factor(self):-
16        tau = 1./(self.adaptive_scale_factor * self.proposal_sd)**2-
17        cur_val = self.stochastic.value-
18        last_val = self.stochastic.last_value-
19
20        lp_for = pymc.truncnorm_like(cur_val, last_val, tau, \
21            self.low_bound, self.up_bound)-.
22        lp_bak = pymc.truncnorm_like(last_val, cur_val, tau, \
23            self.low_bound, self.up_bound)-.
24
25        if self.verbose > 1:-
26            print self._id + ': Hastings factor %f'%(lp_bak - lp_for)-.
return lp_bak - lp_for
```

Line: 26 Column: 31

Python



Soft Tabs:

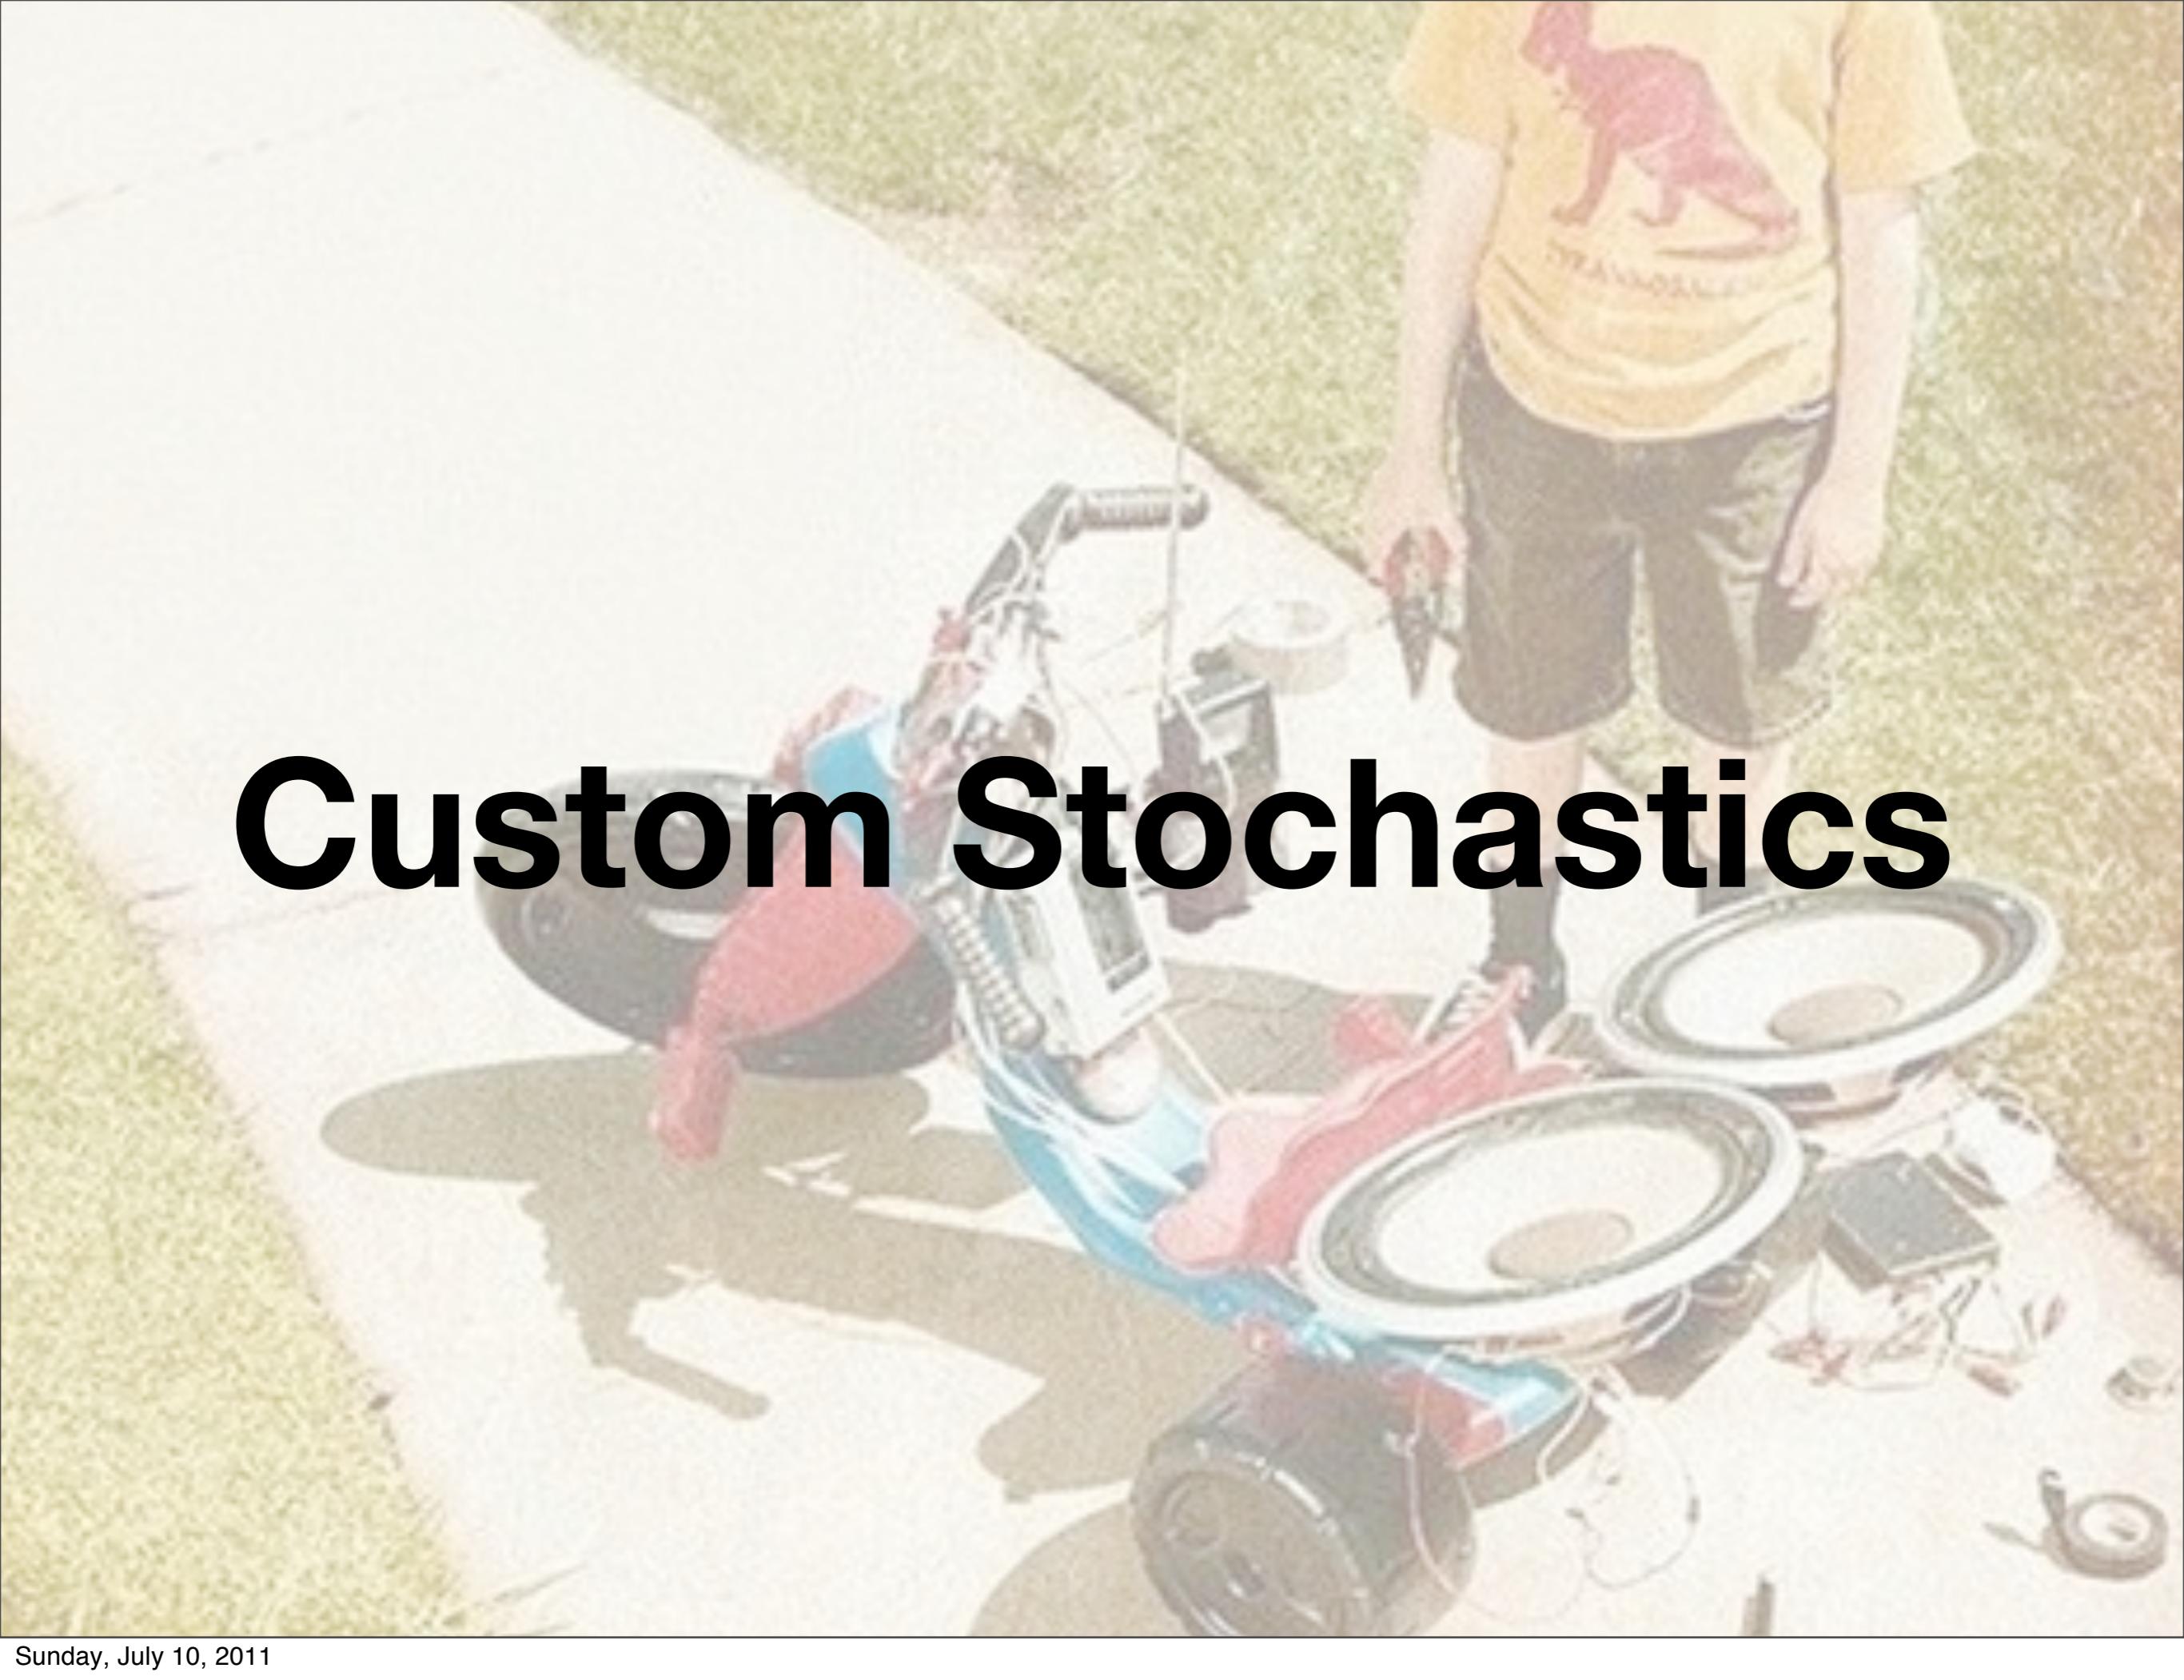
4



hastings_factor(self)



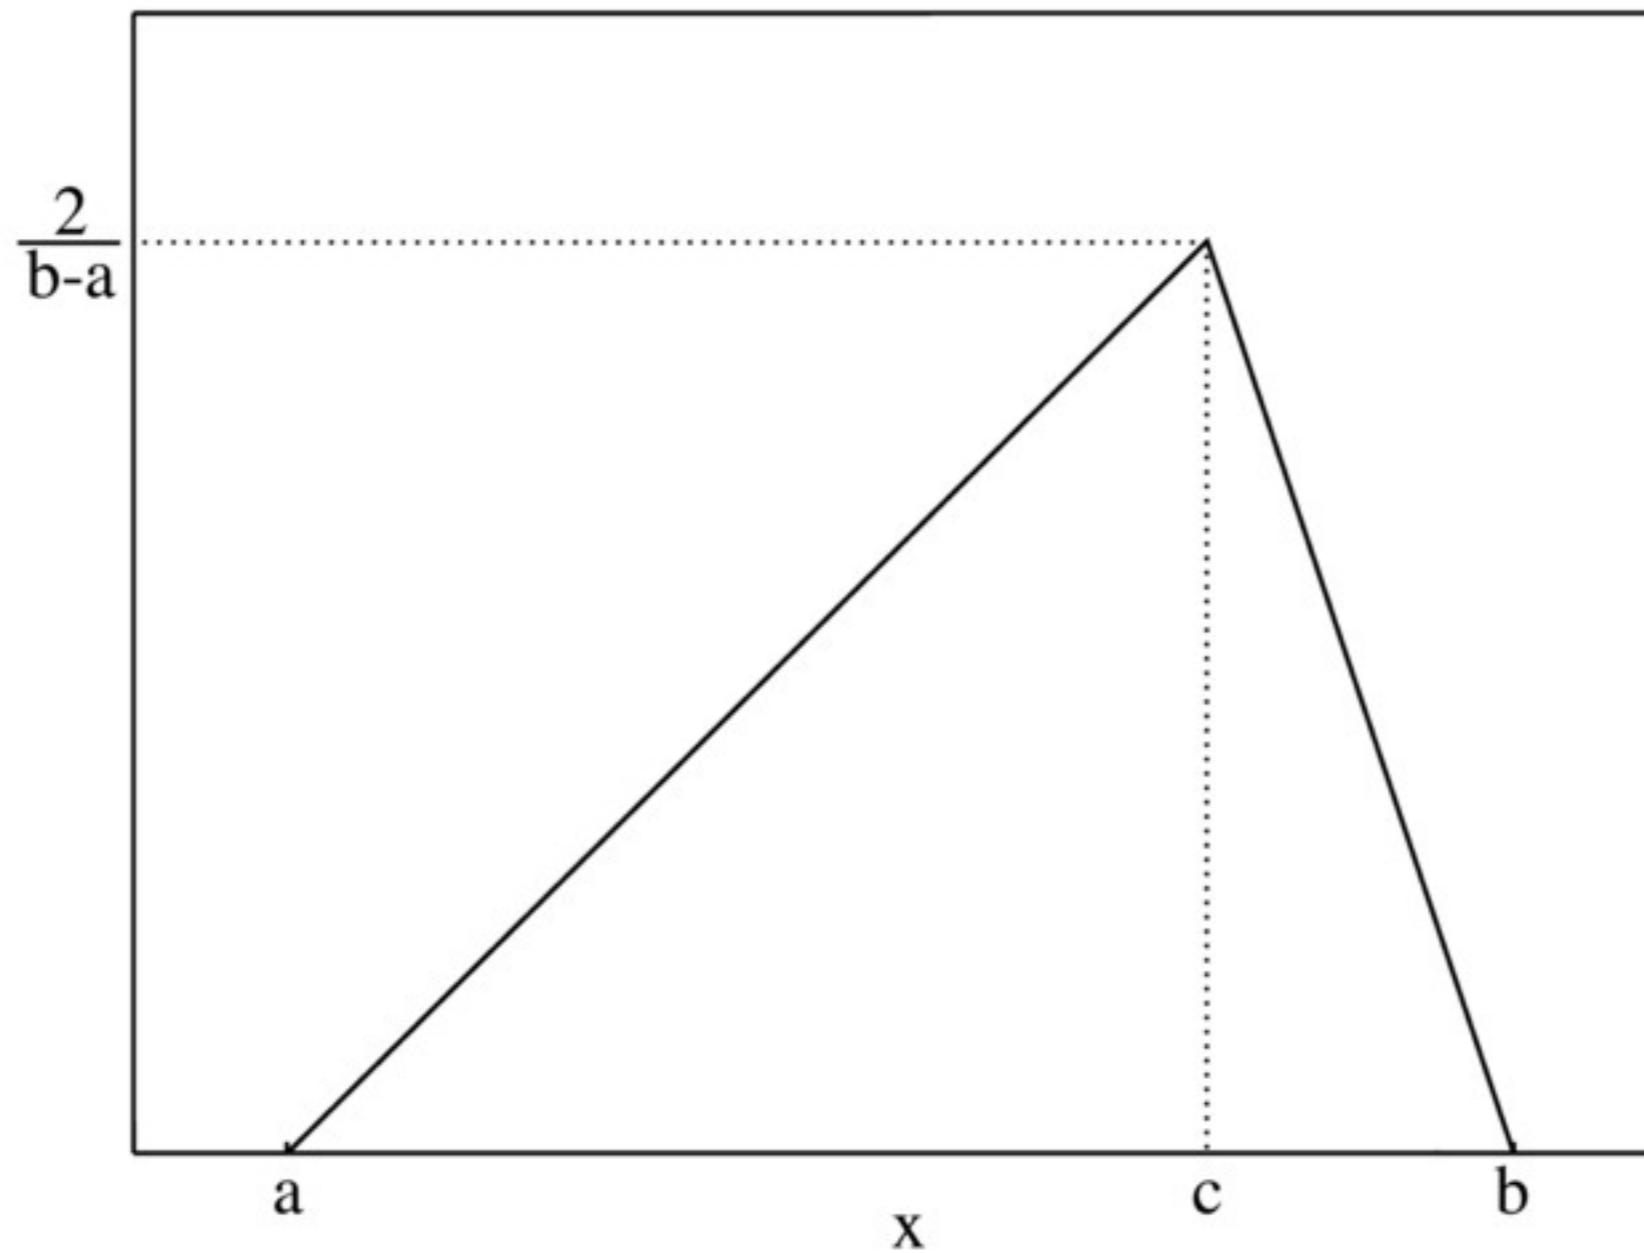
Custom Stochastics



Stochastic attributes

- (1) name
- (2) logp method
- (3) random method
- (4) dtype
- (5) mv flag

example



triangle distribution

triangle distribution

$$f(x) = \begin{cases} \frac{2(x - \min)}{(\text{mode} - \min)(\max - \min)} & \text{if } \min \leq x \leq \text{mode} \\ \frac{2(\max - x)}{(\max - \min)(\max - \text{mode})} & \text{if } \text{mode} \leq x \leq \max \end{cases}$$

distribution function

```
triangular.py
1 from numpy import log, random, sqrt, zeros, atleast_1d
2
3 def triangular_like(x, mode, minval, maxval):
4     """Log-likelihood of triangular distribution"""
5
6     x = atleast_1d(x)
7
8     # Check for support
9     if any(x<minval) or any(x>maxval): return -inf
10
11     # Likelihood of Left values
12     like = sum(log(2*(x[x<=mode] - minval)) - log(mode - minval) - log(maxval - minval))
13
14     # Likelihood of right values
15     like += sum(log(2*(maxval - x[x>mode])) - log(maxval - minval) - log(maxval - mode))
16
17     return like
18
19 def rtriangular(mode, minval, maxval, size=1):
20     """Generate triangular random numbers"""
21
22     # Uniform random numbers
23     z = atleast_1d(random.random(size))
24
25     # Threshold for transformation
26     threshold = (mode - minval)/(maxval - minval)
27
28     # Transform uniforms
29     u = atleast_1d(zeros(size))
30     u[z<=threshold] = minval + sqrt(z[z<=threshold]*(maxval - minval)*(mode - minval))
31     u[z>threshold] = maxval - sqrt((1 - z[z>threshold])*(maxval - minval)*(maxval - mode))
32
33     return u
```

Line: 6 Column: 21

Python

Soft Tabs: 4 triangular_like(x, mode, minval, maxval)

 triangular.py

```
from numpy import log, random, sqrt, zeros, atleast_1d\n\n\ndef triangular_like(x, mode, minval, maxval):\n    """Log-likelihood of triangular distribution"""\n\n    x = atleast_1d(x)\n\n    # Check for support-\n    if any(x<minval) or any(x>maxval): return -inf\n\n    # Likelihood of left values-\n    like = sum(log(2*(x[x<=mode] - minval)) - log(mode - minval) - log(maxval - minval))\n\n    # Likelihood of right values\n    like += sum(log(2*(maxval - x[x>mode]))) - log(maxval - minval) - log(maxval - mode)\n\n    return like\n\n\ndef rtriangular(mode, minval, maxval, size=1):\n    """Generate triangular random numbers"""\n\n    # Uniform random numbers-\n    z = atleast_1d(random.random(size))
```

```
# Likelihood of right values
like += sum(log(2*(maxval - x[x>mode])) - log(maxval - minval) -
log(maxval - mode))-
-
return like

def rtriangular(mode, minval, maxval, size=1):
    """Generate triangular random numbers"""

    # Uniform random numbers
    z = atleast_1d(random.random(size))-

    # Threshold for transformation
    threshold = (mode - minval)/(maxval - minval)-

    # Transform uniforms
    u = atleast_1d(zeros(size))-#
    u[z<=threshold] = minval + sqrt(z[z<=threshold]*(maxval - minval)*(mode - minval))-#
    u[z>threshold] = maxval - sqrt((1 - z[z>threshold])*(maxval - minval)*(maxval - mode))-#
    -
    return u
```

6 Column: 21 | Python | Soft Tabs: 4 | triangular_like(x, mode, minval, maxval)

```
>>> Triangular = pymc.stochastic_from_dist  
('Triangular', logp=triangular_like,  
random=rtriangular, dtype=float)
```

```
>>> x = Triangular('x', mode=3, minval=1,  
maxval=10)  
  
>>> x.value  
array(6.0565769362029522)  
  
>>> x.logp  
-2.0779384017901741  
  
>>> x.random()  
array(5.2428089474233568)
```

Gaussian Processes



SciPy 2011

Having observed data
 $\{\mathbf{X}, \mathbf{Y}\}$, we want to find
suitable function \mathbf{f}

$$Y = f(X) + \epsilon$$

Gaussian process

“a collection of random variables, any finite number of which have a joint Gaussian distribution”

“a distribution over functions”

Gaussian process

$$f \sim \mathcal{GP}(m(\mathbf{x}), C(\mathbf{x}, \mathbf{x}))$$

Gaussian process

$$f \sim \mathcal{GP}(m(\mathbf{x}), C(\mathbf{x}, \mathbf{x}))$$

mean function

covariance function

The diagram illustrates the components of a Gaussian process. It shows the mathematical expression $f \sim \mathcal{GP}(m(\mathbf{x}), C(\mathbf{x}, \mathbf{x}))$. A vertical dashed line with two dots at its ends connects the term $m(\mathbf{x})$ to the term $C(\mathbf{x}, \mathbf{x})$, indicating they are the mean function and covariance function respectively.

$$P(f(\mathbf{x}))|y_N, X_N) = \frac{P(y_N|f(\mathbf{x}), X_N)P(f(\mathbf{x}))}{P(y_N|X_N)}$$

GPP

$$P(f(\mathbf{x}))|y_N, X_N) = \frac{P(y_N|f(\mathbf{x}), X_N)P(f(\mathbf{x}))}{P(y_N|X_N)}$$



Mean function

A screenshot of a Python code editor window titled "gpexample.py". The code defines a function "quadfun" which returns a quadratic expression, and then creates a "Mean" object using this function with parameters $a = 1.$, $b = .5$, and $c = 2.$. The code editor interface includes a toolbar at the top, a status bar at the bottom showing "Line: 7 Column: 42", and tabs for "Python" and "quadfun(x, a, b, c)" at the bottom right.

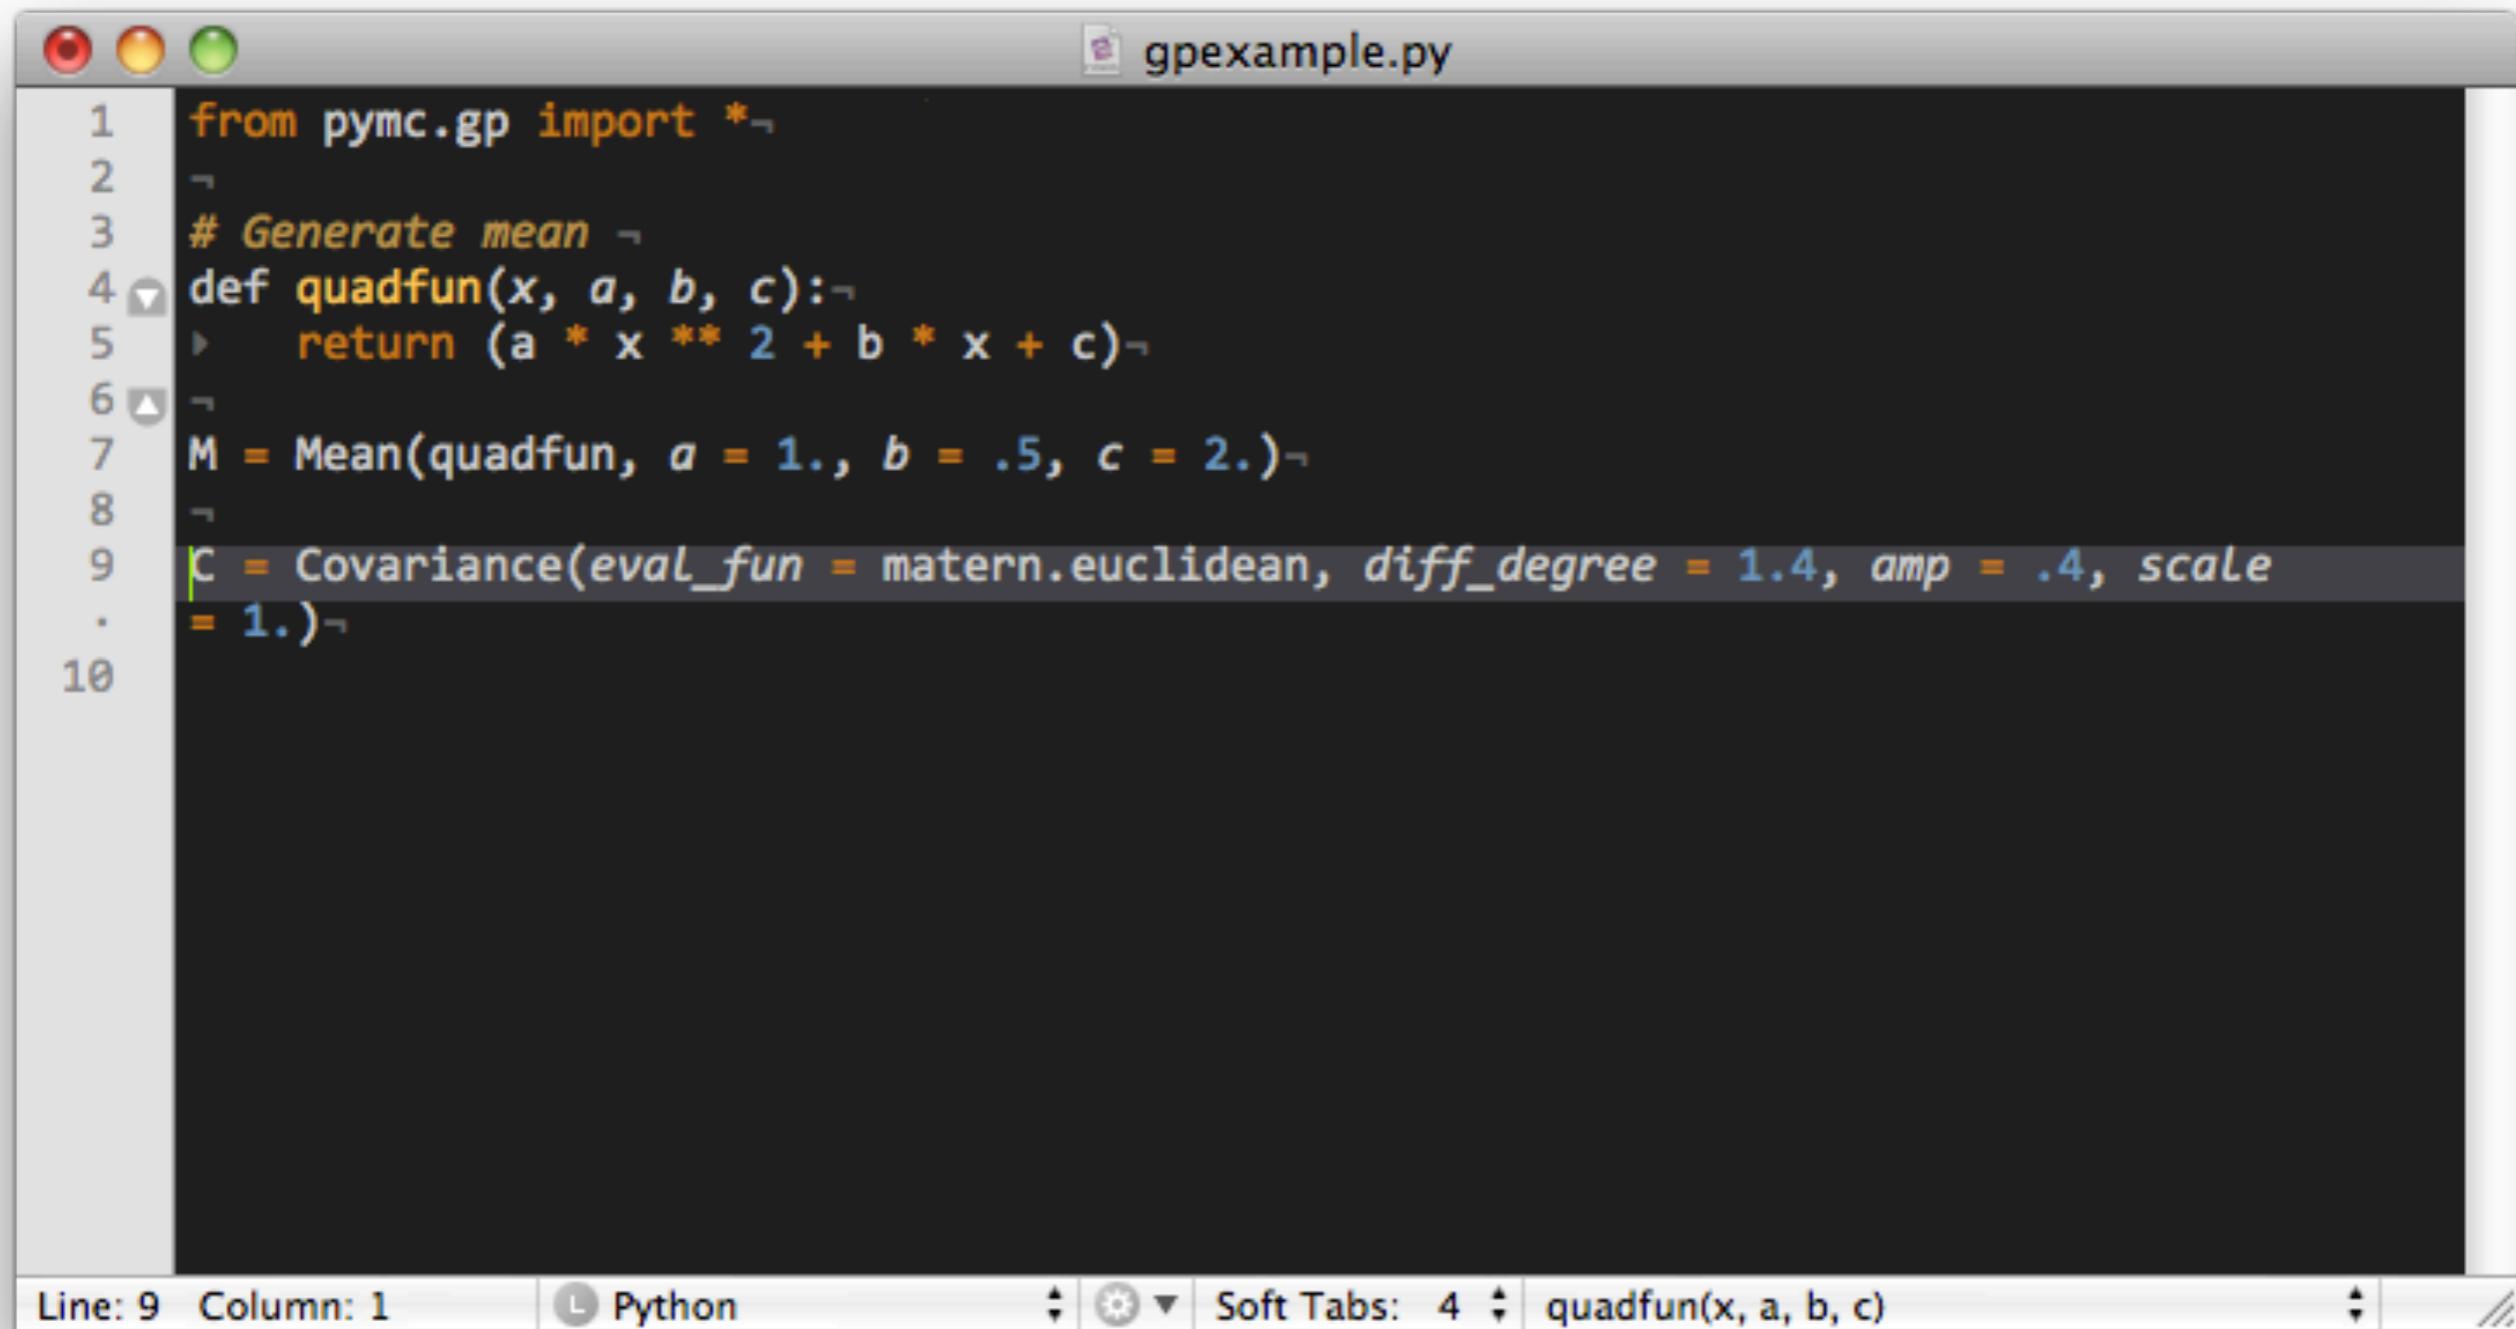
```
1 from pymc.gp import *
2
3 # Generate mean
4 def quadfun(x, a, b, c):
5     return (a * x ** 2 + b * x + c)
6
7 M = Mean(quadfun, a = 1., b = .5, c = 2.)
```

Covariance Function

$$C(d) = \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}d}{l} \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}d}{l} \right)$$

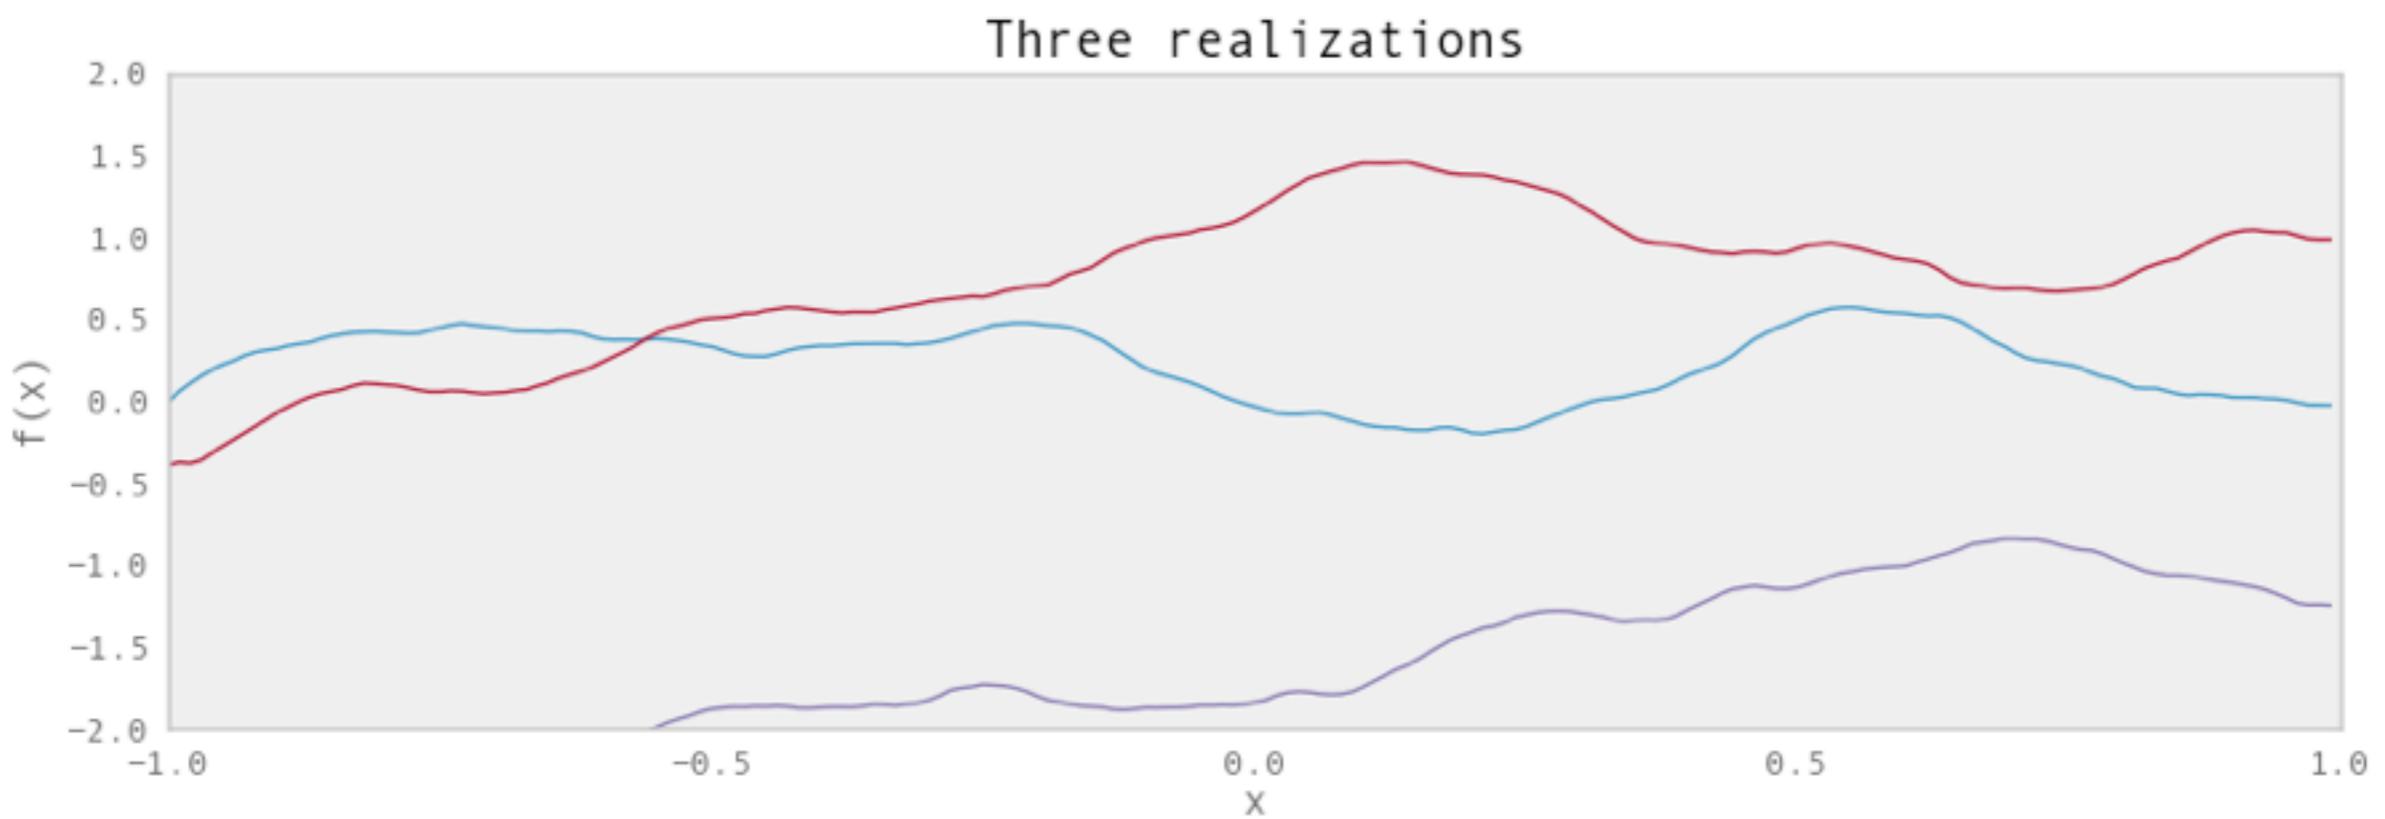
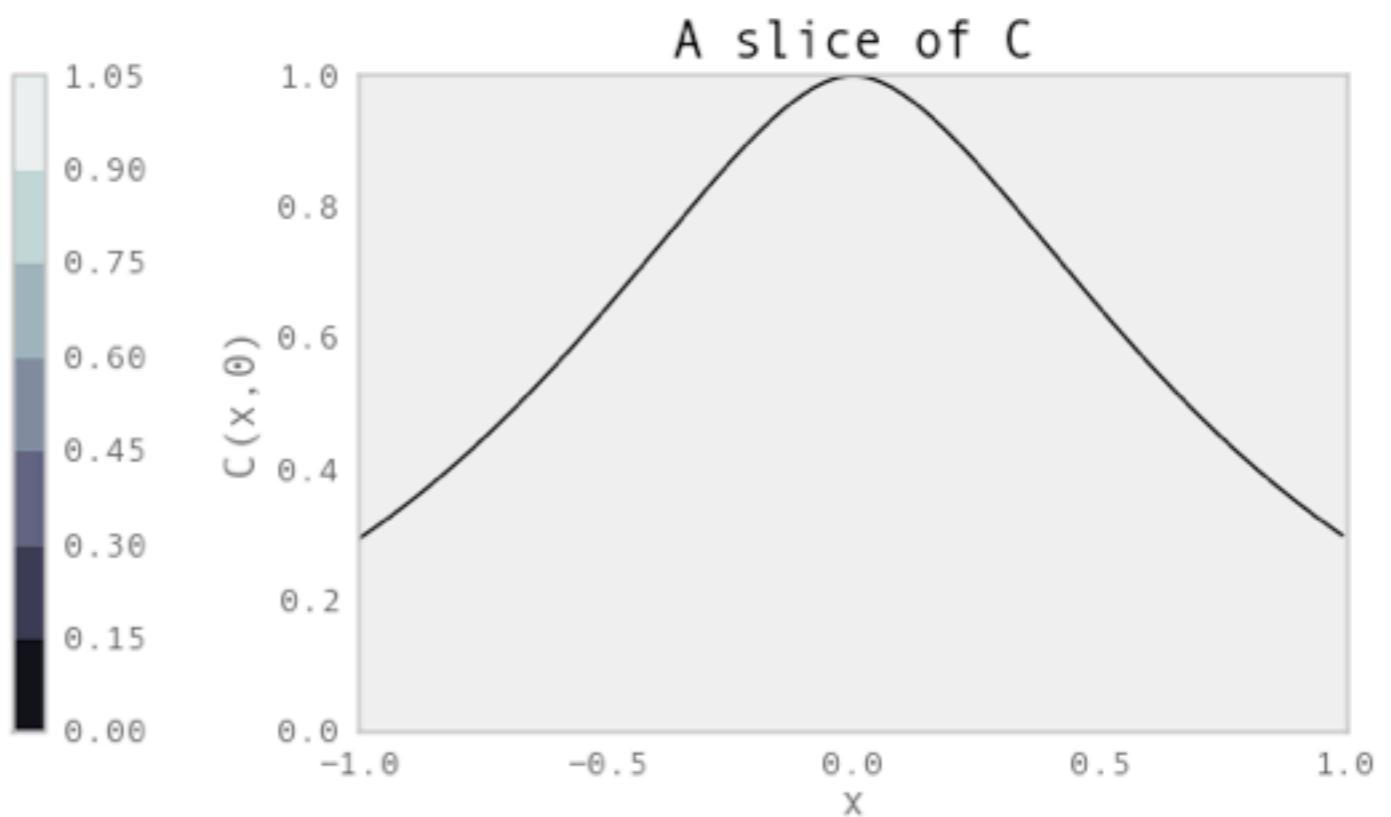
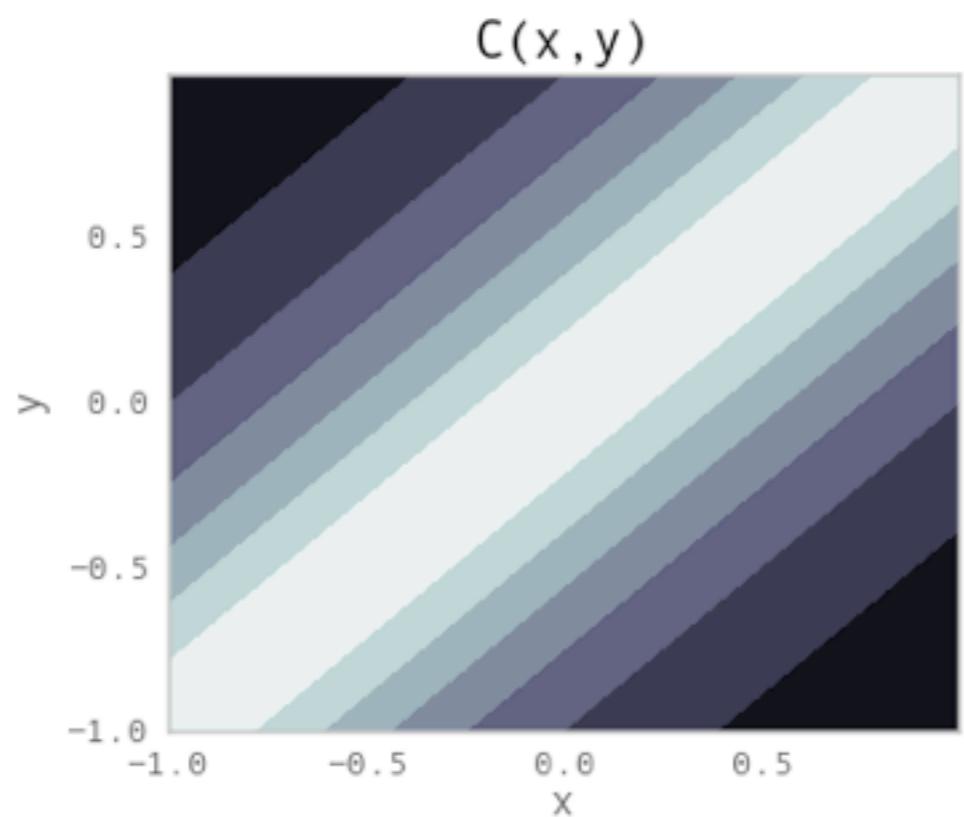
Matérn covariance

Covariance Function

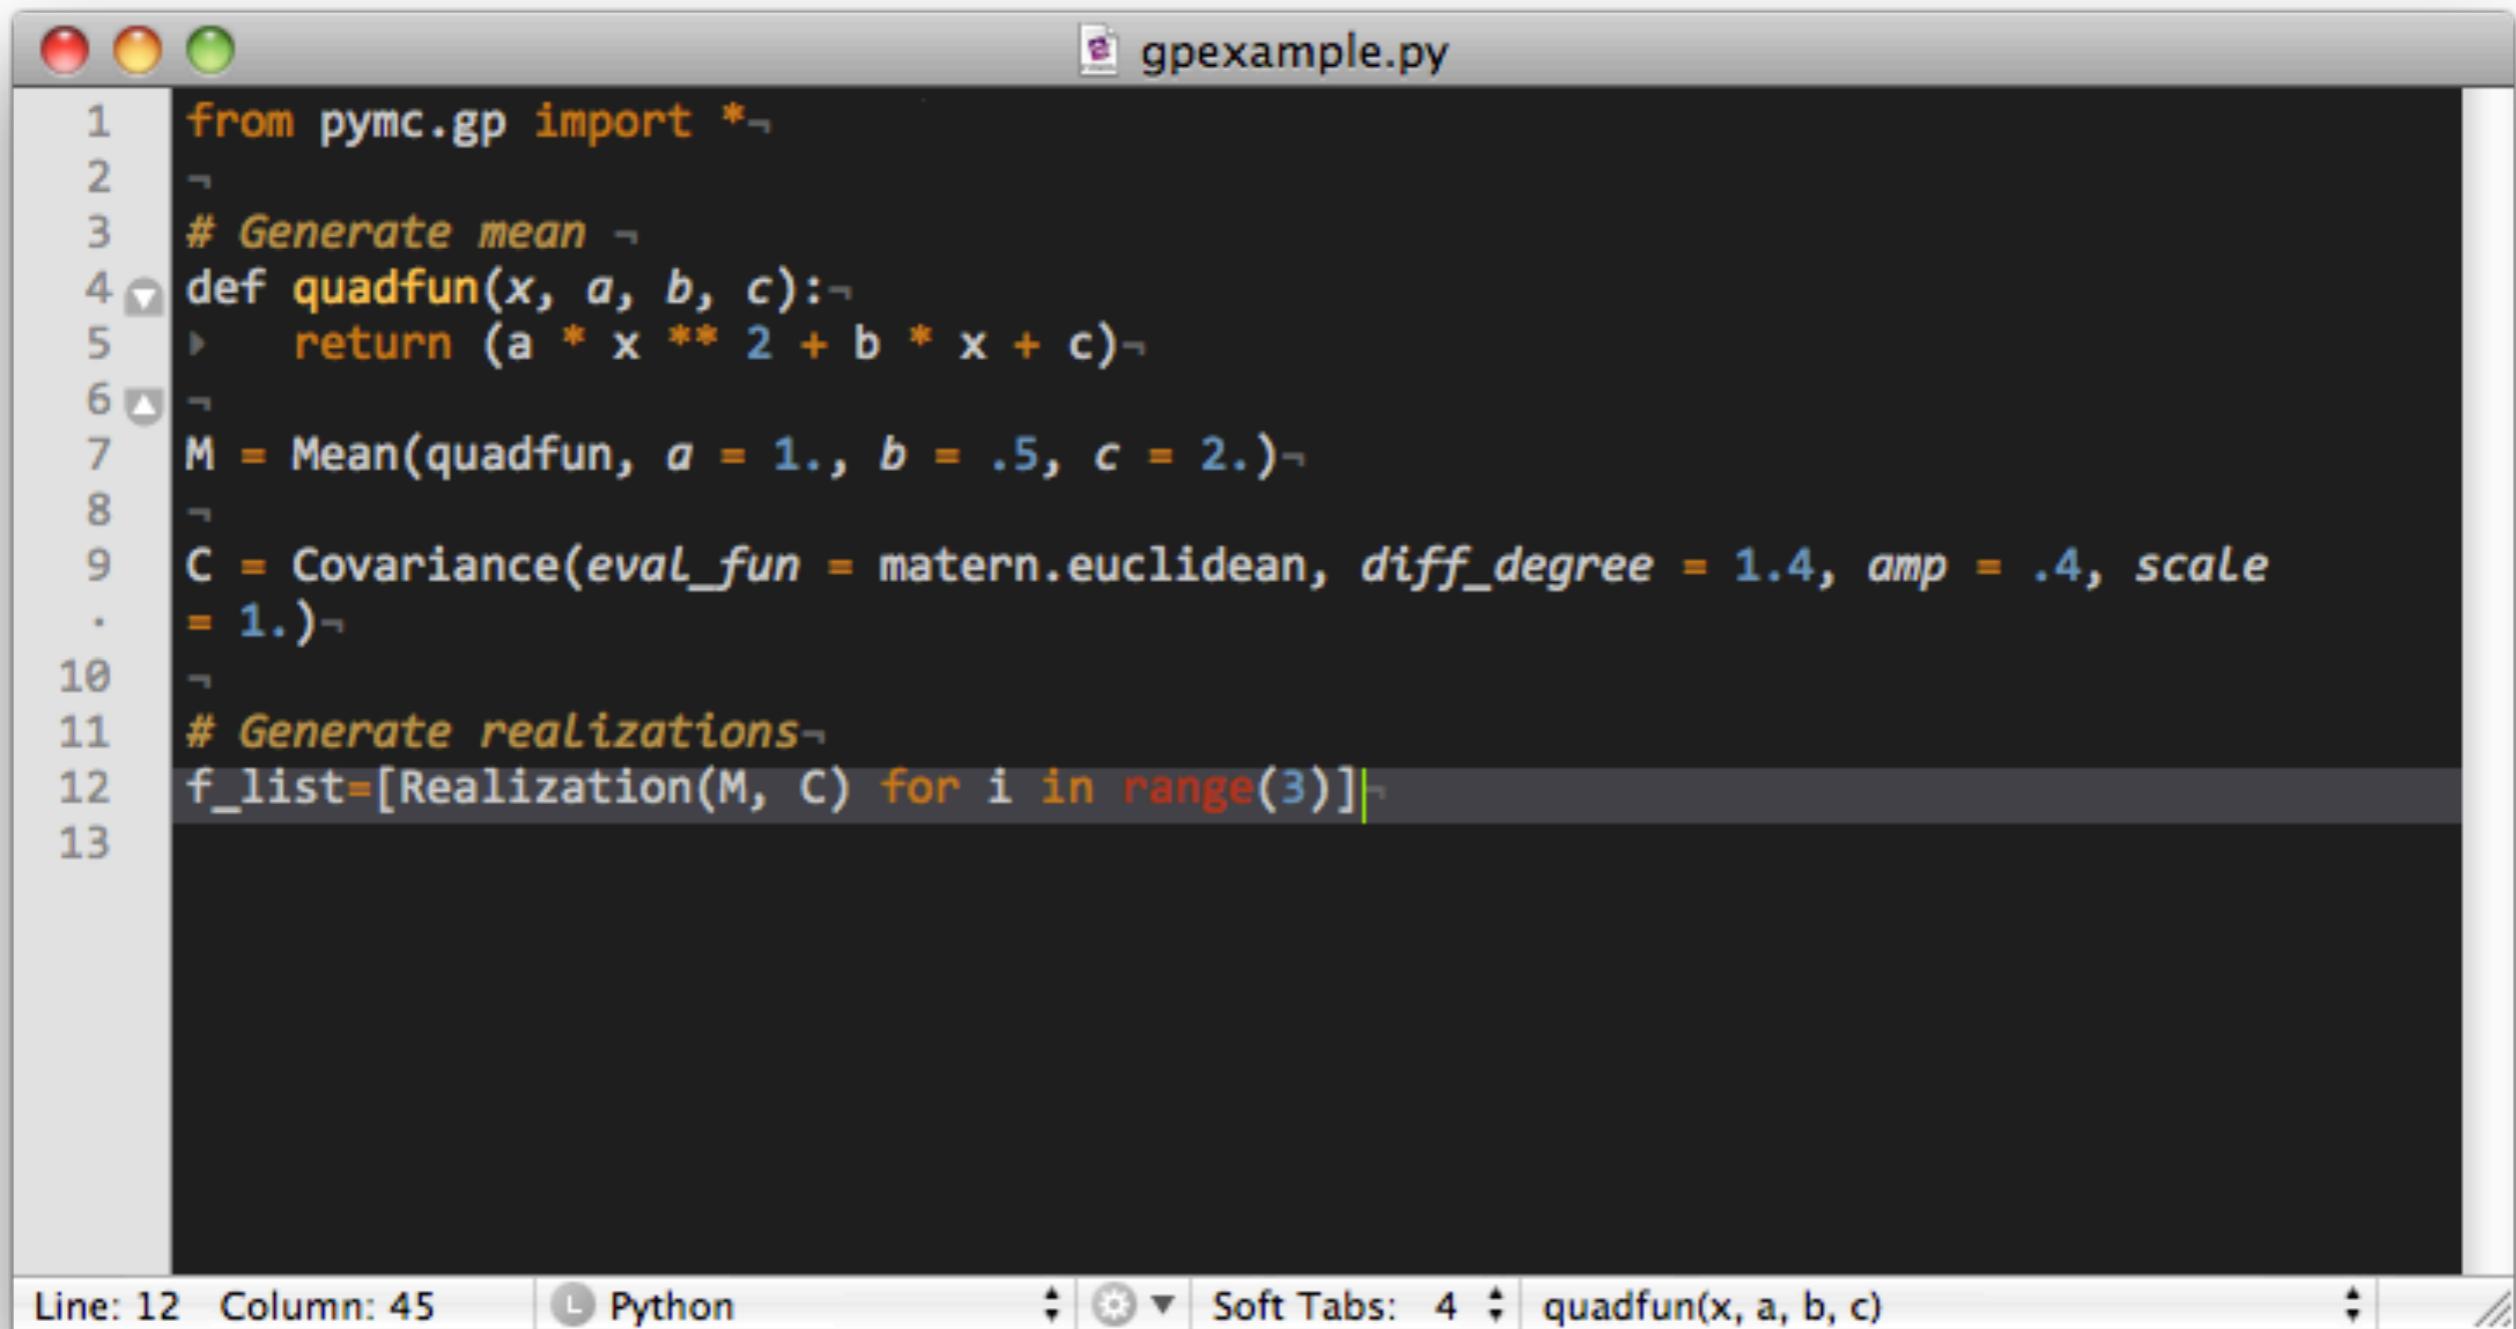


```
gpexample.py
1 from pymc.gp import *
2
3 # Generate mean
4 def quadfun(x, a, b, c):
5     return (a * x ** 2 + b * x + c)
6
7 M = Mean(quadfun, a = 1., b = .5, c = 2.)
8
9 C = Covariance(eval_fun = matern.euclidean, diff_degree = 1.4, amp = .4, scale
. = 1.)
10
```

Line: 9 Column: 1 Python Soft Tabs: 4 quadfun(x, a, b, c)



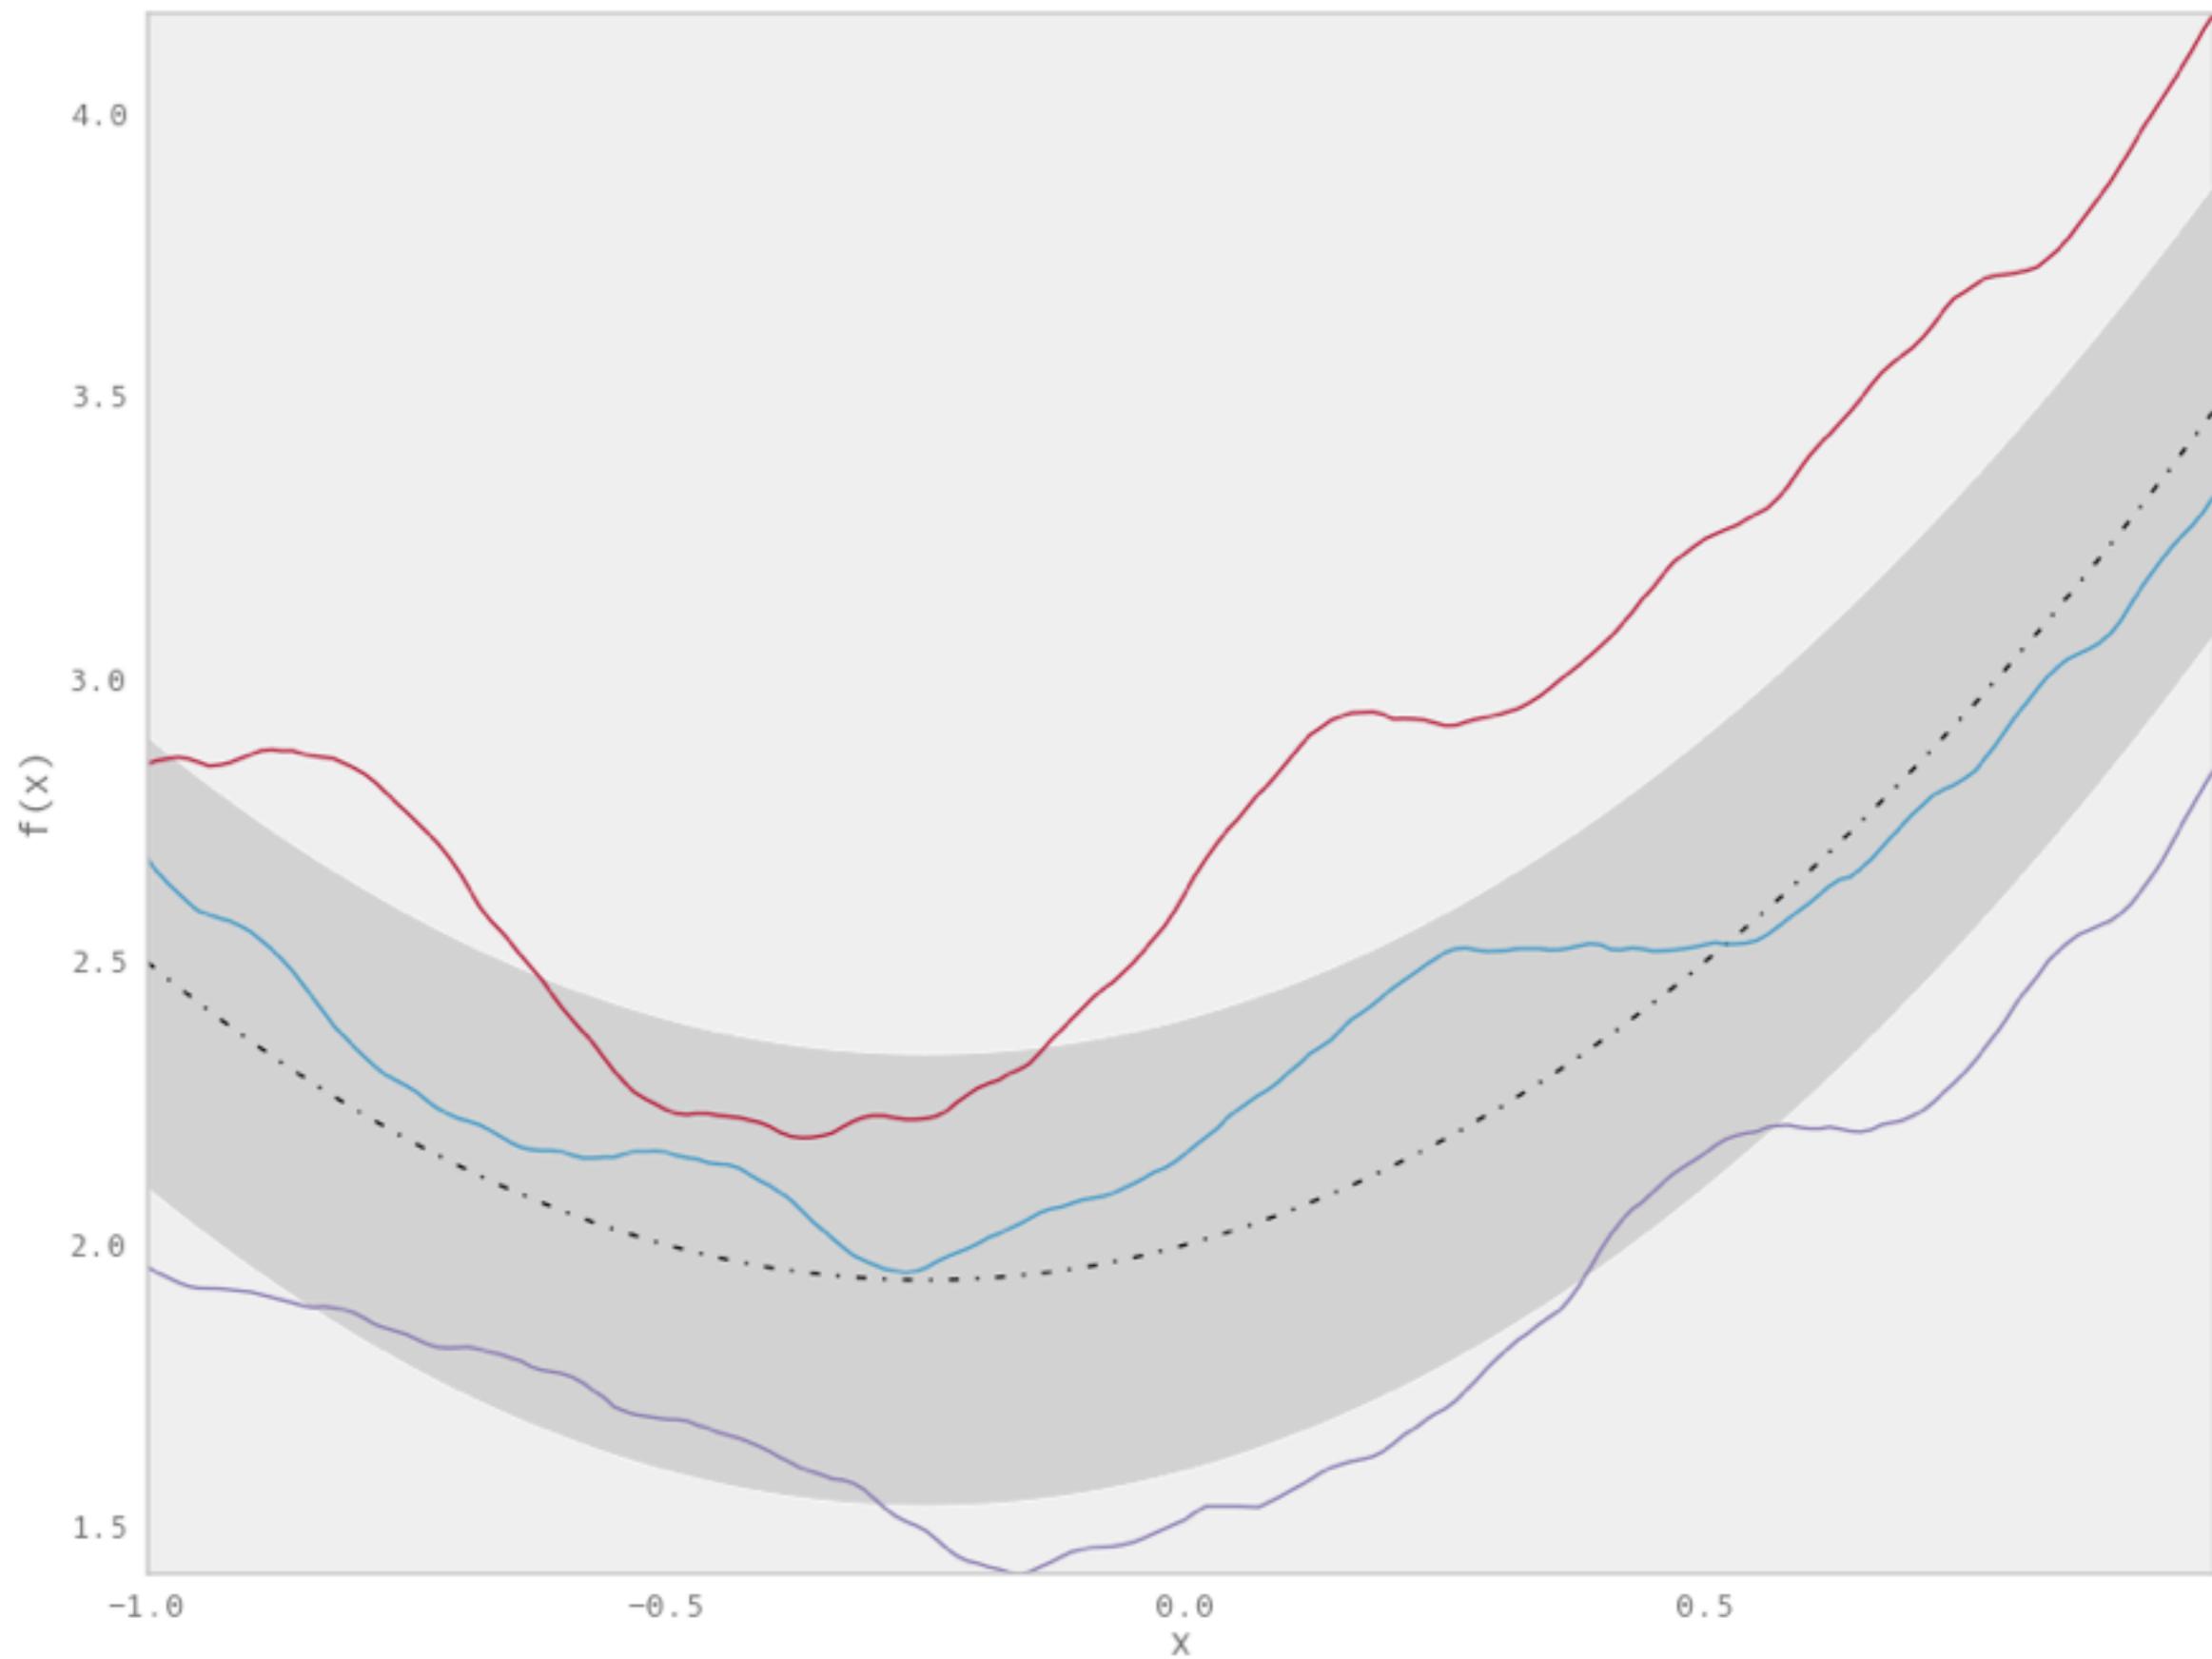
Gaussian process



```
gpexample.py
1 from pymc.gp import *
2
3 # Generate mean
4 def quadfun(x, a, b, c):
5     return (a * x ** 2 + b * x + c)
6
7 M = Mean(quadfun, a = 1., b = .5, c = 2.)
8
9 C = Covariance(eval_fun = matern.euclidean, diff_degree = 1.4, amp = .4, scale
. = 1.)
10
11 # Generate realizations
12 f_list=[Realization(M, C) for i in range(3)]
13
```

Line: 12 Column: 45 Python Soft Tabs: 4 quadfun(x, a, b, c)

Three realizations of the GP



observations

$\{x = -0.5, y = 3.1\},$
 $\{x = 0.5, y = 2.9\}$



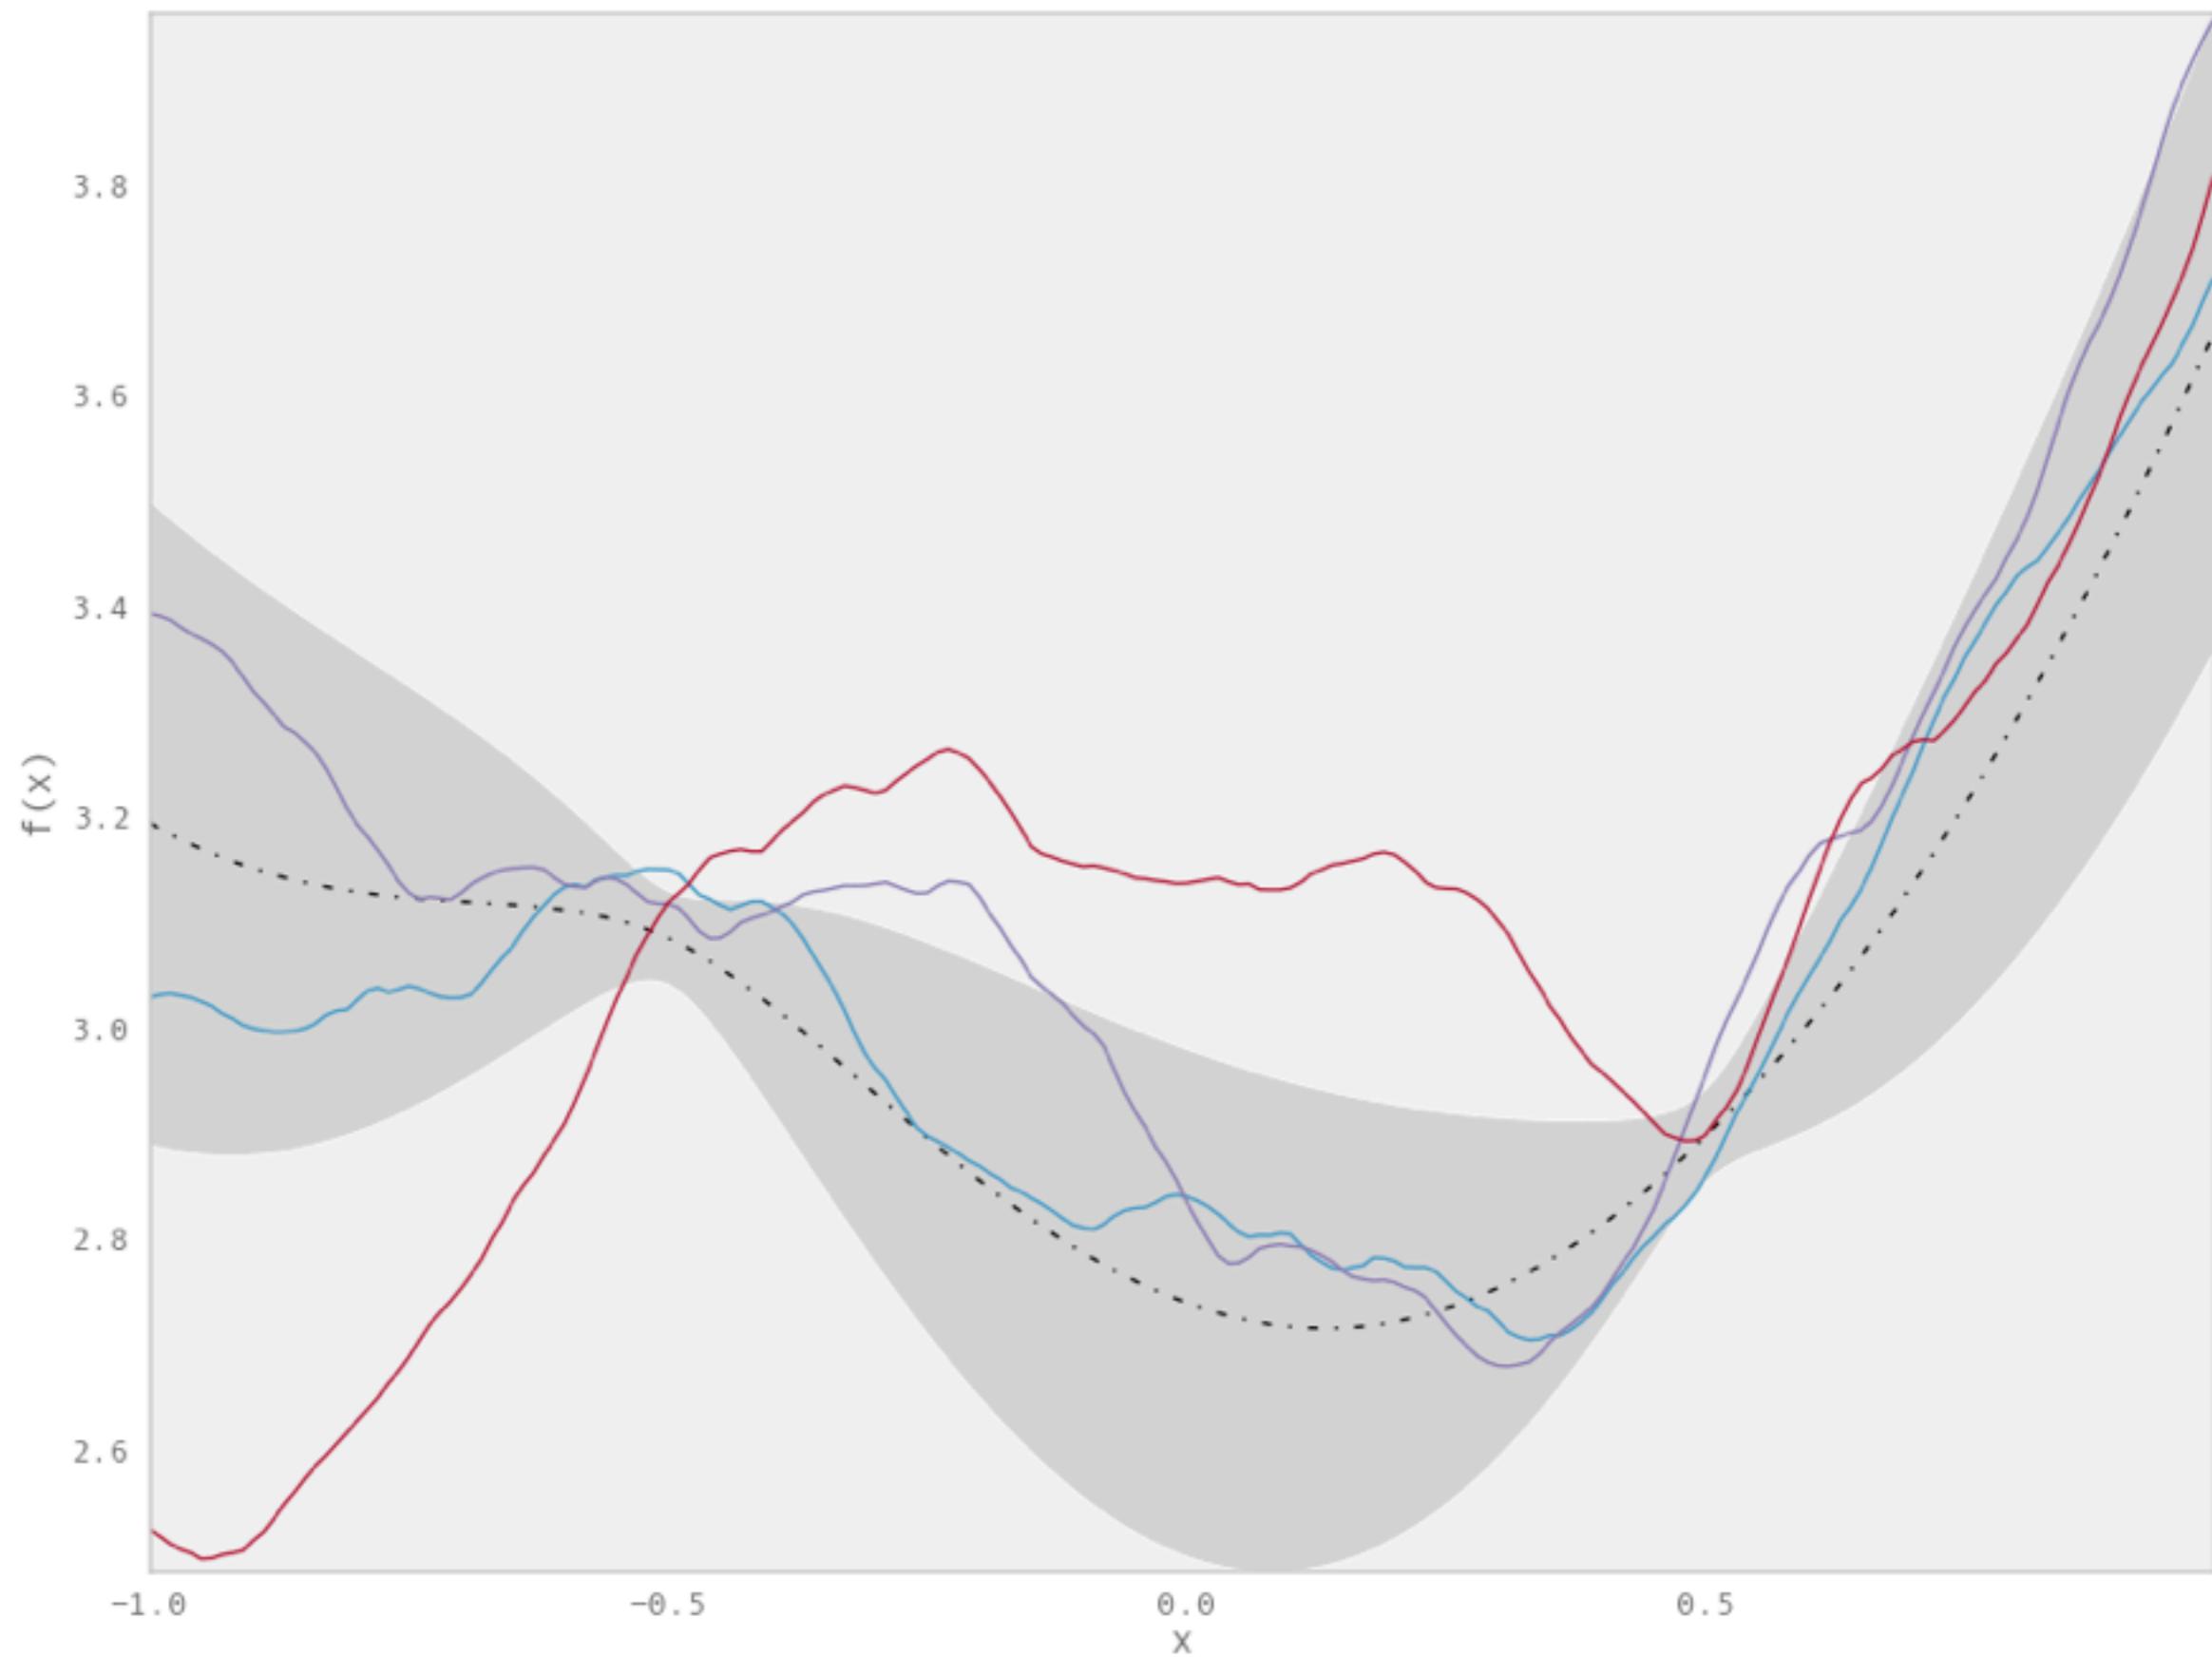
observations

The screenshot shows a window titled "obs.py" containing Python code. The code imports mean and cov from the mean and cov modules respectively, and observes data points at specific coordinates with a given covariance matrix. It then generates three realizations of the function f.

```
1 # Import the mean and covariance-
2 from mean import M-
3 from cov import C-
4 from pymc.gp import *-
5 from numpy import *-
6 -
7 # Impose observations on the GP-
8 o = array([- .5, .5])-#
9 V = array([.002, .002])-#
10 data = array([3.1, 2.9])-#
11 observe(M, C, obs_mesh=o, obs_V = V, obs_vals = data)-#
12 -
13 # Generate realizations-
14 f_list=[Realization(M, C) for i in range(3)]-
15
```

Line: 15 Column: 1 Python Soft Tabs: 4

Three realizations of the observed GP



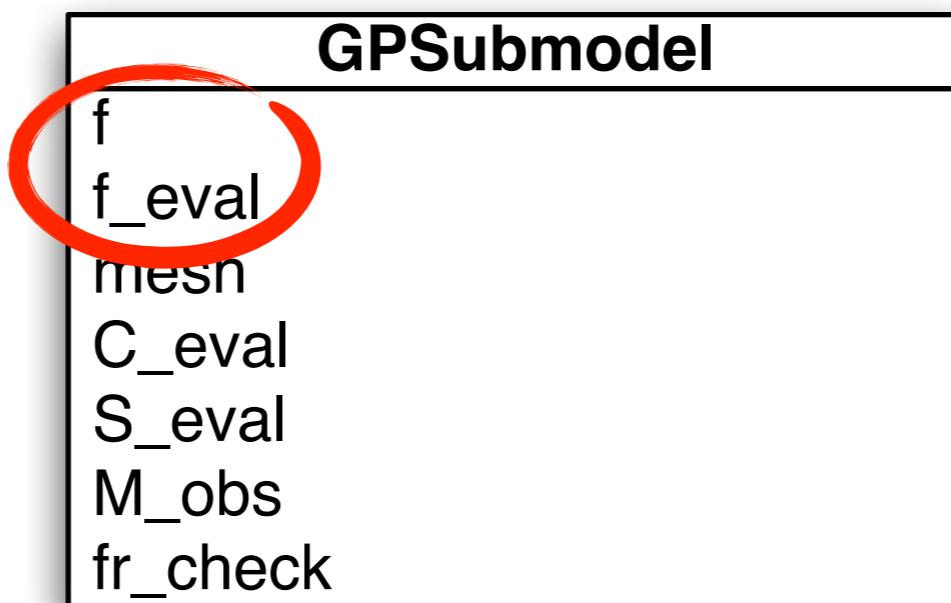
realization

$$f \sim \mathcal{GP}(m(\mathbf{x}), C(\mathbf{x}, \mathbf{x}))$$

Gaussian Process submodel

GPSubmodel
f
f_eval
mesh
C_eval
S_eval
M_obs
fr_check

Gaussian Process submodel



```
>>> sm  
<pymc(gp(gp_submodel.GPSubmodel at  
0x11796a810)>
```

```
>>> sm.f  
<pymc(gp(gp_submodel.GaussianProcess  
'sm_f' at 0x11796ab90)>
```

```
>>> sm.f_eval  
<pymc(gp(gp_submodel.GPEvaluation  
'sm_f_eval' at 0x11796a6d0)>
```

```
>>> sm.f.value
<pymc.gp.Realization.StandardRealization
at 0x10218a210>

>>> sm.f_eval.value
array([ 1.25731172,  1.19401778,
       1.14649327,  1.0654201 ,  0.96870045,
       0.9051113 ,  0.82684987,
       0.73403466,  0.68371817,  0.65486055,
       0.60000131,  0.52193754,
       0.43451894,  0.34834653,  0.28722153,
       0.2158233 ,  0.18256193
       0.66528437,  0.75319381,  0.81996597,
       0.87137425])
```

PyMCmodel.py

```
11
12     # Prior parameters of C-
13     nu = pm.Uniform('nu', 1., 3, value=1.5)-
14     phi = pm.Lognormal('phi', mu=.4, tau=1, value=1)-#
15     theta = pm.Lognormal('theta', mu=.5, tau=1, value=1)-#
16
17     # The covariance dtrm C is valued as a Covariance object.-#
18     @pm.deterministic-
19     def C(eval_fun = gp.matern.euclidean, diff_degree=nu, amp=phi,
20     scale=theta):-
21         return gp.NearlyFullRankCovariance(eval_fun, diff_degree=diff_degree,
22     amp=amp, scale=scale)-#
23
24     # Prior parameters of M-
25     a = pm.Normal('a', mu=1., tau=1., value=1)-#
26     b = pm.Normal('b', mu=.5, tau=1., value=0)-#
27     c = pm.Normal('c', mu=2., tau=1., value=0)-#
28
29     # The mean M is valued as a Mean object.-#
30     def linfun(x, a, b, c):-
31         return a * x ** 2 + b * x + c-
32
33     @pm.deterministic-
34     def M(eval_fun = linfun, a=a, b=b, c=c):-
35         return gp.Mean(eval_fun, a=a, b=b, c=c)-#
36
37     # The actual observation Locations-
38     actual_obs_locs = np.linspace(-.8,.8,4)-#
39
40     if fmesh_is_obsmesh:-
41         o = actual_obs_locs-
42         fmesh = o-
43     else: -
44         # The unknown observation Locations-
45         o = pm.Normal('o', actual_obs_locs, 1000., value=actual_obs_locs)-#
46         fmesh = np.linspace(-1,1,n_fmesh)-#
47
48     # The GP parameters-
```

```
30
31     @pm.deterministic-
32     def M(eval_fun = linfun, a=a, b=b, c=c):-
33         return gp.Mean(eval_fun, a=a, b=b, c=c)-
34
35     # The actual observation Locations-
36     actual_obs_locs = np.linspace(-.8,.8,4)-
37
38     if fmesh_is_observemesh:-_
39         o = actual_obs_locs-
40         fmesh = o-
41     else: _-
42         # The unknown observation Locations-
43         o = pm.Normal('o', actual_obs_locs, 1000., value=actual_obs_locs)-_
44         fmesh = np.linspace(-1,1,n_fmesh)-_
45
46     # The GP submodel-
47     sm = gp.GPSubmodel('sm',M,C,fmesh)-_
48
49     # Observation variance-
50     V = pm.Lognormal('V', mu=-1, tau=1, value=.0001)-_
51     observed_values = pm.rnormal(actual_obs_locs**2,10000)-_
52
53     # The data d is just array-valued. It's normally distributed about
54     # GP.f(obs_x).-
55     d = pm.Normal('d',mu=sm.f(o), tau=1./V, value=observed_values,
      observed=True)-
```

Line: 20 Column: 65

Python



Soft Tabs: 4

C(eval_fun = gp.matern.euclidean, diff_deg...)





GP

Step Method

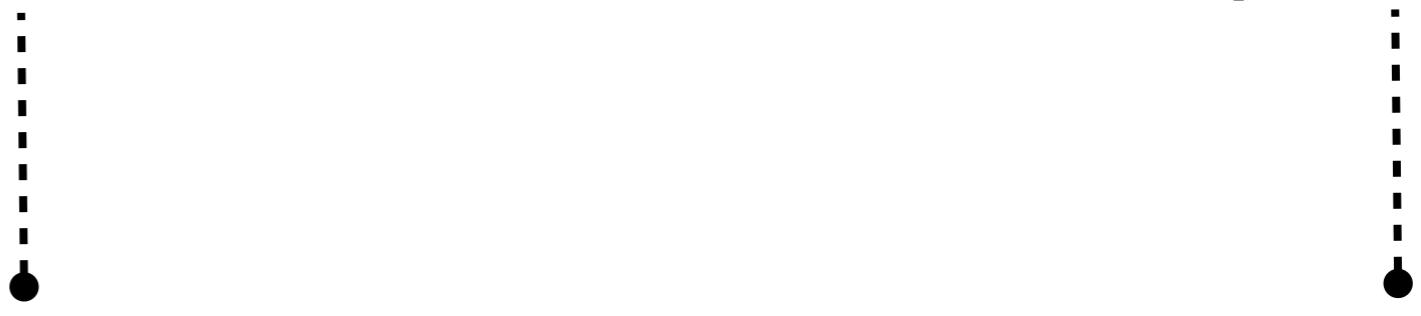
$$\alpha = \frac{p(K|f')p(f'|P')q(P)}{p(K|f)p(f|P)q(P')}$$

Metropolis acceptance rate

$$\alpha = \frac{p(K|f') p(f'|P') q(P)}{p(K|f) p(f|P) q(P')}$$

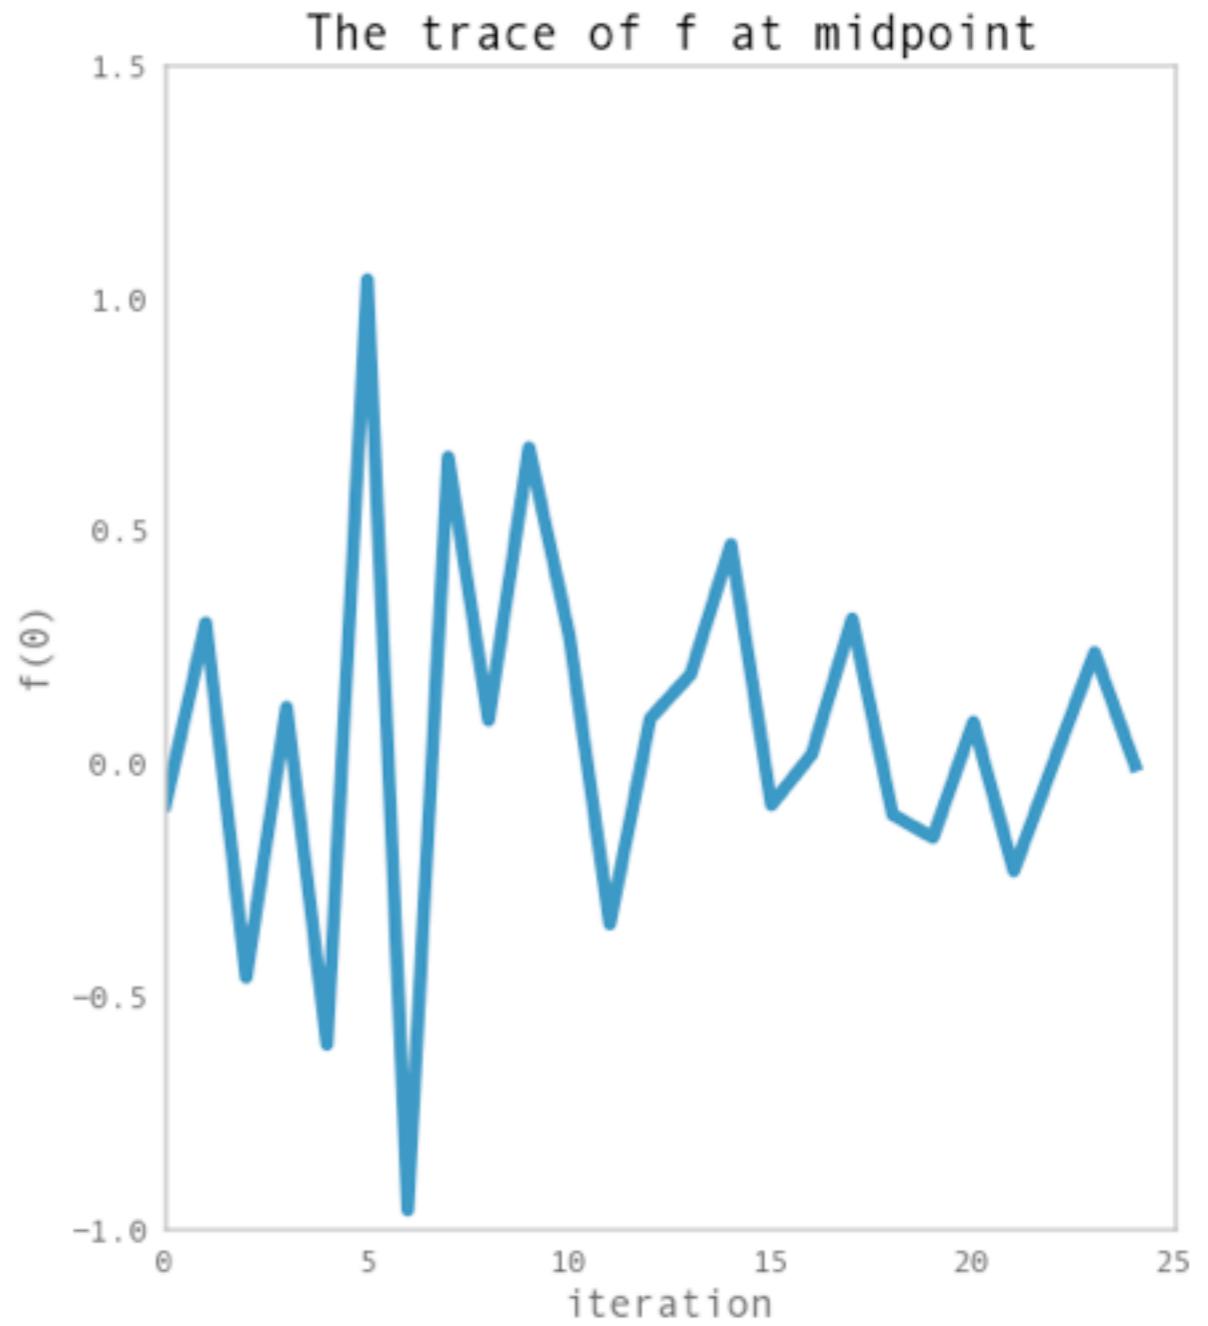
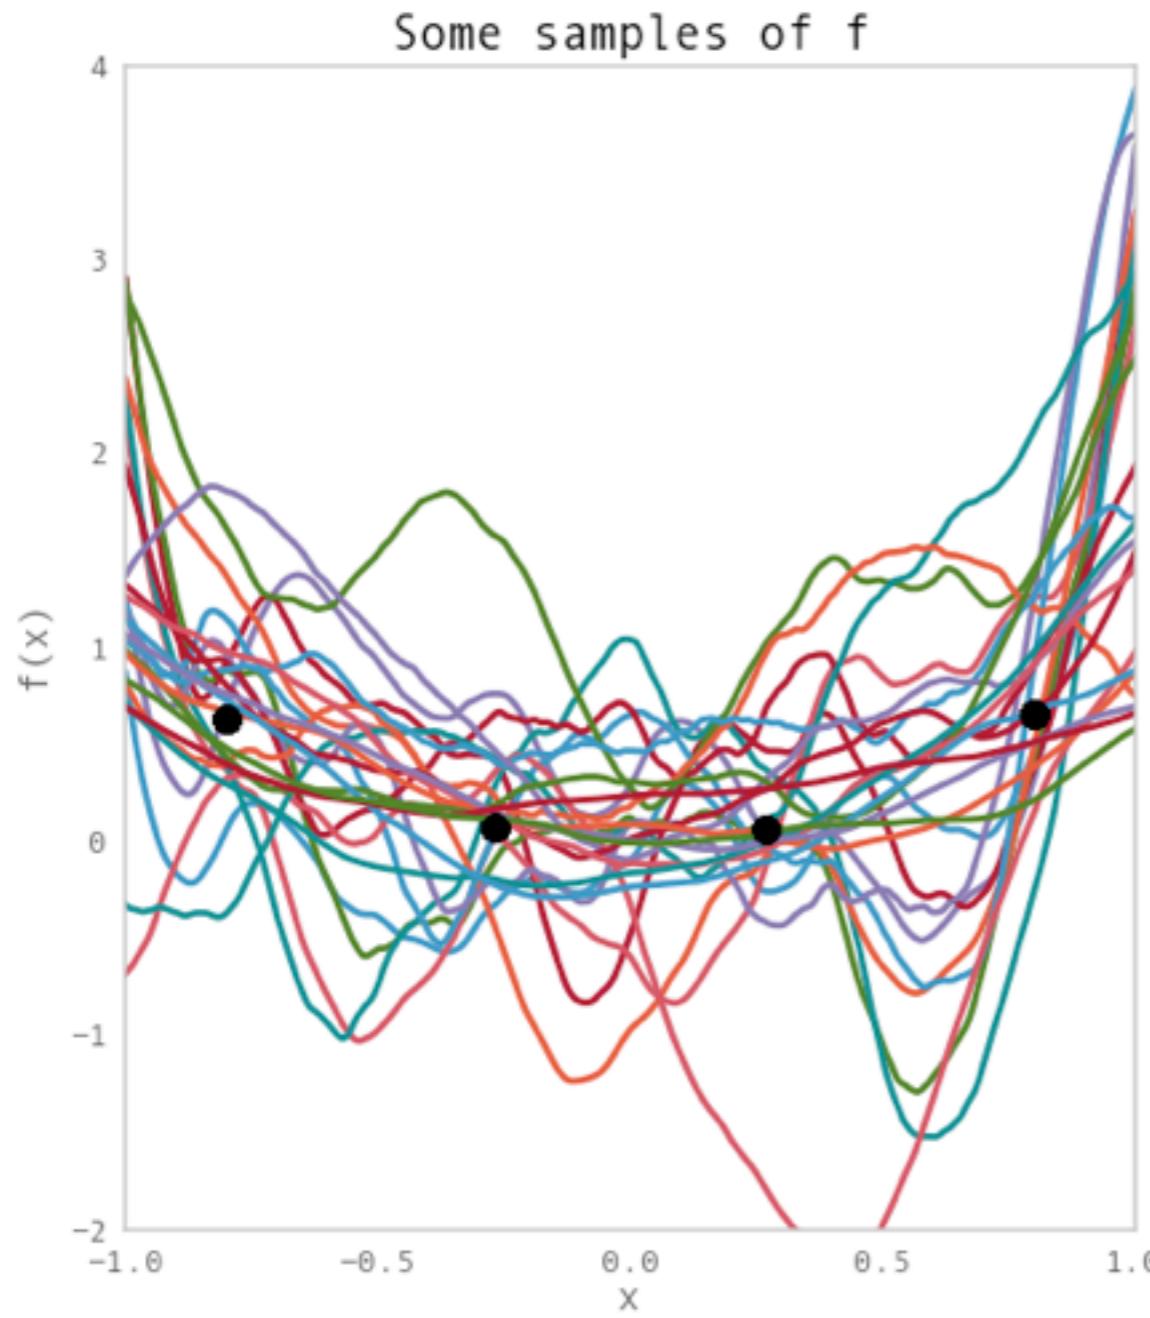
children

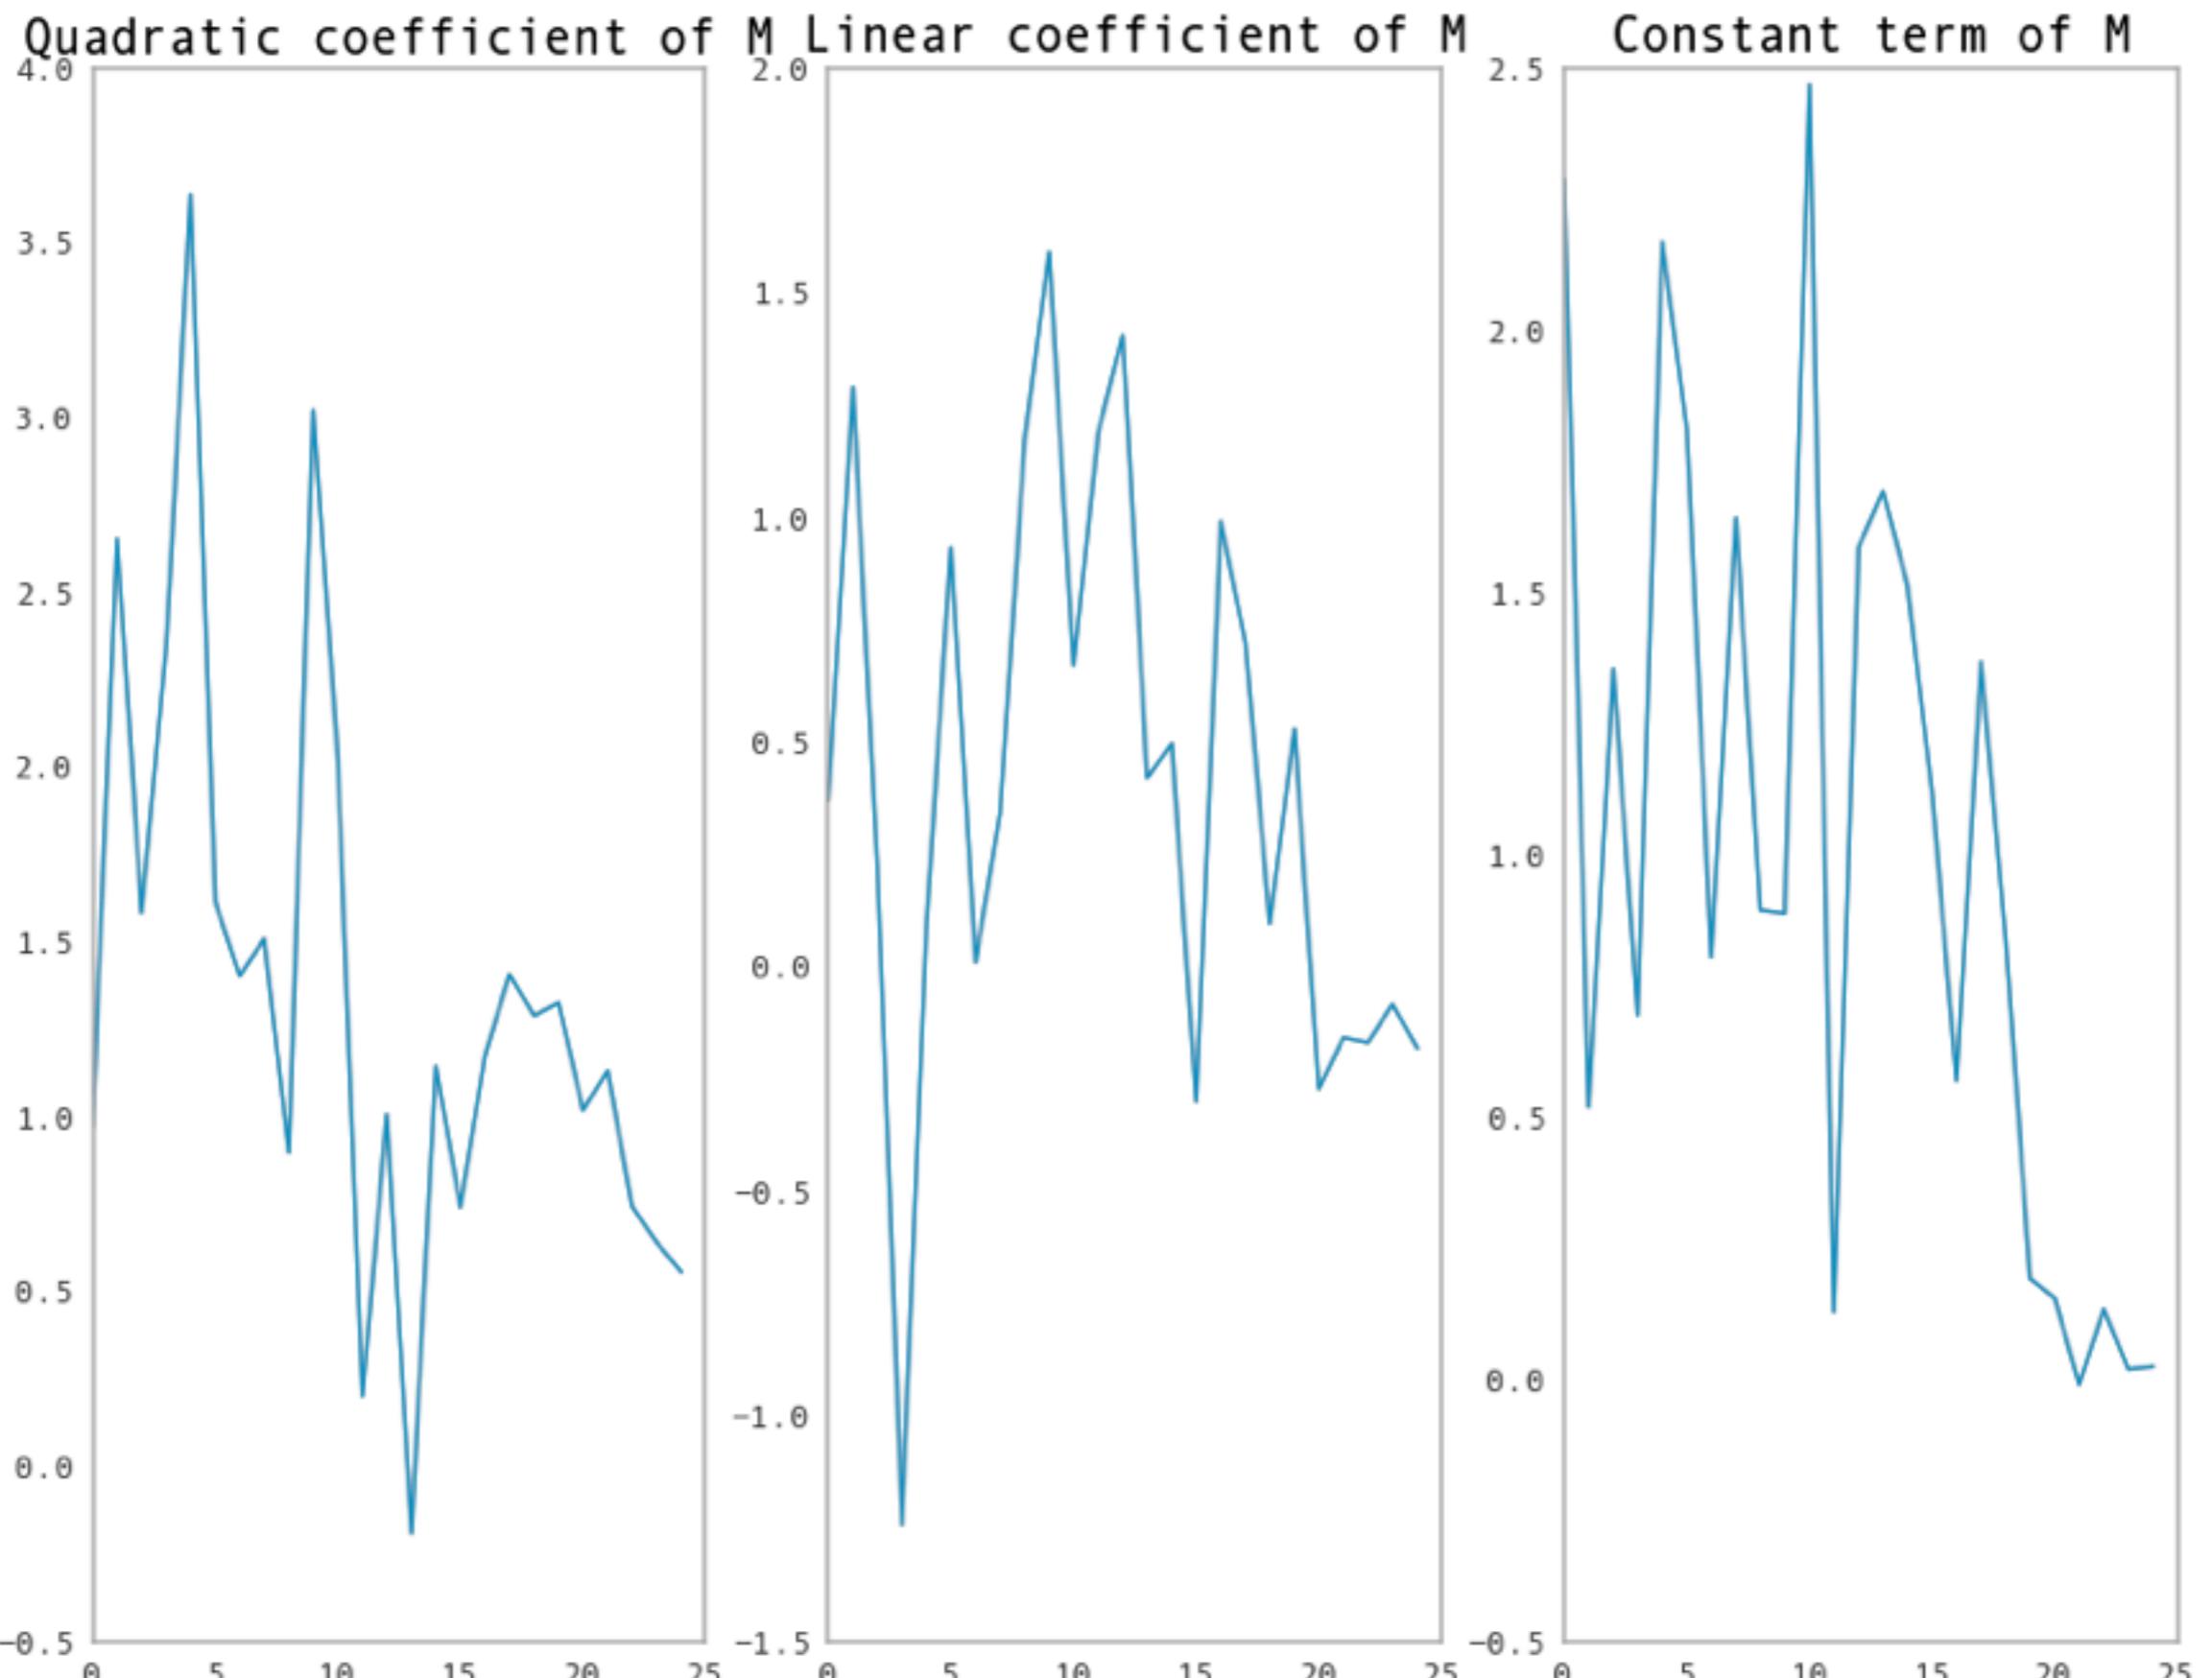
parents

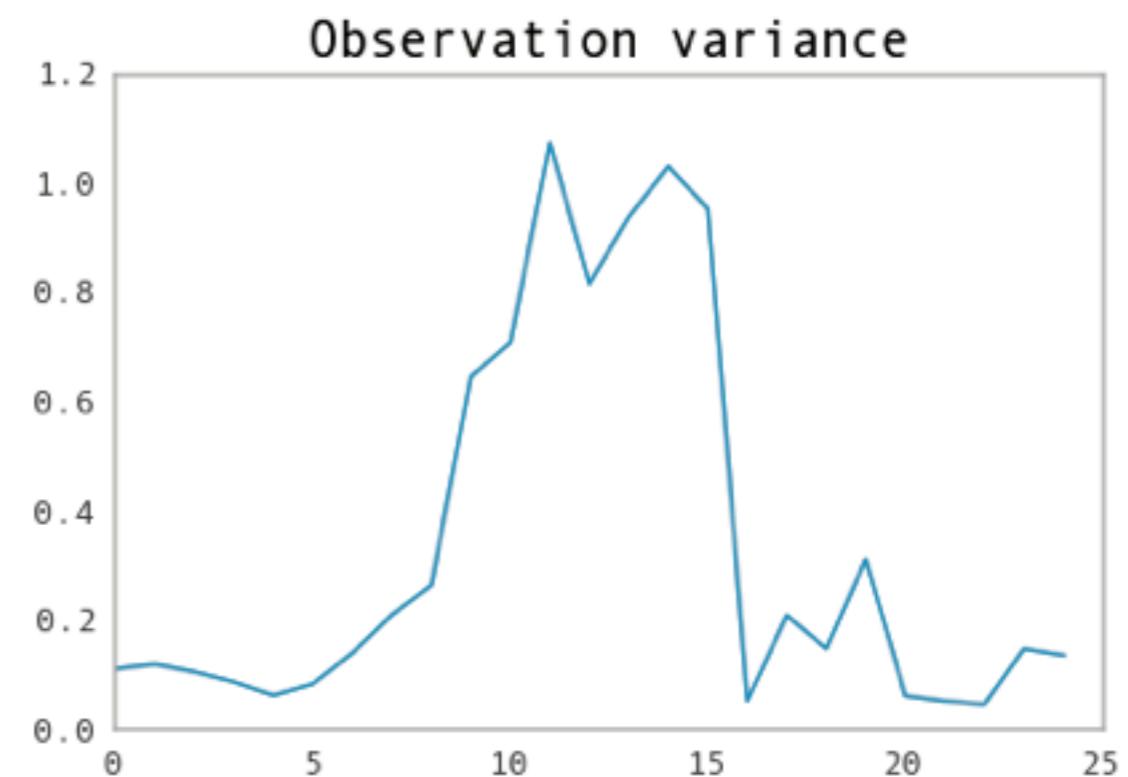
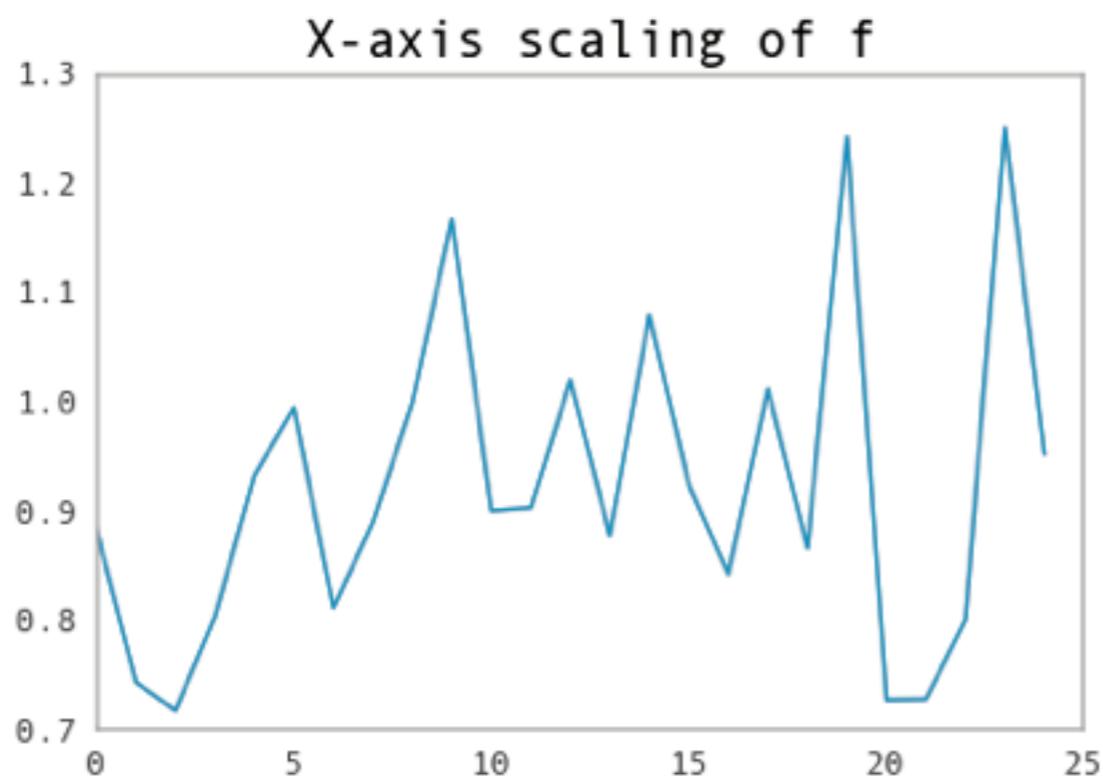
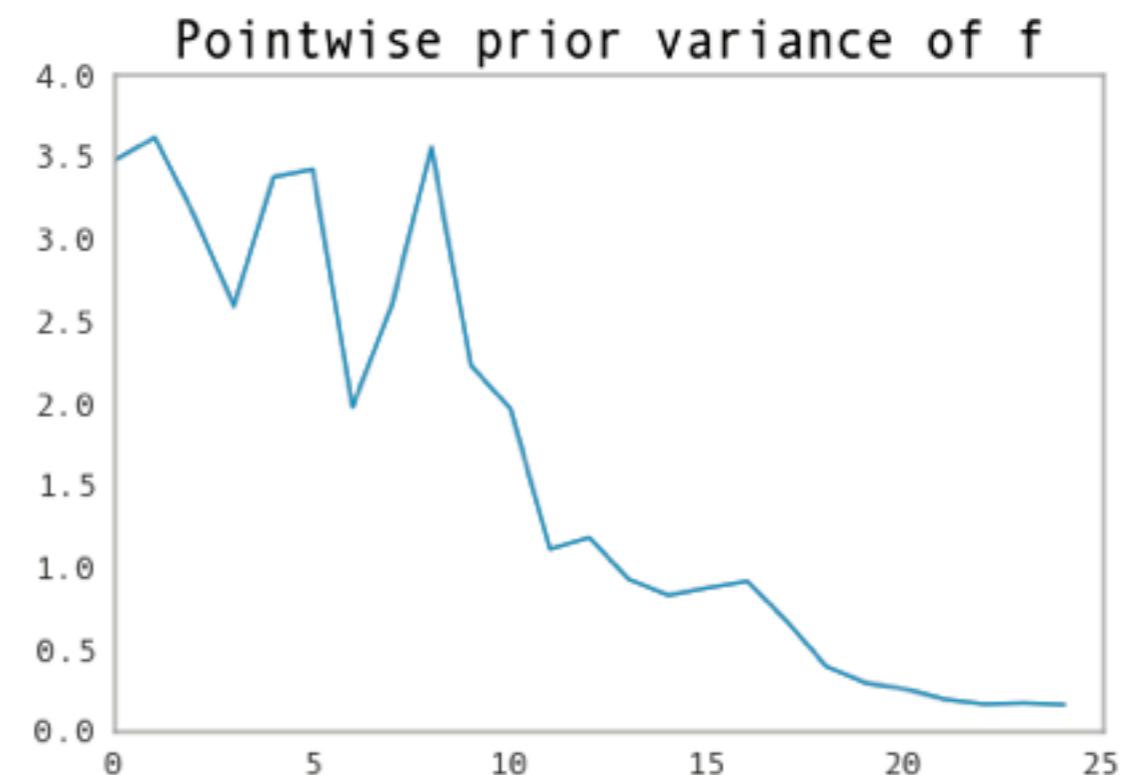
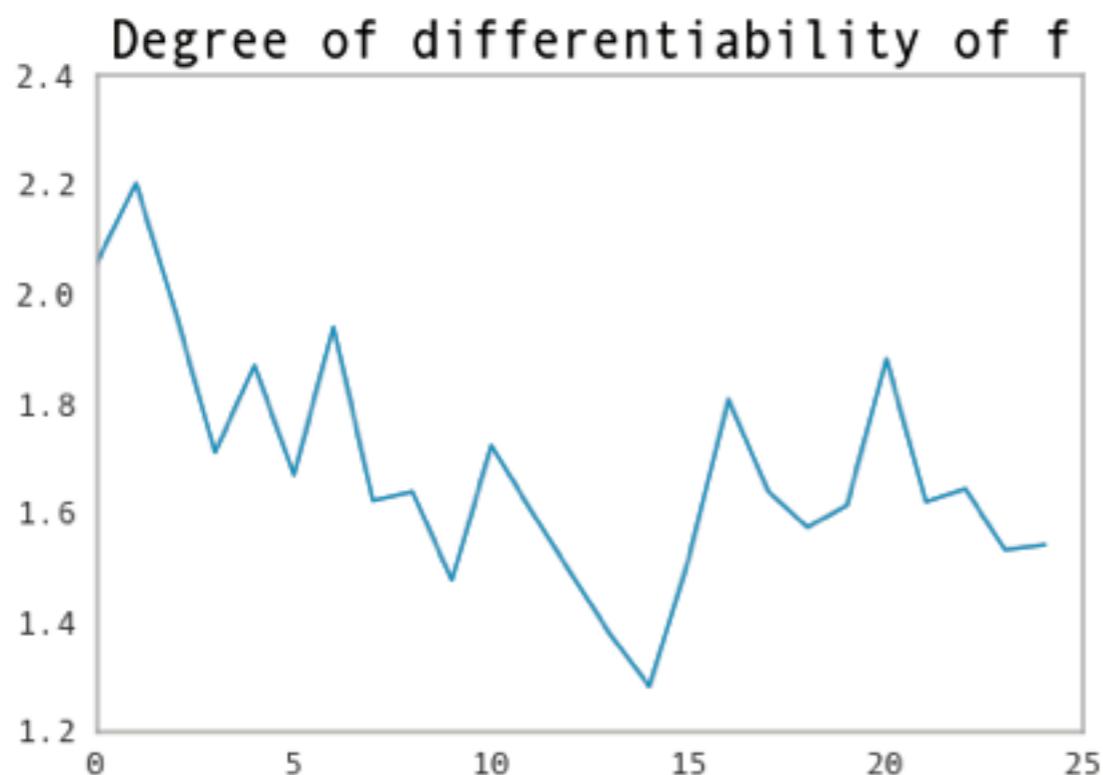


Metropolis acceptance rate

wrap_metropolis_for_gp_parents(some_Metropolis)







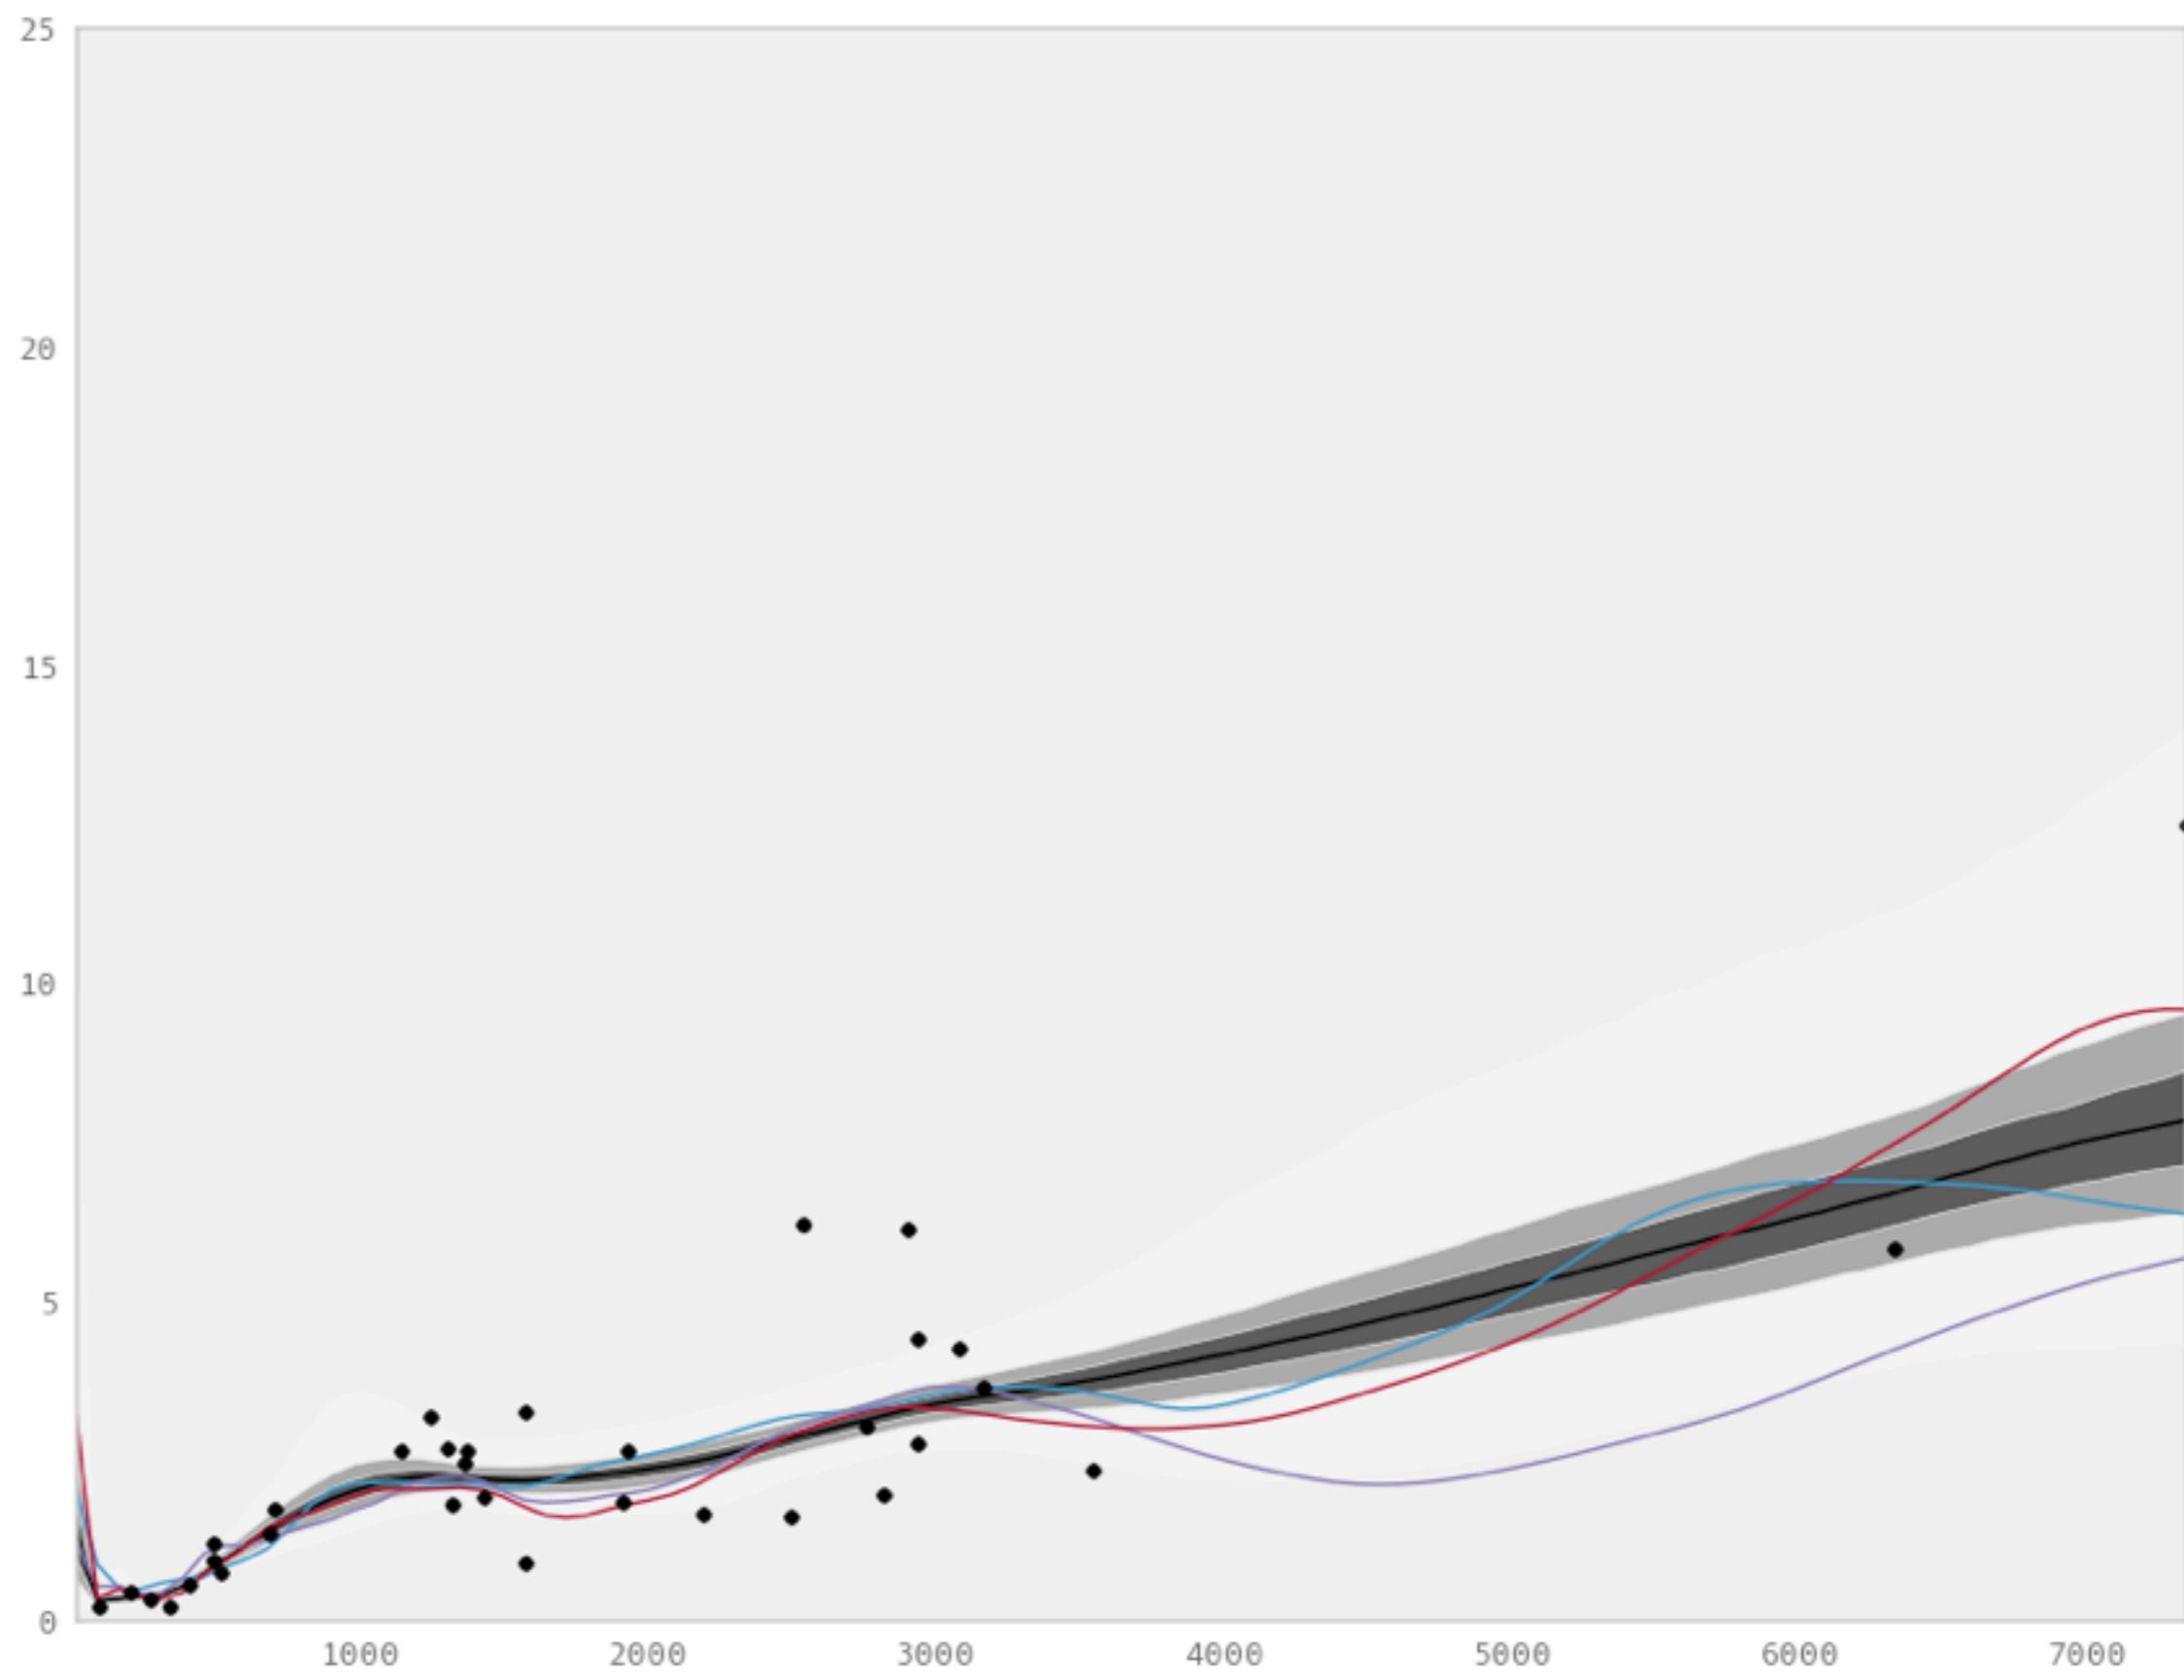


Example

salmon recruitment

Beaverton-holt model

$$n_{t+1} = \frac{R_0 n_t}{1 + n_t / M}$$



Gradient-based Methods



SciPy 2011

 jsalvatier / gradient_samplers Watch[Source](#) [Commits](#) [Network](#) [Pull Requests \(0\)](#) [Issues \(0\)](#) [Graphs](#)[Switch Branches \(2\) ▾](#) [Switch Tags \(0\)](#) [Branch List](#)[samplers for pymc that require gradients — Read more](#)

github.com/jsalvatier/gradient_samplers

[HTTP](#)[Git Read-Only](#)https://github.com/jsalvatier/gradient_samplers [Read-Only access](#)

added error checking to hessian estimation



jsalvatier (author)

March 27, 2011

[commit](#)
[tree](#)
[parent](#)

gradient_samplers /

name	age	message
 gradient_samplers/	March 27, 2011	added error checking to hessian estimation [jsalvatier]
 README	February 13, 2011	first commit [jsalvatier]
 setup.py	March 03, 2011	added better hessian approximation [jsalvatier]

In the sandbox

multiprocessing
multi-model inference
RJMCMC
Dirichlet processes
EM

